


# CMPE 282 Cloud Services

## ***NoSQL***

Instructor: Kong Li

# Content

- CAP
- Eventual Consistency
- NoSQL
- Cassandra 
- HBase
- MongoDB
- CouchDB
- Lab: Cassandra, MongoDB


# Big Data

- Characteristics
  - Volume: MB, GB, TB, PB, etc
  - Variety: different forms or types
  - Velocity: batch, near realtime, realtime
- Search for actionable insights
  - Regardless of structured, semi-structured, or unstructured data
  - Q: How to analyze structured, semi-structured, and unstructured data?
- Evolution: Batch ➔ real time ➔ prediction
- Tools
  - Generic: NoSQL, SQL, search
  - Batch: MapReduce, Hive, Pig, etc.
  - Real time / streaming: Spark (streaming), Storm, etc
  - Machine learning: Mahout, Spark ML, etc
- Q: how to use the right tool for the job?
  - <http://www.slideshare.net/AmazonWebServices/aws-november-webinar-series-architectural-patterns-best-practices-for-big-data-on-aws>



# Brewer's CAP Theorem

- Three desirable properties when designing distributed sys
  - **C**onsistency: All copies have the same value
  - **A**vailability: System can run even if parts have failed
  - **P**artition-tolerant: Survive network partitioning
- It is *impossible to achieve all three* in
  - Async networks: no clock; node makes decision based on msg received and local computation
  - Partially sync networks: each node has local clock; all clocks increase at the same rate; clocks are not synchronized
- *Any two of these three can be achieved*
- Examples
- Large systems will partition at some point → availability or consistency?
  - Traditional DB chooses consistency
  - Most web apps choose availability (exception: order processing, etc.)
- *CAP theorem only matters when there is a partition*

# Eventual Consistency

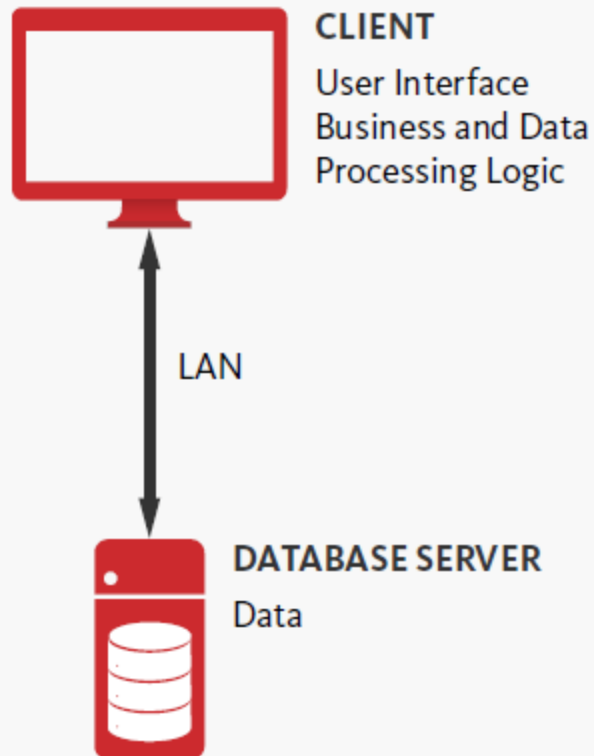
- **Eventual consistency**: when no updates occur for a long period of time, eventually all updates will propagate through the system and all the nodes will be consistent
  - For a given accepted update and a given node, eventually either the update reaches the node or the node is removed from service
  - You may **not** know how long it  may take
- Known as **BASE (Basically Available, Soft state, Eventual consistency)**, as opposed to ACID
  - **Soft state**: copies of a data item may be inconsistent
  - **Eventually Consistent** – copies becomes consistent at some later time if there are no more updates to that data item
- Used by most **NoSQL**
- **Tradeoffs**: consistency, availability, latency/performance

# NoSQL

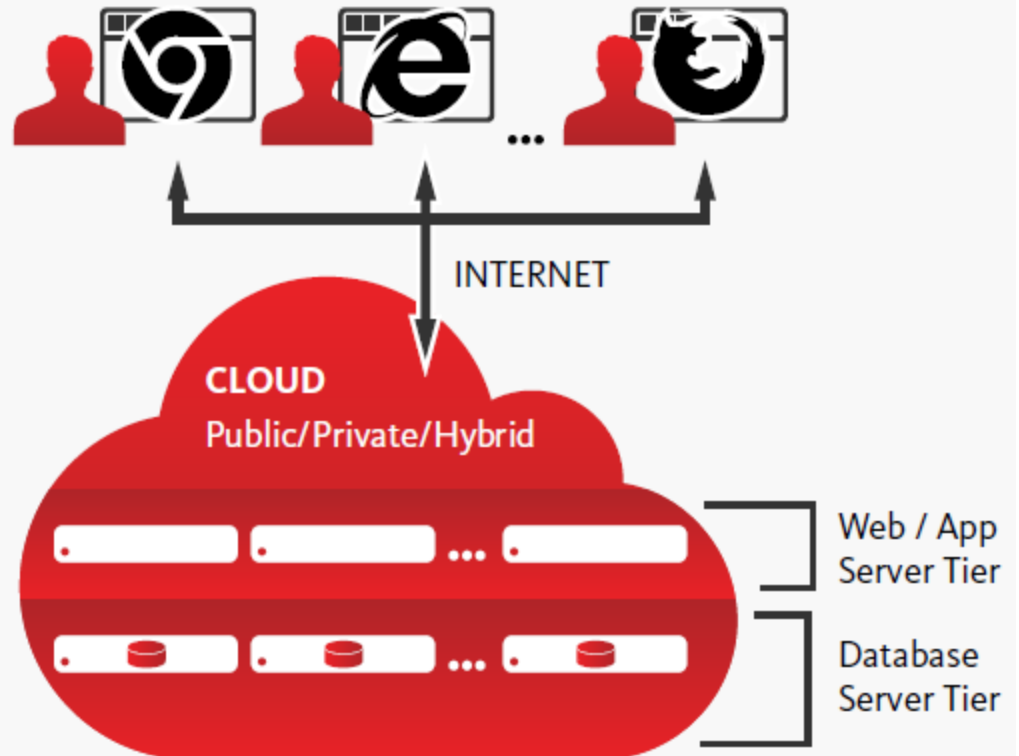
- **Schema-less** data stored as some form of **key-value** pair DB
  - APIs: get(key), set(key, value)
- Simpler functionality 
- Scalability: **scale out**
- **Eventual consistency (BASE)**
- Each addresses a specific set of issues that RDBMS do not
  - none provides a panacea for all issues
- Some gradually brings back functionalities from RDBMS
- “No to SQL” → “Not only SQL”
- Categories 
  - Key-value: AWS DynamoDB, Azure Tables, Google Cloud Datastore, Riak, Redis
  - Column-family/BigTable: Cassandra, HBase, Hypertable
  - Document: MongoDB, CouchDB, AWS DynamoDB
  - Graph: Neo4j

# The Cloud

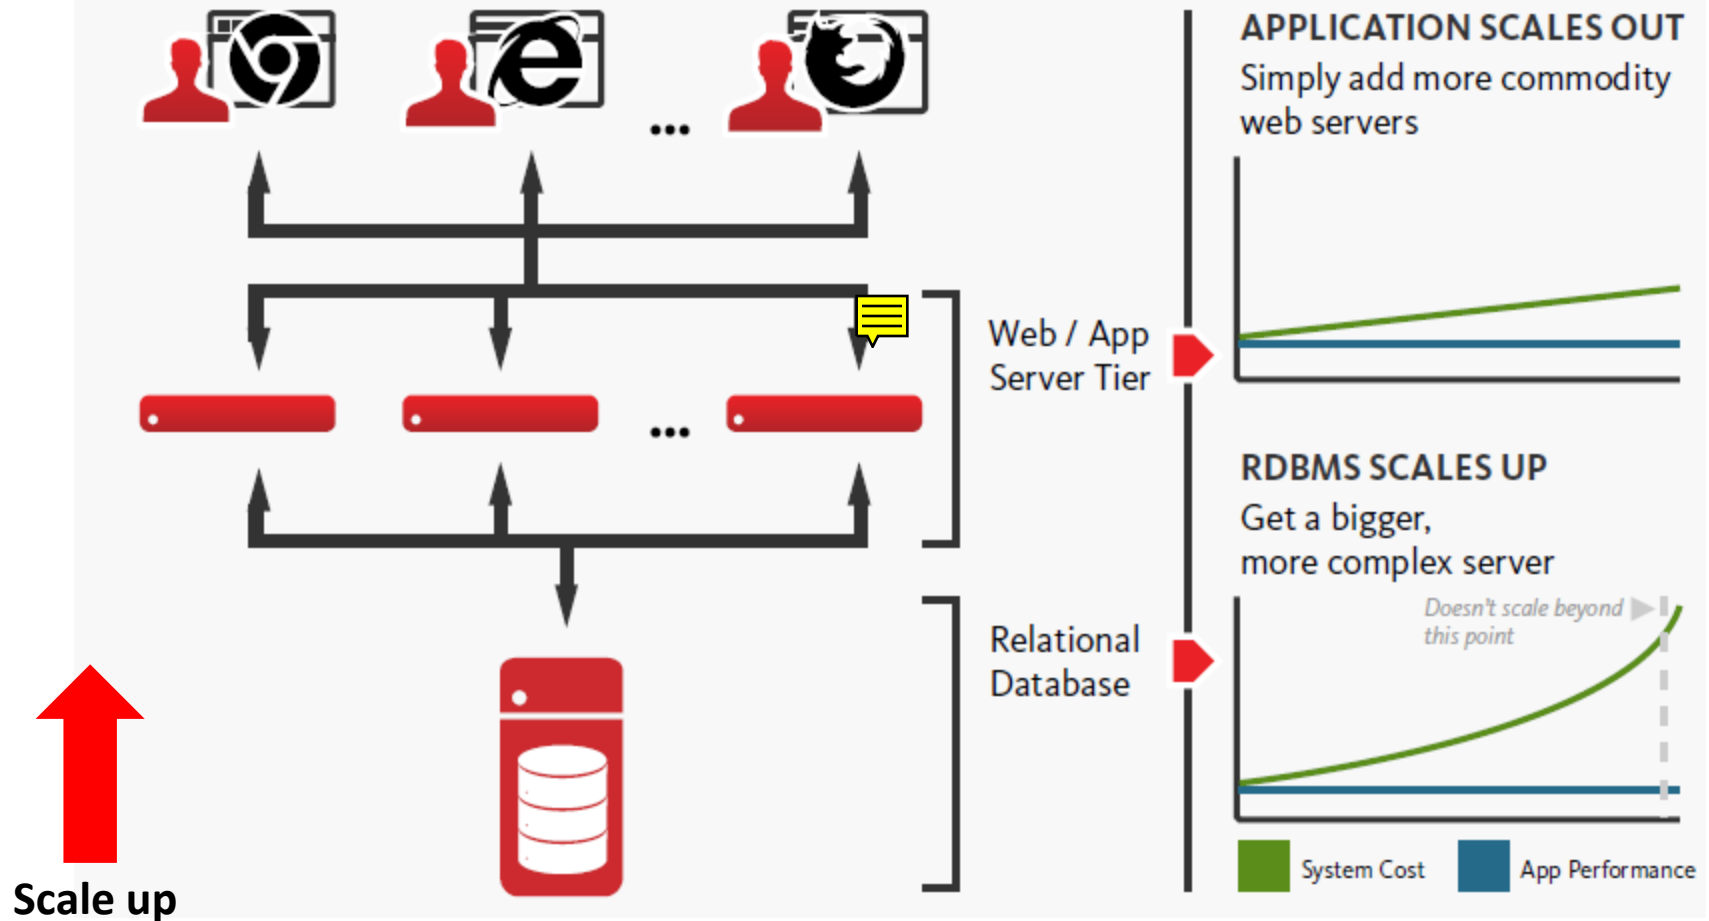
OLD CLIENT/SERVER ARCHITECTURE



NEW 3-TIER INTERNET ARCHITECTURE

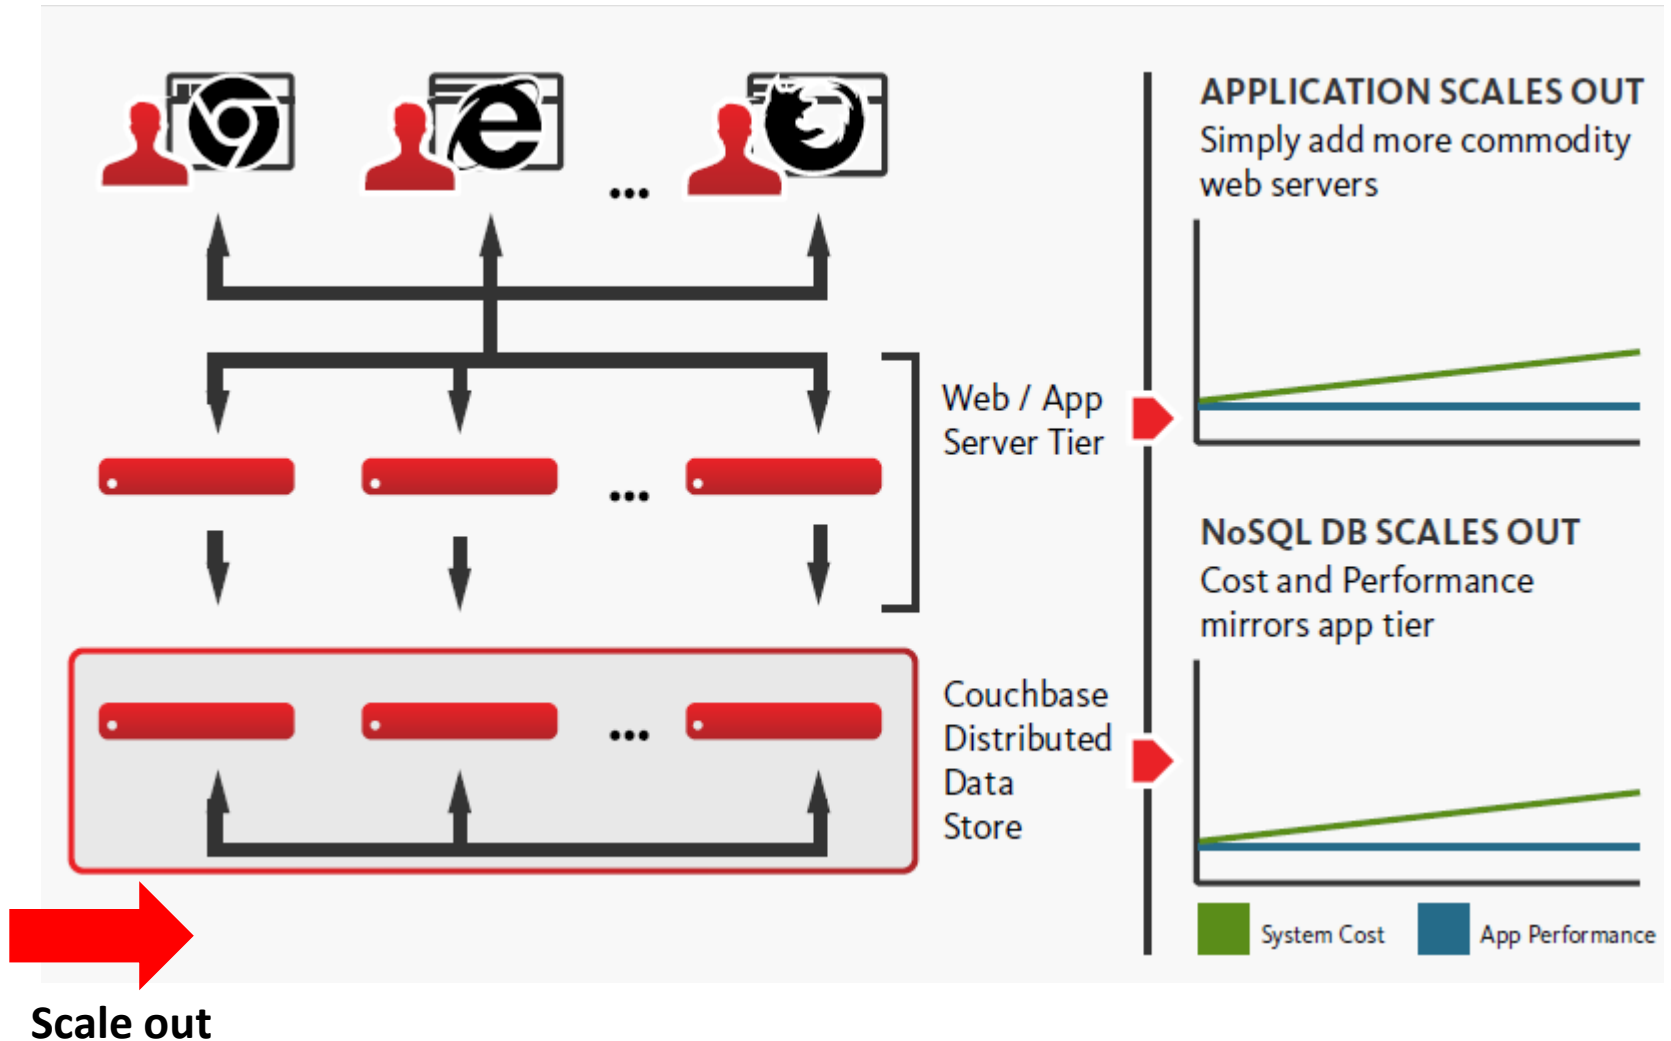


# Scalability of RDBMS





# Scalability of NoSQL



# RDBMS vs NoSQL

	RDBMS	NoSQL
	Schema: tables/columns	Schema-less: key-value pairs
Data	Structured data	Semi- and un-structured data
Process	Requirements → Data model (ERD, etc) → tables	Store data first → (Transform data) → get/put
Data association	Referential integrity, foreign keys	You are on your own
Functionality	Complex	Simple (→ Better performance?)
Properties	ACID	BASE
Pros	(strong) consistency Transactions Joins	Schema-less Scalability, Availability Performance
Cons	Scalability Availability	Eventual consistency Tricky join No transactions

# NoSQL – Key-value Pair

- Simple architecture
  - a unique identifier (key) maps to an obj (value)
  - DB itself does not care about obj
- Good for
  - Simple data model
  - Scalability
- Bad for
  - “Complex” datasets
- Examples:
  - AWS DynamoDB, Azure Tables, Google Cloud Datastore, Riak, Redis, Azure CosmosDB

# NoSQL – Column-family/BigTable

- Key-value pair + grouping
  - Key – {a set of values}
  - E.g., time series data from sensors, languages of web page
- Good for
  - More complex data set (than simple key-value pair)
  - Scalability
- Examples
  - Cassandra
  - HBase
  - Hypertable
  - Azure CosmosDB

# NoSQL – Document

- DB knows the nature of the data
  - Document – JSON
  - Need index to improve performance
- Scalability – clustering, replication, some w/ partition-tolerant
- Good for
  - Systems already using document-style data
  - Scalability
- Examples: MongoDB, CouchDB, AWS DynamoDB, Azure CosmosDB

# NoSQL – Graph

- How data is related and what calculation is to be performed
- Usually also has “transaction”
- Good for
  - Interconnected Data and non-tabular
    - Geospatial problems, recommendation engine, network analysis, bioinformatics
  - “Closeness” of relationship
    - How vulnerable a company is to “bad news” for another company
- Bad for
  - Simple tabular data
- Examples: Neo4j, Azure CosmosDB

# Cassandra

- Scalable **HA** NoSQL, key-value pair + column family ( $\approx$  table in RDBMS)
- Cluster spanning multiple data centers (DCs): linear scalability
- Data model: a partitioned row store
  - Rows are organized into tables
  - Partition: **consistent hashing** - decides storage node ( $\approx$  AWS DynamoDB)
    - adding/removing nodes in cluster only reshuffles keys among *neighboring* nodes
  - **Tunable consistency: decided by client apps, per operation**
    - Consistency level: 1, 2, 3, **quorum**, all, any, etc.
    - **W: sent to all replica, R: decided by consistency level**
- Replication and multi-DC replication
  - asynchronous **masterless** (peer-to-peer) replication
  - Rack aware, rack unaware, datacenter aware: performance vs availability
- Cassandra Query Language (CQL): SQL like
  - No joins and subqueries
  - materialized views (i.e. pre-computed results of queries)
- Hadoop integration, w/ MapReduce support

controls how many responses the coordinator waits for before responding to the client

quorum  $\approx$  AWS DynamoDB

# NRW: Consistency vs Availability

- R and W: **read quorum**  $Q_R$  and **write quorum**  $Q_W$

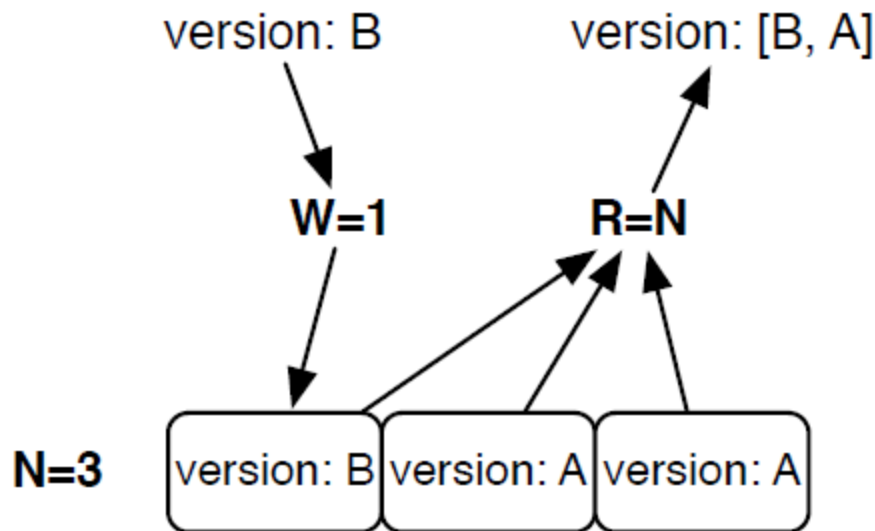


Figure 11—Consistency by reads:  $W=1, R=N$

- Both are special case of consistency by quorum




Figure 10—Consistency by writes:  $W=N, R=1$



# NRW: Consistency vs Availability (cont'd)

- R and W: **read quorum**  $Q_R$  and **write quorum**  $Q_W$

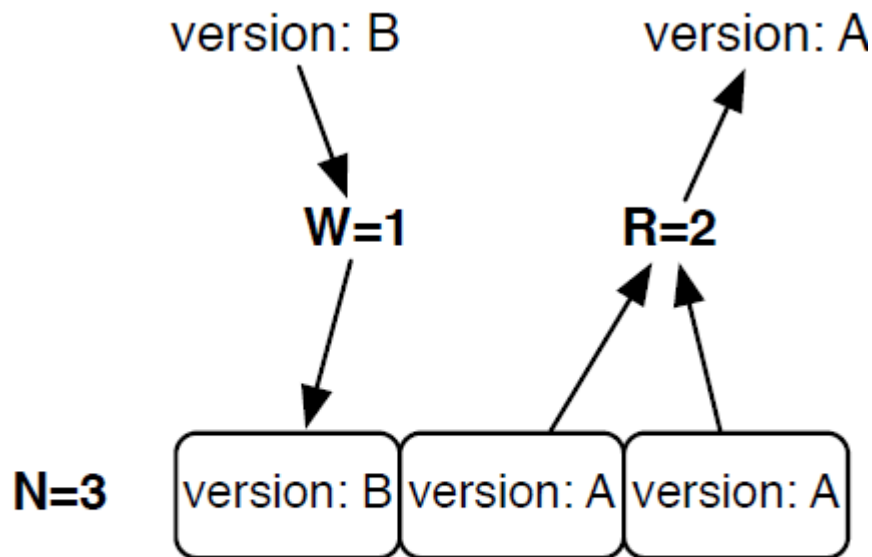


Figure 9—Eventual consistency:  $W+R \leq N$

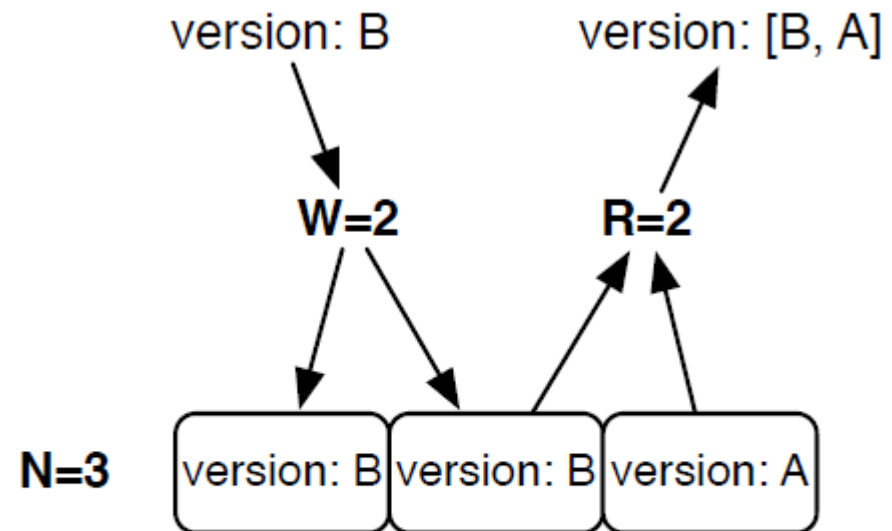


Figure 12—Consistency by quorum:  $W+R > N$

AWS DynamoDB: R – strong/eventual consistent; W – unconditional/conditional  
Cassandra: tunable consistency; per operation  
Riak: tunable on a per-bucket, or per-request, basis

# HBase

- Column-oriented DB based on Google BigTable
- “table”, “row”, and “column”: different from RDBMS
- Distributed, (timestamp) versioned, NoSQL DB on top of Hadoop and HDFS
- Cluster: scale out, failover
- CRUD: shell
- **Strong consistent reads and writes: *not* eventual consistency**
  - All operations are **atomic at the row level**, regardless of # of columns
  - Write ahead logging (WAL): protect from server failure

	row keys	column family "color"	column family "shape"
row	"first"	"red": "#F00" "blue": "#00F" "yellow": "#FF0"	"square": "4"
row	"second"		"triangle": "3" "square": "4"

first/color:red = #F00

# HBase (cont'd)

- Region-based storage architecture:
  - Region: a set of rows, [starting-row-key, ending-row-key), w/ a region server
    - Regions never overlap
  - *Auto sharding* of tables:
    - tables distributed on cluster via regions, regions automatically split & re-distributed
- HBaseMaster: **master** node
  - Assign regions to region servers
  - Automatic failover between *region servers*
  - Upon HBaseMaster failure, one region server takes over as master
- HDFS and HBase: rack-aware, data replication within & between DC racks
- Strength
  - Scale out architecture, rack-aware replication
  - Built-in versioning
- weakness
  - Complexity: + Hadoop + etc
  - Designed for Big, not for simple/small problem: minimum 5 nodes

# MongoDB

- Scalable distributed document DB
  - Nested JSON document (binary form of JSON – BSON)
  - A document: a JSON obj w/ key-value pairs  $\approx$  row in RDBMS
- CRUD: JavaScript, **ACID at document level**
- Emulate SP in RDBMS: executing a server-side function (.js)
- No joins: still allow to retrieve data via relationship
- B-tree based Indexing: performance
- map(), reduce(), and finalize()

\$lookup: left outer join for pipelined data aggregation

12 bytes

- globally unique
- Can be override

## JSON syntax:

- Key-value: {"key": "value"}
- , separated: {...}, {...}
- {} for obj: {"key": {"k2": "v2"}}
- [] for array: {"key": [...], {...]}

```
> printjson( db.towns.findOne({"_id" : ObjectId("4d0b6da3bb30773266f39fea")) )
```

```
{
  "_id" : ObjectId("4d0b6da3bb30773266f39fea"),
  "country" : {
    "$ref" : "countries",
    "$id" : ObjectId("4d0e6074deb8995216a8309e")
  },
  "famous_for" : [
    "beer",
    "food"
  ],
  "last_census" : "Thu Sep 20 2007 00:00:00 GMT -0700 (PDT)",
  "mayor" : {
    "name" : "Sam Adams",
    "party" : "D"
  }
}
```

Collection

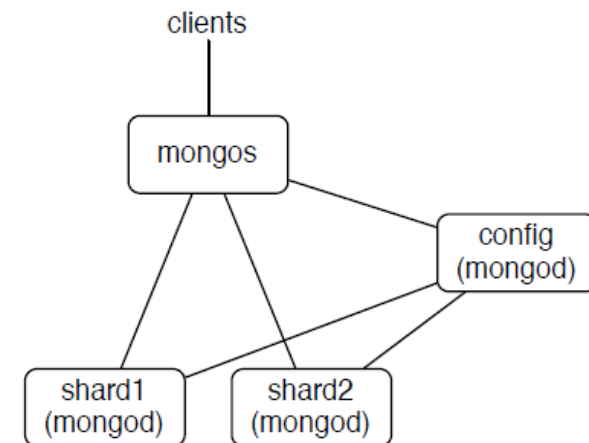
Database

Identifier

Document

# MongoDB (cont'd)

- **Replica Sets**
  - One server as primary (**master**), others as secondary (w/ replicated data)
  - Client only reads from or writes to master
  - When master is down, one of secondary server becomes primary (master)
  - *majority of nodes that can still communicate make up the network*
    - Quorum - require odd number of nodes in replica set, for consistency
- **Auto sharding** – range-based, hash-based, user-defined
  - mongod (config server): track how sharding is done
  - mongos: query router - entry point for clients
  - Data distribution: splitting, balancing
  - App transparent
- GridFS: distributed file system
- Strength
  - Scale out, sharding, replication, easy to use, similar query capability as RDBMS (w/o joins)
- Weakness: Too flexible? simple typo could cause debugging headache



# CouchDB

- JSON- and REST-based document-oriented DB
- Each document
  - A set of key-val, val is any JSON structure nested to any depth (like MongoDB)
  - Unique immutable `_id`: auto-gen or explicit assigned per doc
  - `_rev`: revision per doc, starting from 1, **auto-modified after doc changes**
  - **Each update/delete must specify *both* `_id` and `_rev` (why?)**
  - Can be nested (like MongoDB)
- CRUD: Futon web interface
  - RESTful (GET: read, POST: create, PUT: update or create, DELETE: delete)
  - Update: overwrite the *entire* doc (unlike MongoDB - modify *in place*)
- **No transaction, no locking**
- View: ordered list of key-value pairs, indexed, key and val can be any JSON
  - Node.js (server-side JavaScript); can be persisted as a design document
  - **Content generated by mapreduce functions incrementally** (~ materialized view)

# CouchDB (cont'd)

- **Changes API** – monitor *any* changes to documents
  - polling, long-polling, and continuous
  - Can develop non-blocking event-driven client apps
- Replication
  - All or nothing; no sharding
  - **Multi-master replication: HA**
  - Conflict resolution: same `_id` + `_rev` – all nodes choose the same winner
- Strength
  - Variety of deployment: smart phone to DC
  - Indexed view built by mapreduce incrementally
  - Multi-master replication, HA
- Weakness
  - Mapreduce-based view: limited when compared with views in RDBMS
  - All or nothing replication; no sharding


# NoSQL Pros and Cons

- Pros
  - schema-less
  - scalability (scale out), availability
  - performance
  - simpler change mgmt (“agile”, faster dev)
- Cons
  - schema-less
  - eventual consistency
  - tricky (server-side) join
  - No transaction
- Trade-offs (w/ RDBMS): functionality, availability, consistency, durability, scalability, performance
- To SQL or NoSQL?
  - It depends; likely co-exist
  - both provide features the other doesn't
    - RDBMS: valuable where data consistency is required



# CAP Again

- MongoDB and HBase
  - Consistent and partition-tolerant (CP)
  - During network partition, unable to respond is possible (availability)
- Cassandra
  - Availability and partition-tolerant (AP)
  - Tunable consistency; per operation
- CouchDB
  - Availability and partition-tolerant (AP)
- Some DB can be configured to modify the level/degree
  - MongoDB can be CA
  - CouchDB can be CP
- Values of N, R, W
  - AWS Dynamo: R - strong/eventual consistent; W – unconditional/conditional
  - Cassandra: tunable consistency
  - Riak: tunable on a per-bucket, or per-request, basis



# Lab: Cassandra

- Server v 3.x or 2.2.4
  - [cassandra.apache.org](http://cassandra.apache.org)
  - Mac, Windows: [www.planetcassandra.org/cassandra](http://www.planetcassandra.org/cassandra)
    - set “cdc\_raw\_directory” in conf/cassandra.yaml
  - Install:  
[cassandra.apache.org/doc/latest/getting\\_started/installing.html](http://cassandra.apache.org/doc/latest/getting_started/installing.html)
- CQL: [cassandra.apache.org/doc/latest/cql/index.html](http://cassandra.apache.org/doc/latest/cql/index.html)
  - Shell: cqlsh
- Java client driver: [downloads.datastax.com/java-driver](http://downloads.datastax.com/java-driver)
- Java client
  - Tutorial: [academy.datastax.com/demos/getting-started-apache-cassandra-and-java-part-i](http://academy.datastax.com/demos/getting-started-apache-cassandra-and-java-part-i)
  - GetStarted.java: [gist.github.com/beccam/06c3283e5ee4a480a555](https://gist.github.com/beccam/06c3283e5ee4a480a555)

# Lab: Cassandra (cont'd)

- Start server (console mode): `cassandra -f`
- Start CQL shell: `cqlsh`

Why does it look like  
RDBMS/SQL?

`$ cqlsh`

`> desc keyspaces;`

`> CREATE KEYSPACE mydb WITH  
    replication={'class': 'SimpleStrategy',  
                  'replication_factor': '1'};`

`> desc mydb ;`

Tunable consistency

`> use mydb ;`

`> CREATE TABLE users (`

`firstname text,`

`lastname text,`

`age int,`

`email text,`

`city text,`

`PRIMARY KEY (lastname));`

Supported data  
types include set,  
tuple, list, map, etc.

`> desc table users;`

`> INSERT INTO users (firstname, lastname, age,  
    email, city) VALUES ('John', 'Smith', 46,  
    'johnsmith@email.com', 'Sacramento');`

`> INSERT INTO users (firstname, lastname, age,  
    email, city) VALUES ('Jane', 'Doe', 36,  
    'janedoe@email.com', 'Beverly Hills');`

`> SELECT * FROM users;`

`> SELECT * FROM users WHERE lastname= 'Doe';`

`> UPDATE users SET city= 'San Jose' WHERE  
    lastname= 'Doe';`

`> DELETE from users WHERE lastname = 'Doe';`

# Lab: Cassandra (cont'd)

\$ nodetool flush mydb

Convert memtables to sstables

\$ sstabledump data/mydb/users-14c79820638311e7bf210b60b8eb8b8b/mc-1-big-Data.db

```
[
  {
    "partition" : {
      "key" : [ "Doe" ],
      "position" : 0
    },
    "rows" : [
      {
        "type" : "row",
        "position" : 17,
        "liveness_info" : { "tstamp" : "2017-07-08T02:14:40.267Z" },
        "cells" : [
          { "name" : "age", "value" : "36" },
          { "name" : "city", "value" : "San Jose", "tstamp" : "2017-07-08T02:15:
45.805Z" },
          { "name" : "email", "value" : "janedoe@email.com" },
          { "name" : "firstname", "value" : "Jane" }
        ]
      }
    ]
  },
  {
    {
```

# Lab: MongoDB

- Server: community edition
  - [www.mongodb.org/downloads](http://www.mongodb.org/downloads)
  - Install: [docs.mongodb.com/manual/installation/](http://docs.mongodb.com/manual/installation/)
  - Case sensitive
- Shell: [docs.mongodb.org/getting-started/shell/client/](http://docs.mongodb.org/getting-started/shell/client/)
- Java client driver
  - [oss.sonatype.org/content/repositories/releases/org/mongodb/mongo-java-driver/3.4.2/mongo-java-driver-3.4.2.jar](http://oss.sonatype.org/content/repositories/releases/org/mongodb/mongo-java-driver/3.4.2/mongo-java-driver-3.4.2.jar)
  - [mongodb.github.io/mongo-java-driver](http://mongodb.github.io/mongo-java-driver)
- Java client:
  - [mongodb.github.io/mongo-java-driver/3.4/driver/getting-started/quick-start/](http://mongodb.github.io/mongo-java-driver/3.4/driver/getting-started/quick-start/)

# Lab: MongoDB (cont'd)

- Start server on Ubuntu: `sudo service mongod start`
- Start server on windows: Windows: `mkdir /data/db/; mongod.exe`
- Start client: `mongo`

```
$ mongo
```

```
> show dbs
```

```
> use mydb
```

```
> db
```

```
> db.inventory.insert(
  {item: "ABC1",
   details: {model: "14Q3", madeby: "XYZ
             Inc"},
   stock: [ { size: "S", qty: 25 }, { size: "M",
             qty: 50 } ],
   category: "clothing"
})
> show collections
> db.inventory.find()
```

```
> db.inventory.insert({ item: "ABC2", ...})
> db.inventory.find( {item: "ABC2"} )
> db.inventory.count()
> db.inventory.update({ item: "ABC2" }, {
    item: "ABC2", category: "cookware" })
> db.inventory.remove({item : "ABC"})
> db.help()
> db.inventory.help()
> db.inventory.remove({})
> show collections
> db.inventory.drop()
> show collections
> db.dropDatabase()
```

# References

- Seven Databases in Seven Days
- Cassandra: <http://cassandra.apache.org/>
- Cassandra then and now
  - <http://docs.datastax.com/en/articles/cassandra/cassandrathenandnow.html>
- MongoDB: <https://www.mongodb.org/>
- CouchDB: <http://couchdb.apache.org/>
- Couchbase: <http://www.couchbase.com/>
- Riak: <http://basho.com/products/#riak>
- HBase: <http://hbase.apache.org/>
- 7 hard truths about the NoSQL revolution
  - <http://www.infoworld.com/article/2617405/nosql/7-hard-truths-about-the-nosql-revolution.html>
- AWS DynamoDB:  
<http://docs.aws.amazon.com/amazondynamodb/latest/developerguide/WorkingWithItems.html#WorkingWithItems.ReadingData>
- [https://d0.awsstatic.com/whitepapers/AWS\\_Comparing\\_the\\_Use\\_of\\_DynamoDB\\_and\\_HBase\\_for\\_NoSQL.pdf](https://d0.awsstatic.com/whitepapers/AWS_Comparing_the_Use_of_DynamoDB_and_HBase_for_NoSQL.pdf)