


CMPE 282 Cloud Services

Cloud-Native Application

Design Patterns

Instructor: Kong Li

Content

- What and why
 - Coffee Shop
 - Decomposition
 - Workload
 - Data/state
 - Component refinement
 - Elasticity and Resiliency
 - Composite Cloud Applications
- 
- The diagram consists of three blue brackets on the right side of the list, each grouping a set of items and labeled with a number. The first bracket groups 'What and why' and 'Coffee Shop' and is labeled '1'. The second bracket groups the sub-items of 'Coffee Shop' and is labeled '2'. The third bracket groups 'Composite Cloud Applications' and is labeled '3'.

Coffee Shop App Design (3)

- Decomposition: How to distribute Application Functionality?
 - Distributed App
- Work load: What workload do components experience?
 - Static
 - Periodic
 - Once-in-a-lifetime
 - Unpredictable
 - Continuously Changing
- **Data (State): Where does the application handle state?**
 - **Stateful**
 - **Stateless**
 - **Strict consistency**
 - **Eventual consistency**
 - **Data Abstractor**
- Component Refinements: How are components implemented?
 - Message-oriented Middleware
 - User Interface Component
 - Processing Component
 - Batch Processing Component
 - Multi-component Image
- Elasticity and Resiliency
 - Elastic Load Balancer
 - Elastic Queue
 - Node-based Availability
 - Environment-based Availability
 - Watchdog


State = session state + app state

- **Session State** - *State of a client's interaction with an app*
 - Ex: customer's shopping card of an online store
- **App State** - *Data handled by an app*
 - Ex: customers shipping information stored by an app



Stateful Component

- Intent: Multiple instances of a scaled-out app component **synchronize their internal state** to provide a unified behavior




- Any storage offerings: block/blob storage, RDBMS, NoSQL, etc.

- Solution: The internal state maintained by app component instances is replicated among all instances
 - A config file stored centrally or config send by clients w/ every request
 - Internal state replication: **strict consistency** vs **eventual consistency**
 - Tradeoffs: consistency, availability, performance, durability



Stateless Component

- Intent: State is handled **external of app components** to ease their scaling-out and to make the app more tolerant to component failures
 - Isolation of component/resource failure




- Web services
- REST
- Shopping baskets in Amazon

- Solution: no internal state for app components. App state and config is stored externally in Storage Offerings or provided to the component with each request
 - Ex: Keep state on client-side in Web app, REST
 - Each request can be handled by arbitrary server




Strict Consistency

- Intent: Data is stored at different locations to improve response time and to avoid data loss in case of failures while **consistency of replicas is ensured at all times**



$$r + w > n$$
$$2 * w > n$$



number of replicas (n) = 2
replicas accessed to write (w) = 2
replicas accessed to read (r) = 1


- Solution: Data replication
 - At least 1 replica w/ the latest data version is accessed during each op
 - Each read op: $\text{read} \geq r$ replicas
 - Each write op: $\text{write} \geq w$ replicas
 - Quorum consensus protocol:** $r + w > n$ and $2 * w > n$
 - Tradeoffs: consistency, availability, performance, durability
 - ACID (**A**tomicity, **C**onsistency, **I**solation, **D**urability)

• Many RDBMS, e.g., MySQL, MS SQL, DB2, etc.




Eventual Consistency

- Intent: Performance and availability of data in case of network partitioning are enabled by ensuring **data consistency eventually and not at all times**



$$r + w \leq n$$



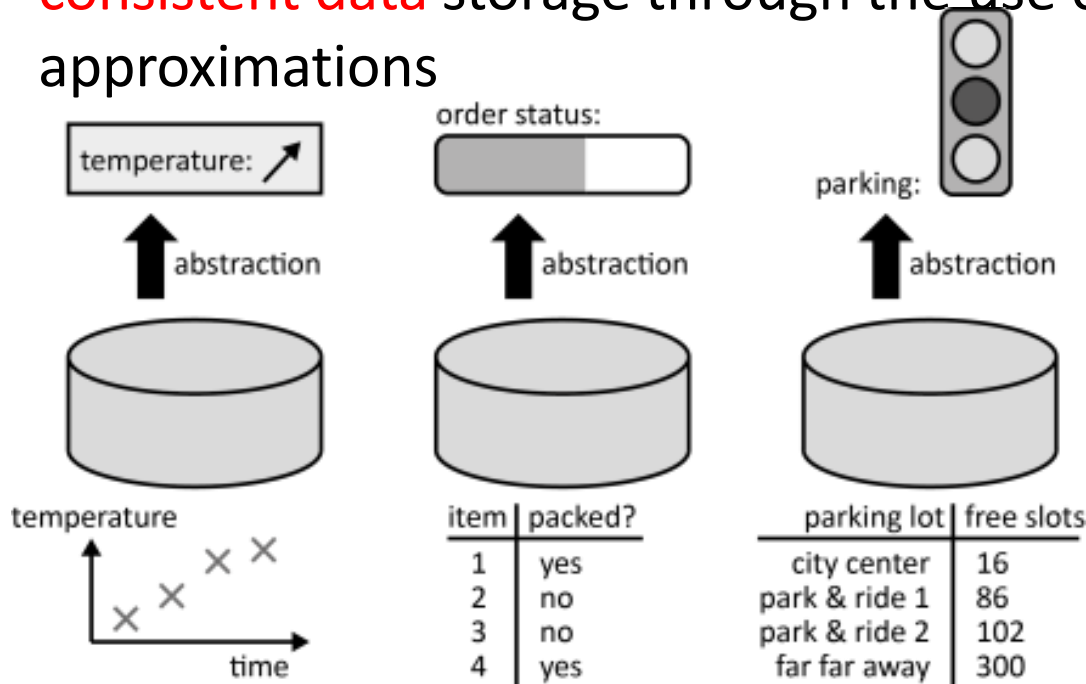
- Solution: trade consistency for availability (**CAP**)
 - reduces # of replicas to be accessed during read and write ops
 - modification eventually replicated to all replicas **asynchronously (msg Q)**
 - $r + w \leq n$
 - **BASE (Basically Available, Soft state, Eventual consistency)**

• NoSQL: AWS DynamoDB, Cassandra, Riak, etc.



Data Abtractor

- Intent: Data is **abstracted to inherently support eventually consistent data** storage through the use of abstractions and approximations

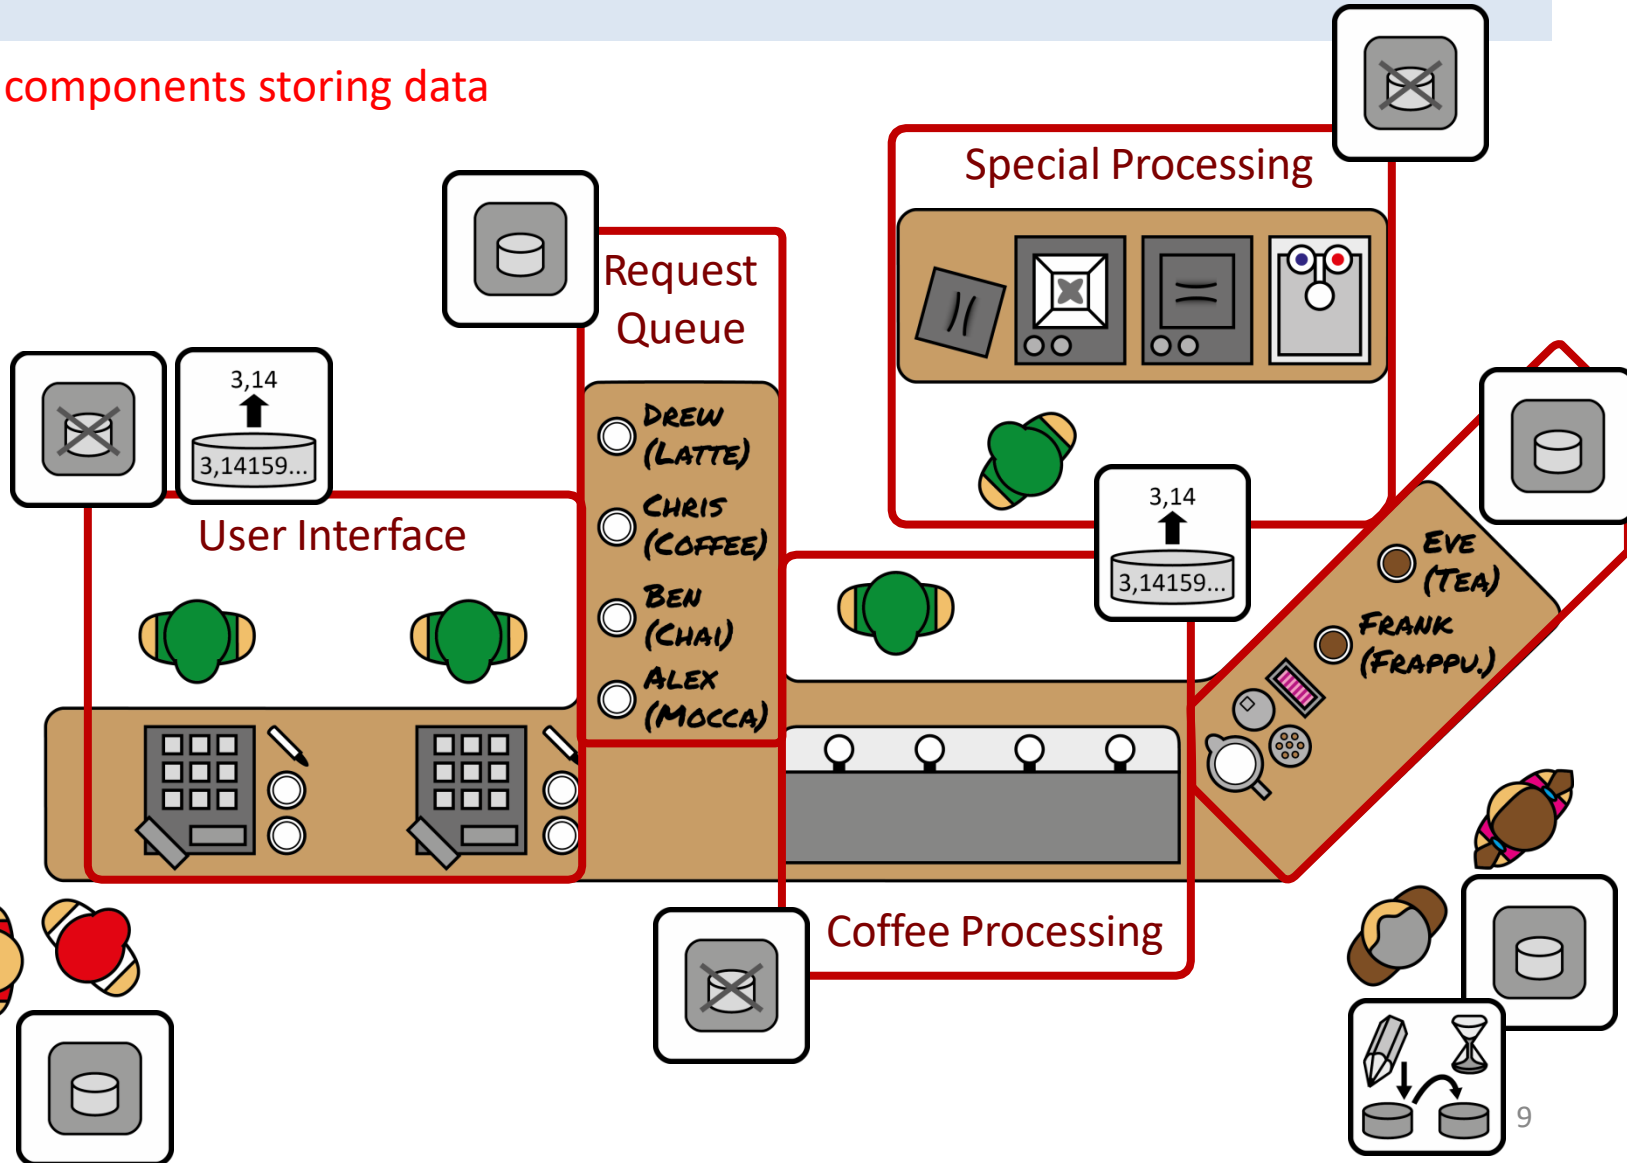


- SFPark.org
- Online stores: "in stock", "limited quantity", and "out of stock"

- Solution: data representation reflects that the consistent state is unknown by approximating values, or abstracting them into more general ones, such as progress bars, traffic lights, or change tendencies (increase / decrease)

Coffee Shop – Data (State)

Identify components storing data



Lessons – Data (State)

- State should **not** be handled by components whenever possible
- Handle session state in
 - requests (has to be provided with every request)
 - provider-supplied storage and communication offerings
- “Lie” about state whenever possible / acceptable


Coffee Shop App Design (4)

- Decomposition: How to distribute Application Functionality?
 - Distributed App
- Work load: What workload do components experience?
 - Static
 - Periodic
 - Once-in-a-lifetime
 - Unpredictable
 - Continuously Changing
- Data (State): Where does the application handle state?
 - Stateful
 - Stateless
 - Strict consistency
 - Eventual consistency
 - Data Abtractor
- **Component Refinements: How are components implemented?**
 - **Message-oriented Middleware**
 - **User Interface Component**
 - **Processing Component**
 - **Batch Processing Component**
 - **Multi-component Image**
- Elasticity and Resiliency
 - Elastic Load Balancer
 - Elastic Queue
 - Node-based Availability
 - Environment-based Availability
 - Watchdog



Message-oriented Middleware

- Intent: **Asynchronous** communication is provided while **hiding complexity of addressing, routing, or data formats** to make interaction robust and flexible



- AWS SQS
- Azure Service Bus
- IBM WebSphere MQ
- Apache ActiveMQ
- Google Cloud Pub/Sub
- Apache Kafka

End-to-end
msg delivery



Exactly-once
Delivery



At-least-once
Delivery



Transaction-based
Delivery

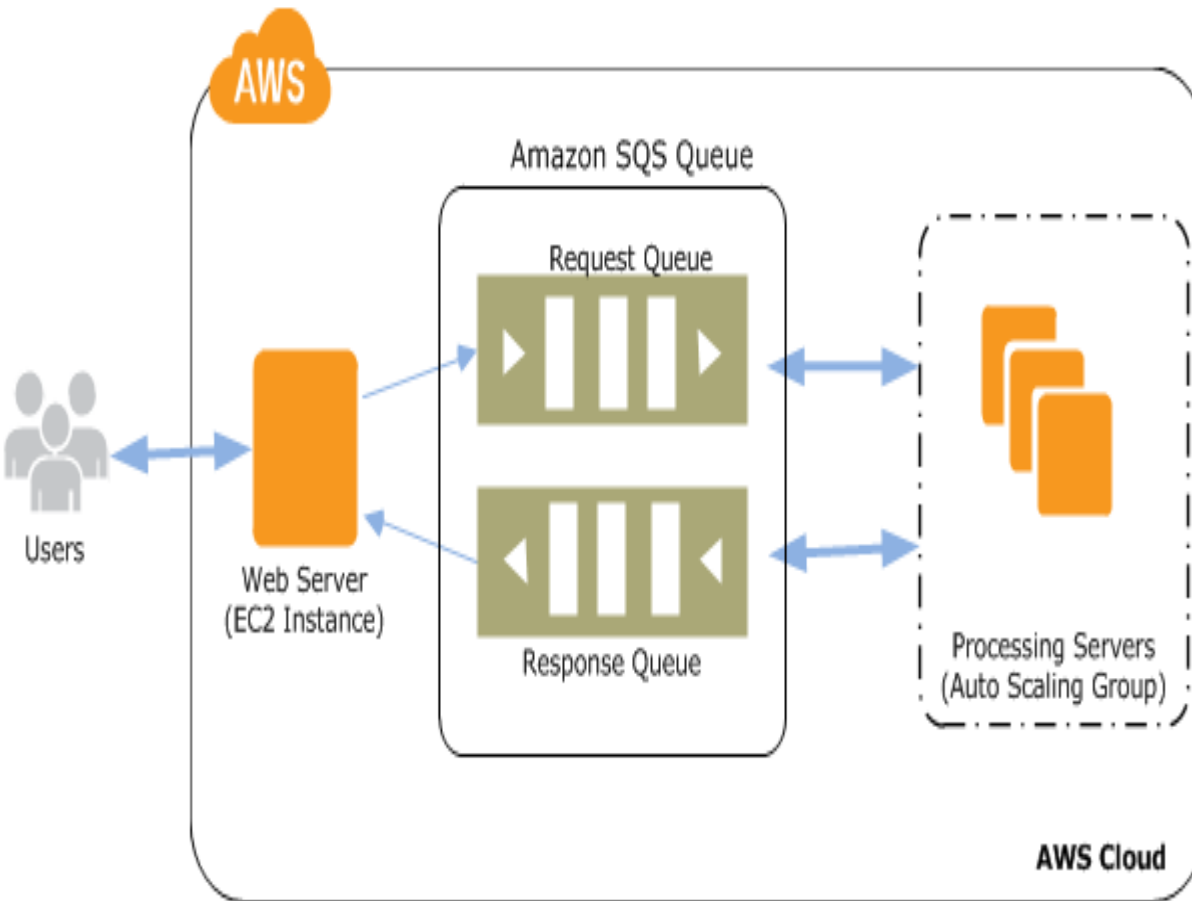


Timeout-based
Delivery

Interacting directly
w/ the receiver

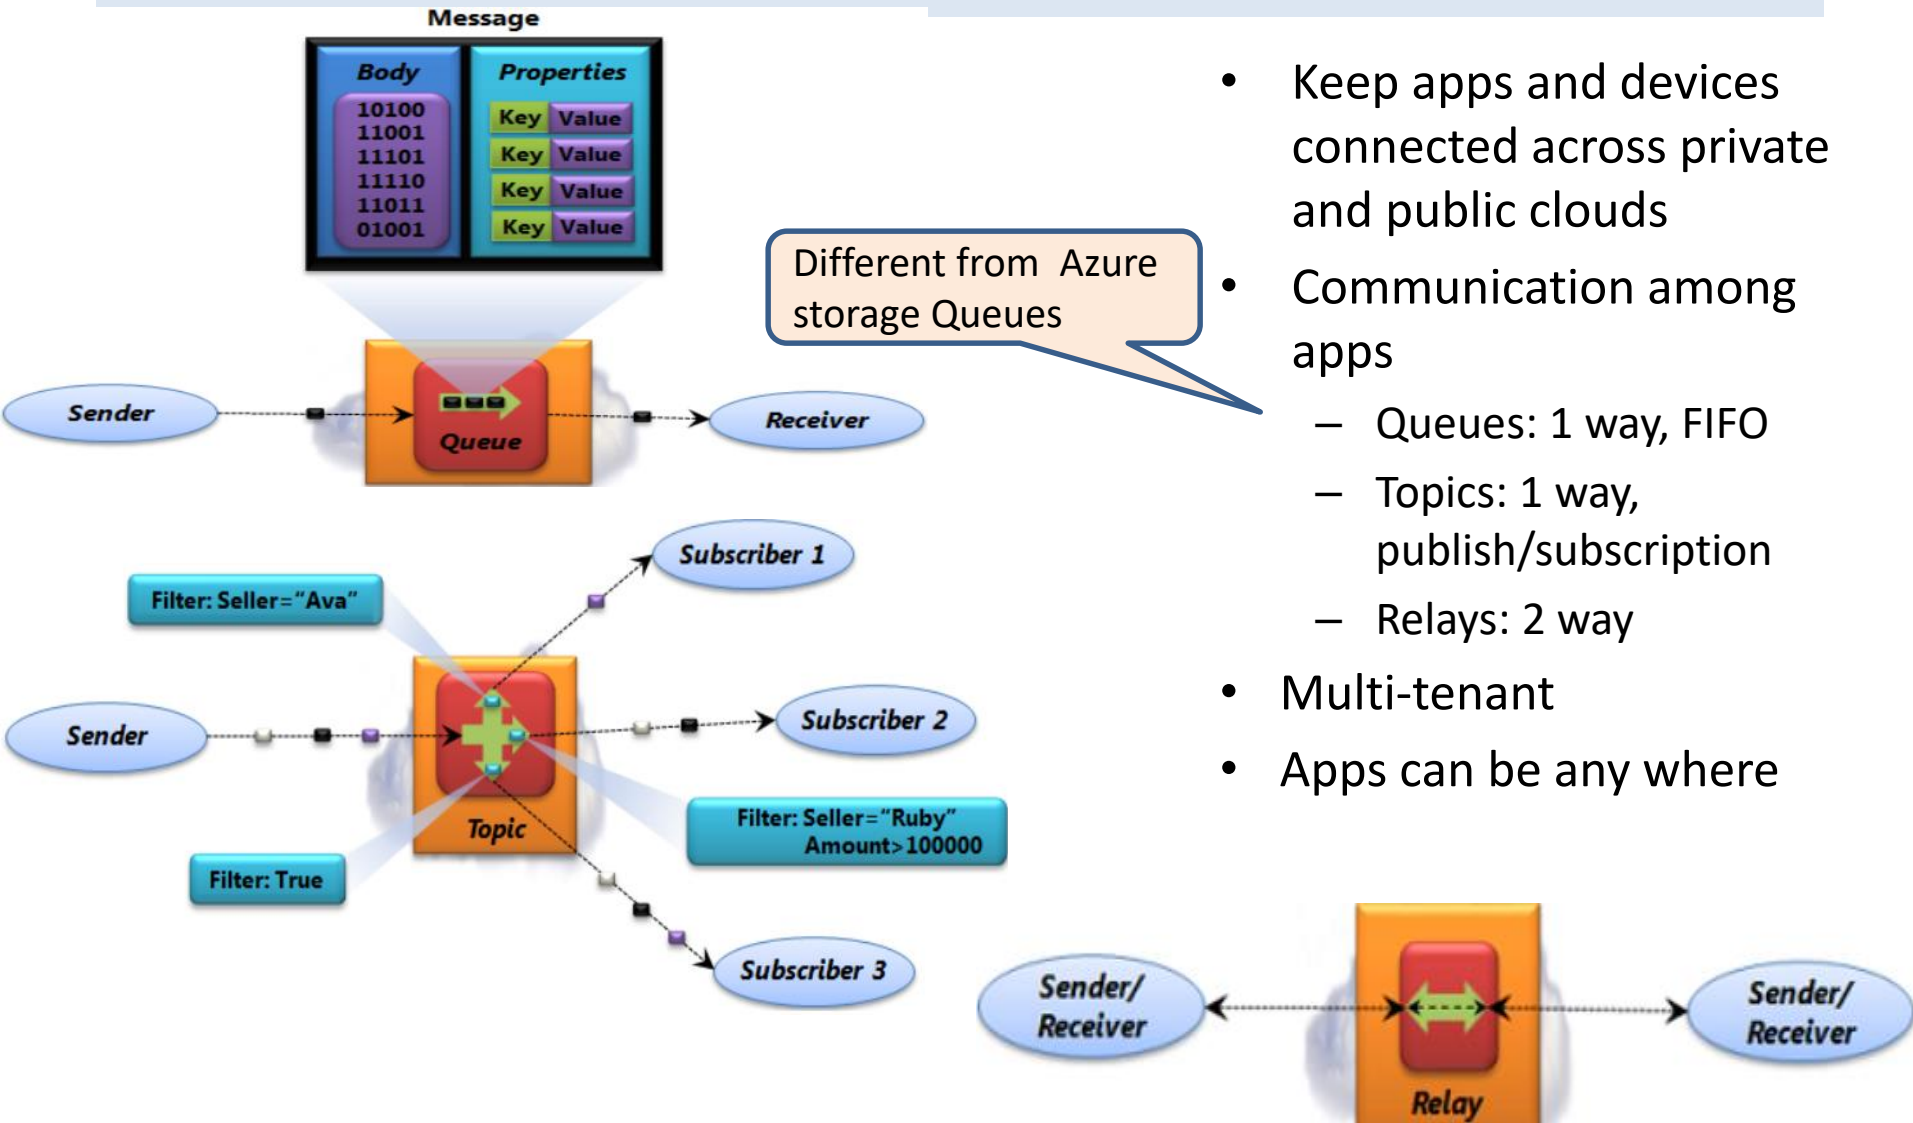
- Solution: message queue and pub-sub channel
 - Assumptions of communication partners reduced (**Loose Coupling**)
 - Platform**: implementation language used
 - Reference**: location of the communication partner (routing)
 - Time**: communication partners are active at different time / speed
 - Format**: message formats can change (transformation)

AWS - Simple Queue Service (SQS)



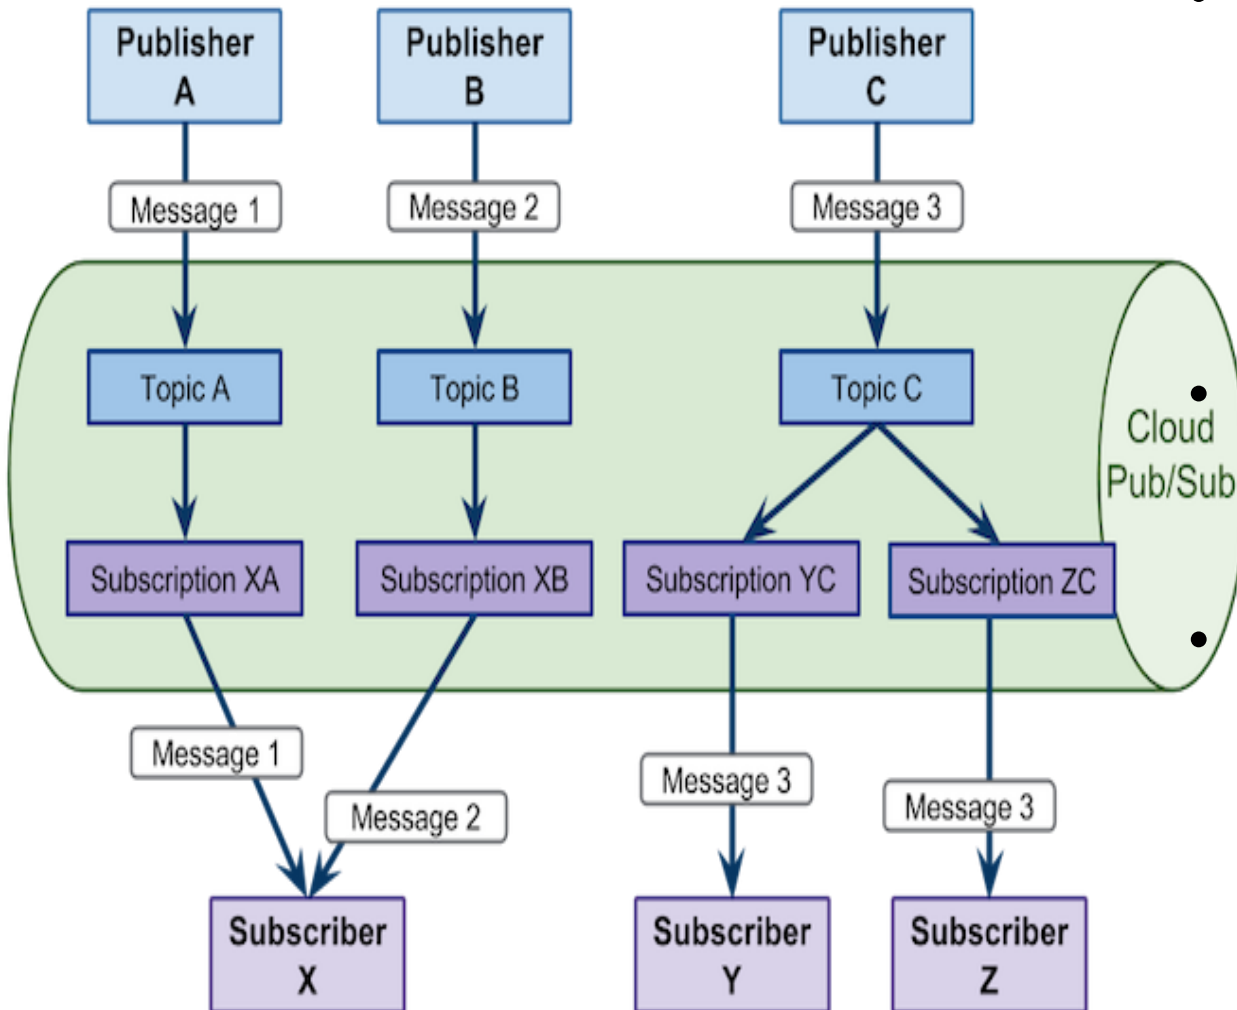
- Message queuing service - reliable msg delivery
- Msg **redundantly** distributed across SQS servers
- While a msg is received and being processed, it remains in Q and is not returned to subsequent receive reqs for the duration of the **visibility timeout**
- CloudWatch: dynamic scaling based on queue length

Azure - App Services - Service Bus



- Keep apps and devices connected across private and public clouds
- Communication among apps
 - Queues: 1 way, FIFO
 - Topics: 1 way, publish/subscription
 - Relays: 2 way
- Multi-tenant
- Apps can be anywhere

Google - Cloud Pub/Sub

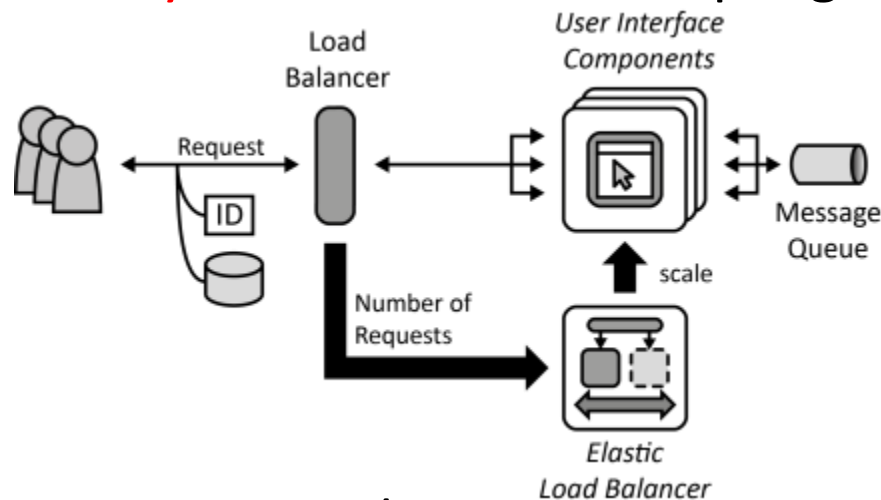


- Message queue
 - Reliable, HA
 - scalability
 - Many to many
 - Async
- Publisher
 - From any apps via https
 - Topic
- Subscriber
 - Any apps
 - modes
 - Pull: https
 - Push: POST over https



User Interface Component

- Intent: **Synchronous** user interfaces are accessed by **humans**, while **application-internal** interaction is realized **asynchronously** to ensure loose coupling



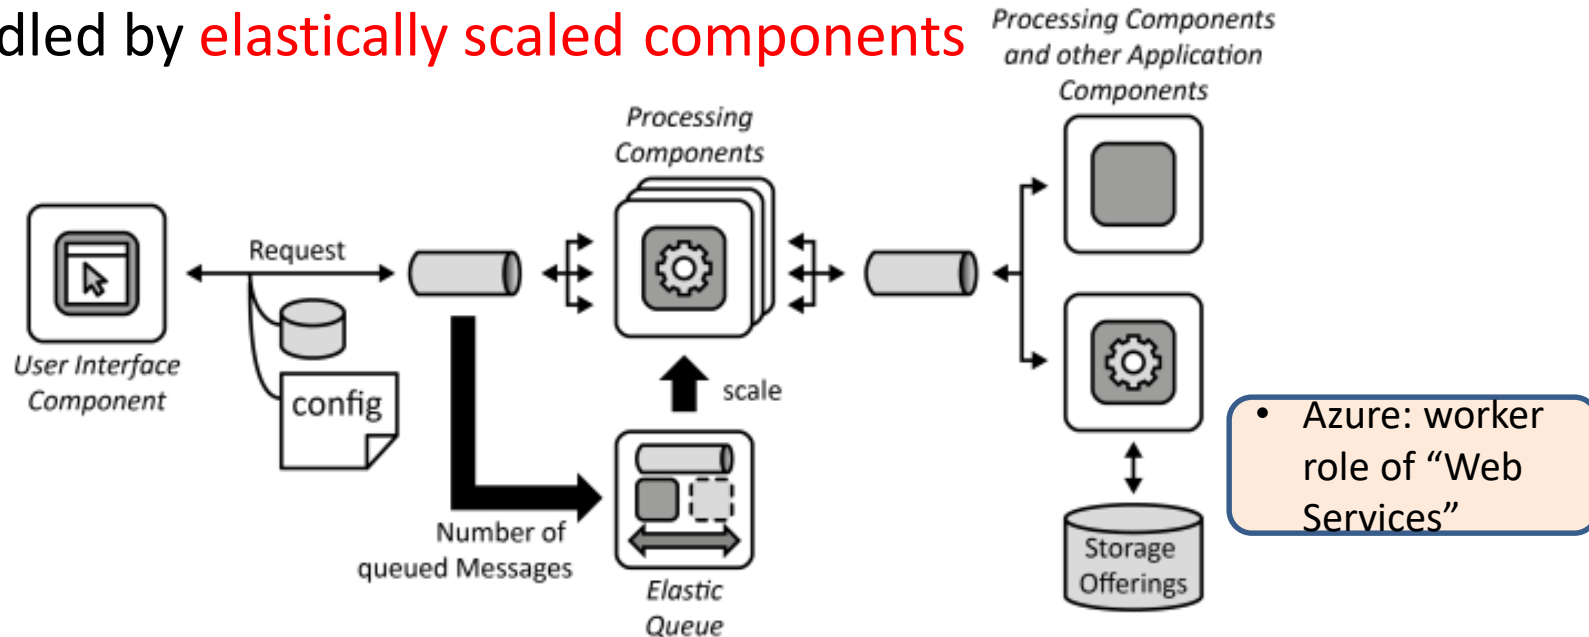
- Google Web Toolkit
- AJAX

- Solution: as a bridge b/w the sync access of the human user and the async communication used w/ other app components
 - State info held externally (Stateless Component pattern) - attached to requests, on user's device, or obtained from external storage
 - instances of UI Components scaled based on # of synchronous requests (Elastic Load Balancer pattern)




Processing Component

- Intent: Possibly long running processing functionality is handled by **elastically scaled components**



- Solution: Processing functionality split into separate function blocks and assigned to independent Processing Components
 - Each processing component: **scaled out independently** and **stateless**
 - Scaling is handled by an **Elastic Queue**
 - Data required for processing provided w/ reqs or by Storage Offerings¹⁷

Azure - Compute - Cloud Services



- PaaS for apps
 - VM roles: web role, worker role
 - Monitor: failure detection (apps, VM, host) & restart
 - **Auto scale, load balancing**
 - HA & FT: VM pool

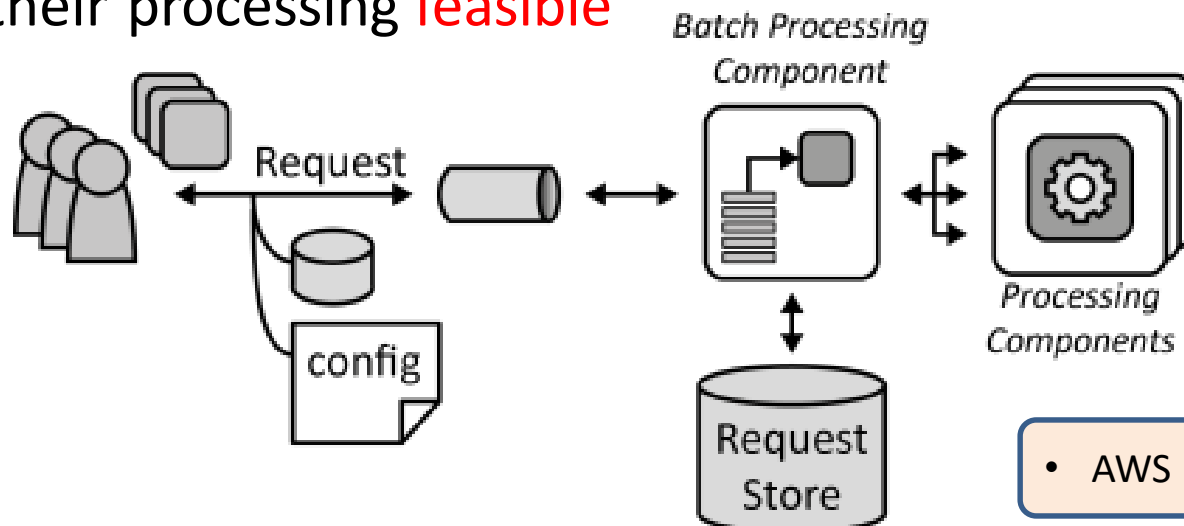
Cloud service apps

- Do not create VM directly (PaaS)
- User deploy apps and specify # of web/worker roles
- apps state: only in SQL DB, blobs, tables, etc. **Not in file systems of VM (why?)**



Batch Processing Component

- Intent: **Requests** are **delayed until** environmental conditions make their processing **feasible**

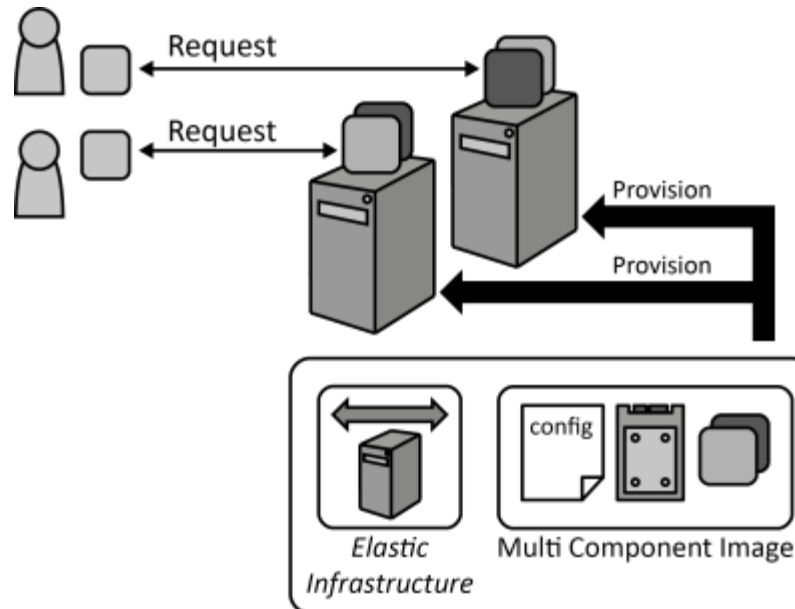


- Solution: Async processing reqs accepted at all times, but stored them until conditions are optimal for their processing
 - Based on # of stored requests, environmental conditions, and custom rules, components are instantiated to handle the requests
 - Requests are only processed under non-optimal conditions if they cannot be delayed any longer



Multi-Component Image

- Intent: **Virtual servers host multiple application components** that may not be active at all times to reduce provisioning and decommissioning operations



vs single image,
tradeoffs?

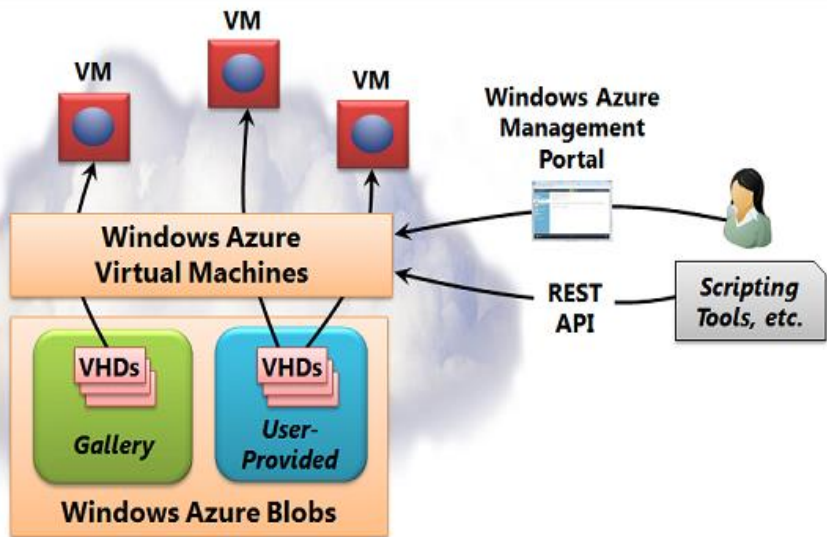
- Solution: Multiple app components (possibly including middleware) are hosted on virtual servers to ensure that running virtual servers may be used for different purposes w/o making provisioning or decommissioning operations necessary
 - Challenges: optimize distribution of app components among images

- VM templates: VMware, AWS EC2, Azure, GCE
- Containers: Docker, AWS ECS, Google Container


AWS - Elastic Computing (EC2)

- IaaS: HVM (Xen)-based VM
- Various instance types
- OS : Linux, Windows, etc.
- Create EC2 instance
 - Select or create an Amazon Machine Image (AMI)
 - Config security, network
 - Choose instance type
 - Decide locations, static IP, etc
- Access EC2 instance
 - DNS name
 - private IP: internal
 - public IP: external, dynamic
 - NAT: internal ↔ external
- Elastic IP addr: static, \$
 - not released when instance is stopped or terminated
 - must be released when not needed
- EBS snapshot: FS volume only; no VM memory; VM up or down
- No live migration, no hot clone
- EC2 SLA (monthly uptime)
 - < 99.95%, > 99%: 10% service credit
 - < 99%: 30% service credit
- UI: AWS Management Console
- AWS SDK libraries and toolkits
 - Java, PHP, C#, etc
 - SOAP (WSDL), REST

Azure - Compute - Virtual Machine





- IaaS
- OS types: Windows, Linux
- Image: portal image, custom
- File: .VHD
- Connect to VM: VPN, ExpressRoute
- **Load balancing** multiple VMs
 - DNS level, network level
- Scalability: **Scale Sets** - autoscale
- HA: **Availability sets** - redundancy
- Grouping VMs as a service (multi-tier apps): Single public IP
 - Load balancing + Scalability + HA



Google Compute Engine (GCE)

- IaaS: KVM-based VM
- OS: various Linux, Windows
- Machine types
 - Standard, small
 - High CPU, high memory
- Instance template
- Load balancing: Replica pool
 - Network, HTTP (cross-region)
- Instance group: pool of VMs
 - Size up/down
- Autoscaler: managed instance group
 - Network, HTTP (cross-region)
 - CPU util%, Cloud metrics
- Authorizing to/from other Google service: OAuth 2.0
- Resource: global / region / zone
 - Zone: isolated location in a region
 - Region-specific resources can be accessed by resources in the region
 - Isolation of failure, redundancy
- Disk snapshot: FS only; no VM memory; VM up or down
- no VM hot clone
- Host maintenance policies
 - [default] VM live migration
 - Terminate and (optionally) restart
- VM crash: [default] auto restart


Docker



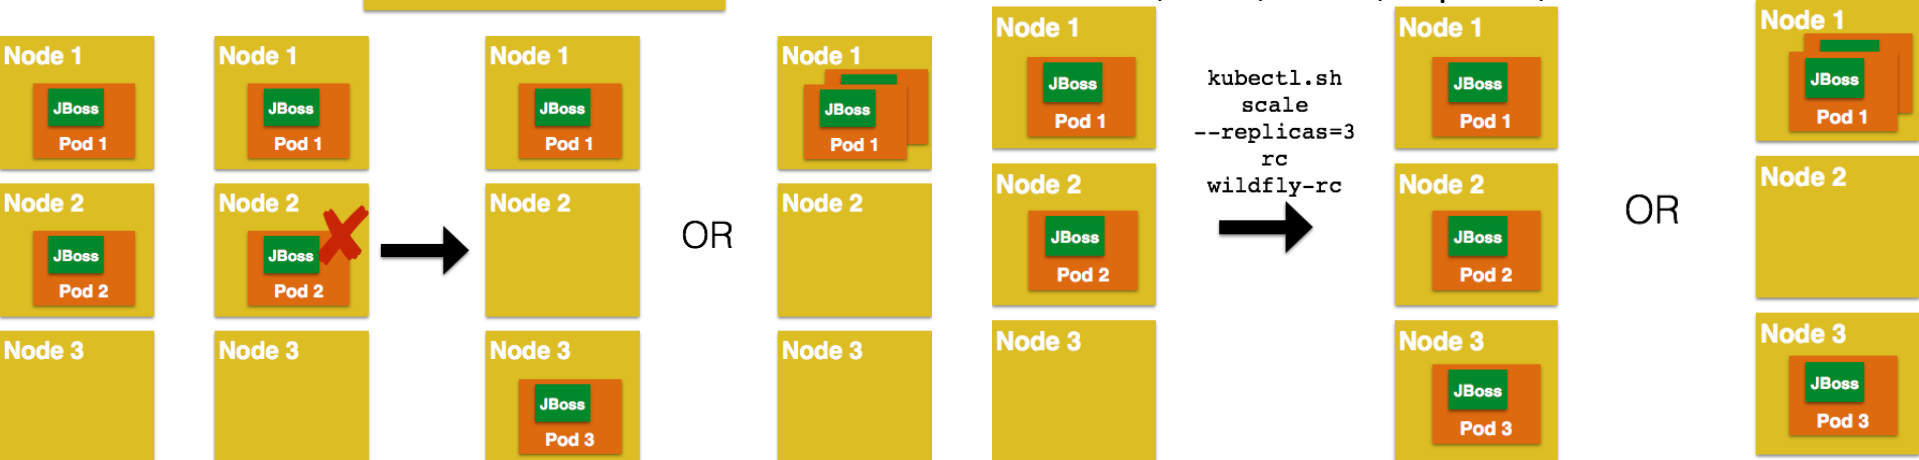
- Docker Engine
 - Based on Linux container LXC
 - Images (left)
 - **Union File system**: Multiple layers of FS + writable layer
 - Share/deploy faster, scale easily
 - Docker Hub, Docker Registry (left)
- Built-in orchestration – Swarm
 - cluster, scheduling, placement
 - Integrated or standalone: since v1.12
- Docker Machine: auto provision hosts across multi-platforms
- Docker Compose: define multi-container apps




Google Container Engine (GKE)

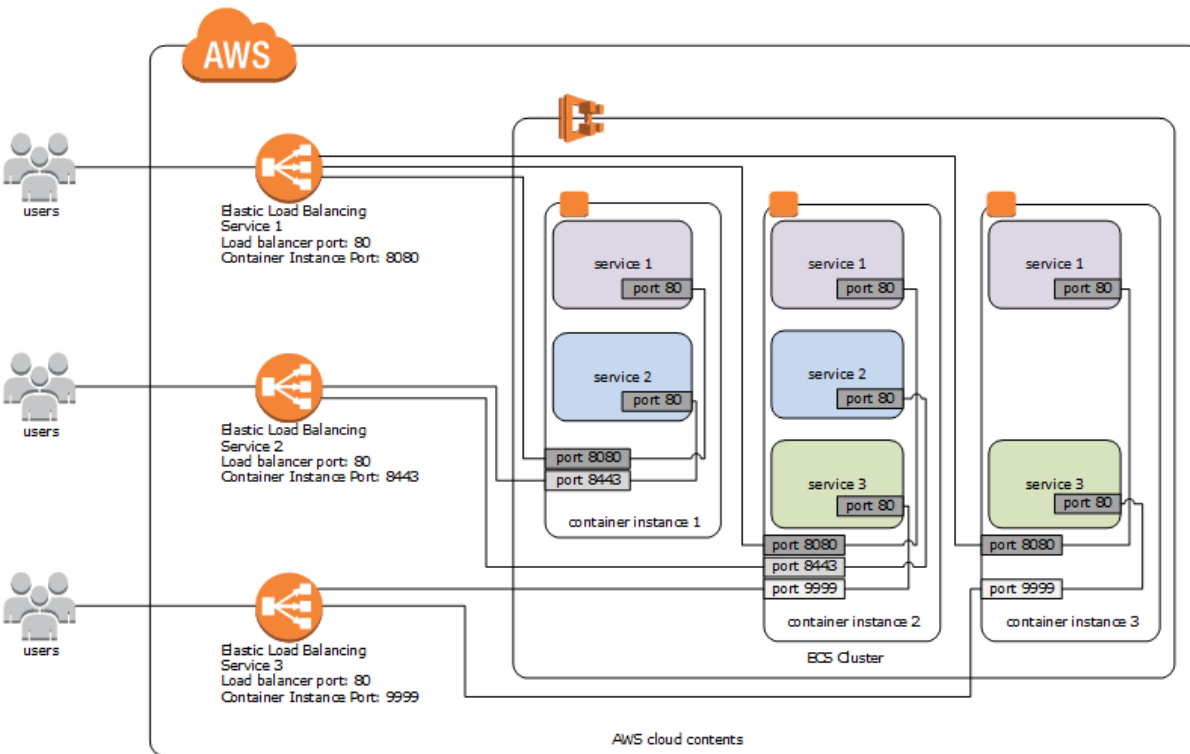


- Docker support
- Run on a cluster of GCE (VMs)
- Kubernetes: managing containerized apps among multi-hosts
 - Logically organize containers into groups
 - Replication: optimal # of pods
 - Rescheduling pods (lower left)
 - Scaling pods (lower right)
 - Cluster Auto-restarting
 - Docker & rkt support
 - GCE, AWS, Azure, vSphere, etc.





AWS EC2 Container Service

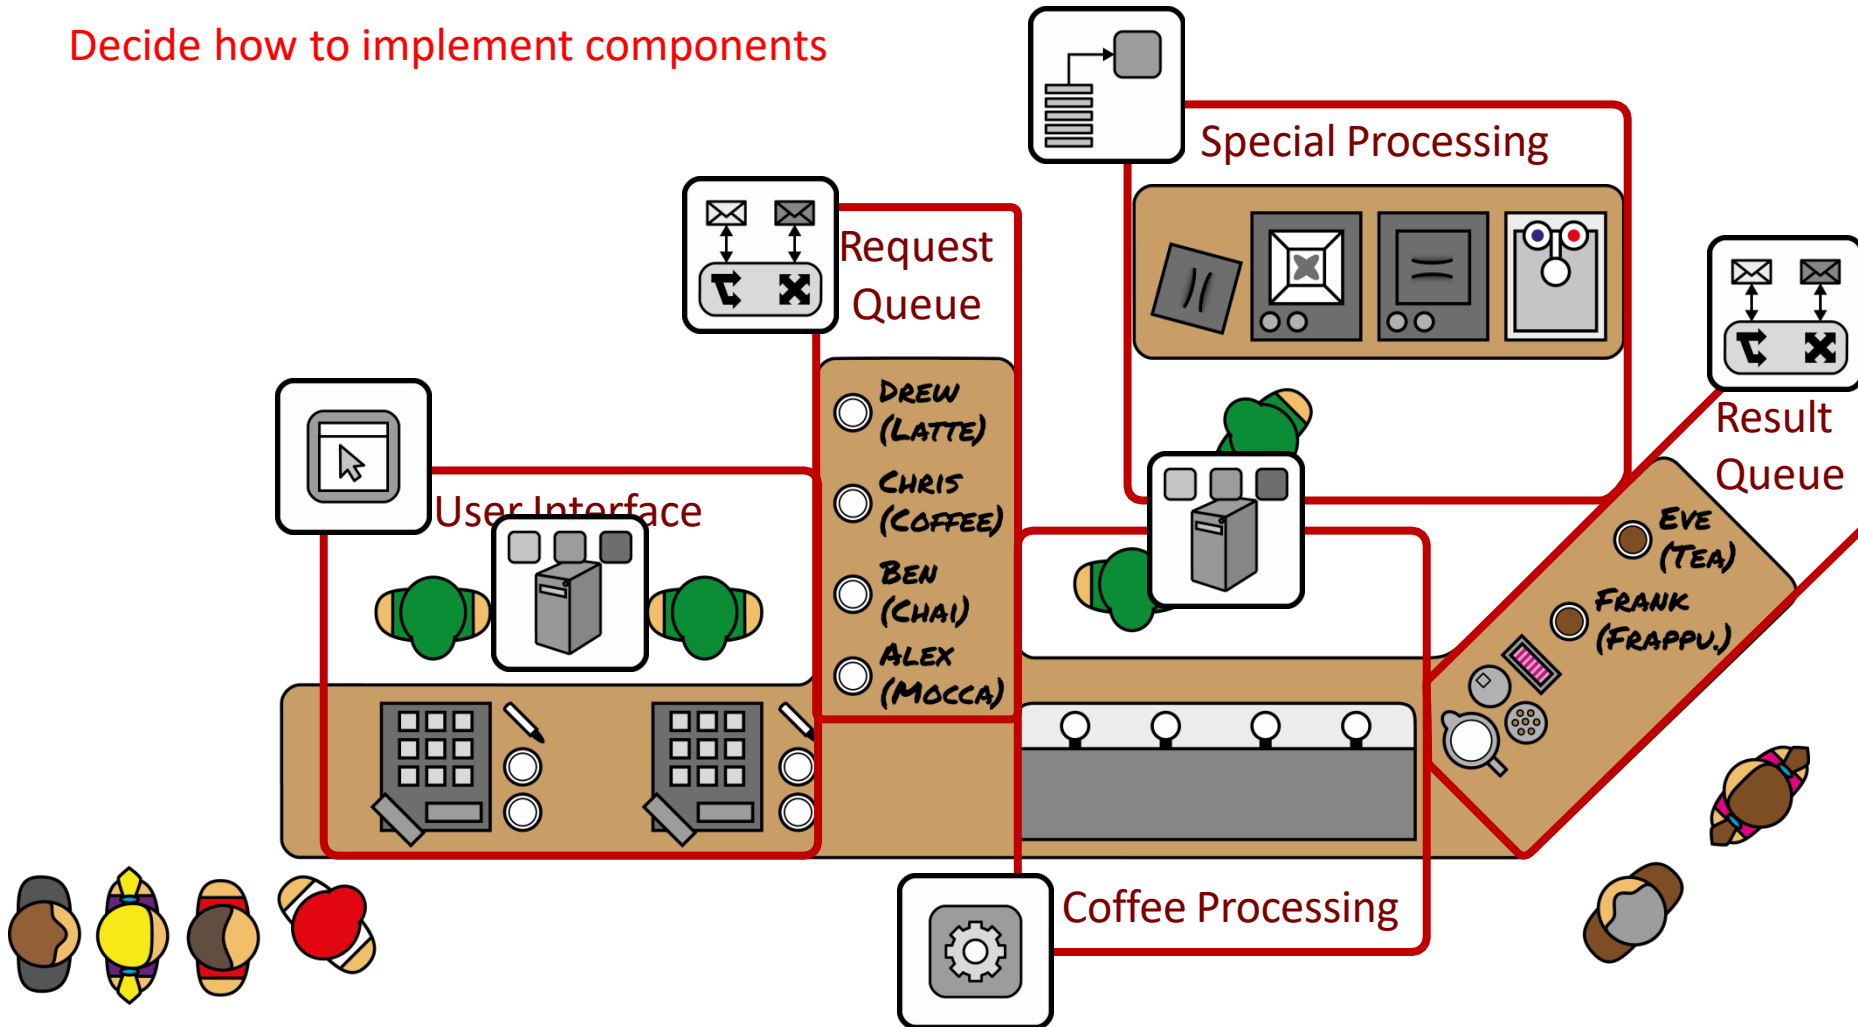


- Docker support
- Across a cluster of EC2 instances
 - App span multiple AZs
 - Placement mgmt: auto or manual control
 - Cluster mgmt and config mgmt
 - Scheduling: balance between resource needs and availability requirements
 - Integrated w/ ELB (left)

• Docker workload can be migrated to/from AWS

Coffee Shop – Refinement of Components

Decide how to implement components



Lessons – Component Refinement

- Components should be integrated with messaging to ensure loose coupling
- Resources supporting functional components should be flexible (multi-component image)

References

- <http://www.cloudcomputingpatterns.org/>
- Christoph Fehling, Frank Leymann, Ralph Retter, Walter Schupeck, and Peter Arbitter, *Cloud Computing Patterns: Fundamentals to Design, Build, and Manage Cloud Applications*. Springer; 2014
 - <http://www.springer.com/978-3-7091-1567-1>
- The coffee shop example is adapted from
 - <https://indico.scc.kit.edu/indico/event/26/session/1/contribution/12/material/slides/0.pdf>
 - <http://www.sei.cmu.edu/library/assets/presentations/retter-saturn2013.pdf>