


CMPE 282 Cloud Services

***Cloud-Native Application***

***Design Patterns***


Instructor: Kong Li

# Content

- What and why
  - Coffee Shop
    - Decomposition
    - Workload
    - Data/state
    - Component refinement
    - Elasticity and Resiliency
  - Composite Cloud Applications
- 
- The diagram consists of three blue brackets on the right side of the list, each grouping a set of items and labeled with a number. The first bracket groups 'What and why' and 'Coffee Shop' and is labeled '1'. The second bracket groups the sub-items of 'Coffee Shop' and is labeled '2'. The third bracket groups 'Composite Cloud Applications' and is labeled '3'.

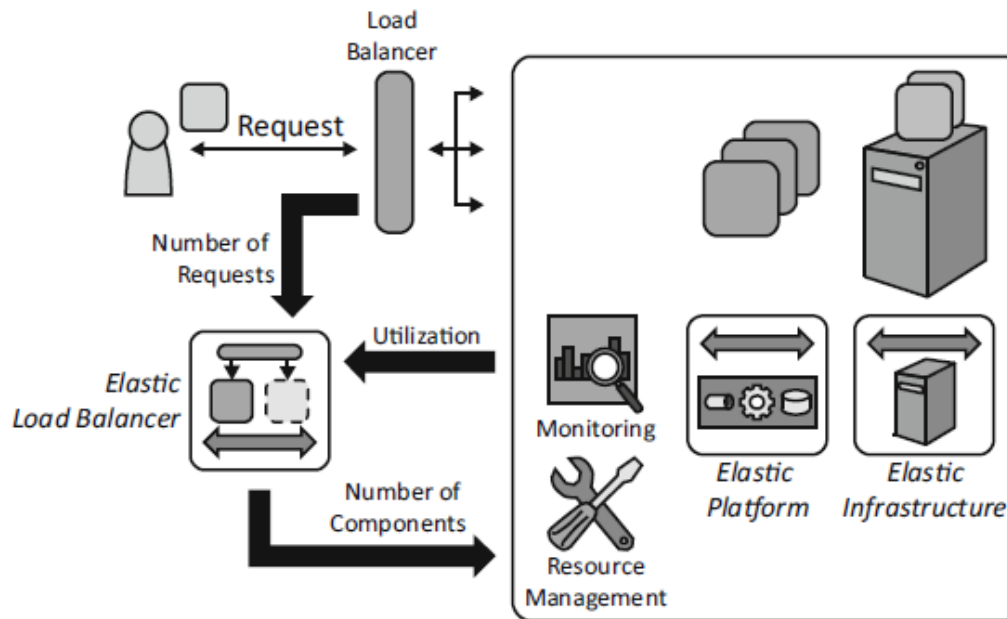
# Coffee Shop App Design (5)

- Decomposition: How to distribute Application Functionality?
  - Distributed App
- Work load: What workload do components experience?
  - Static
  - Periodic
  - Once-in-a-lifetime
  - Unpredictable
  - Continuously Changing
- Data (State): Where does the application handle state?
  - Stateful
  - Stateless
  - Strict consistency
  - Eventual consistency
  - Data Abtractor
- Component Refinements: How are components implemented?
  - Message-oriented Middleware
  - User Interface Component
  - Processing Component
  - Batch Processing Component
  - Multi-component Image
- **Elasticity and Resiliency**
  - **Elastic Load Balancer**
  - **Elastic Queue**
  - **Node-based Availability**
  - **Environment-based Availability**
  - **Watchdog**



# Elastic Load Balancer

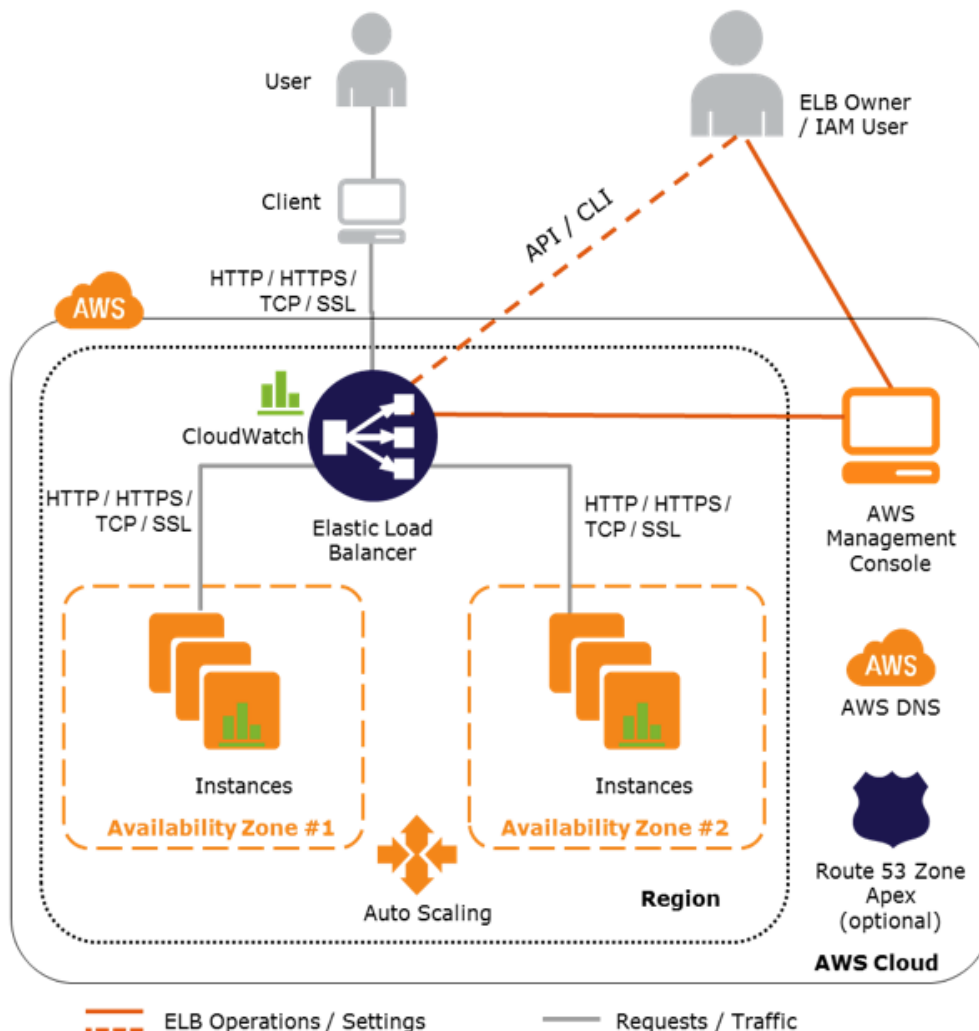
- Intent: # of **sync accesses** to an elastically scaled-out app and its util% is used to determine # of required app component **instances**



- AWS ELB, auto scaling
- Google GCE auto scaler, GAE auto scaling
- Azure Traffic Manager, Cloud Services auto scale

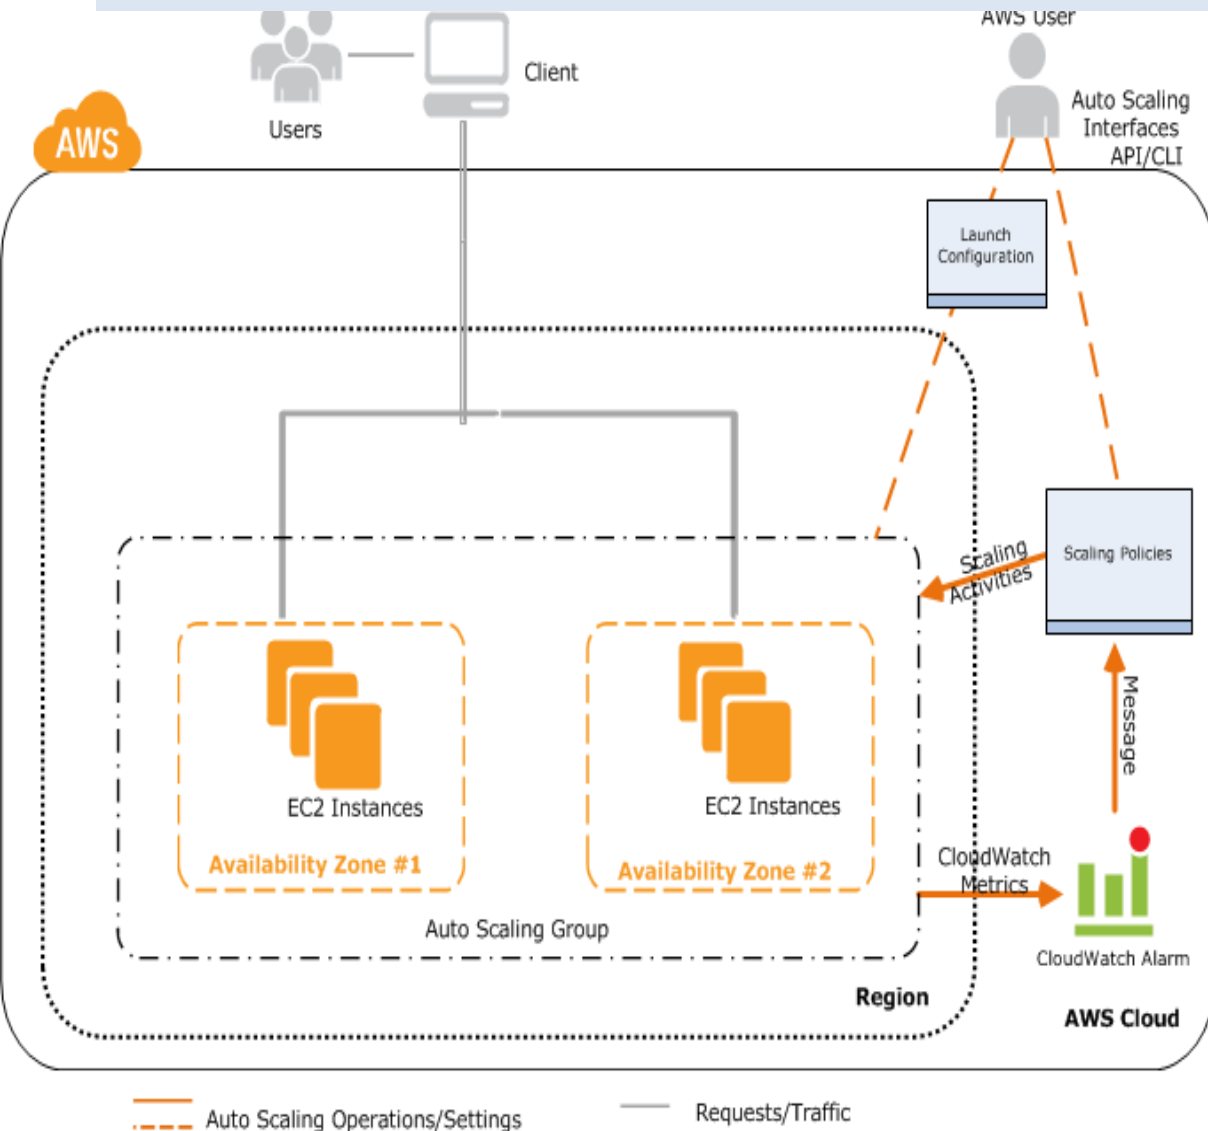
- Solution: Based on # of sync user requests and/or util%, the elastic load balancer determines # of component instances, and performs provisioning and decommissioning on elastic platform or elastic infrastructure
  - Horizontal scale: scale out / scale in
  - Network load balance and workload load balance

# AWS - Elastic Load Balancing (ELB)



- Distribute incoming apps traffic across AZs (each w/ multiple EC2 instances)
  - *Classic LB: Network load balance*
    - to healthy EC2 only
  - *Application LB: content-based routing*
    - across > 1 services or containers on EC2 instances
  - Integration w/ *auto scaling (workload LB)*
- vs. VMware DRS

# AWS - Auto Scaling

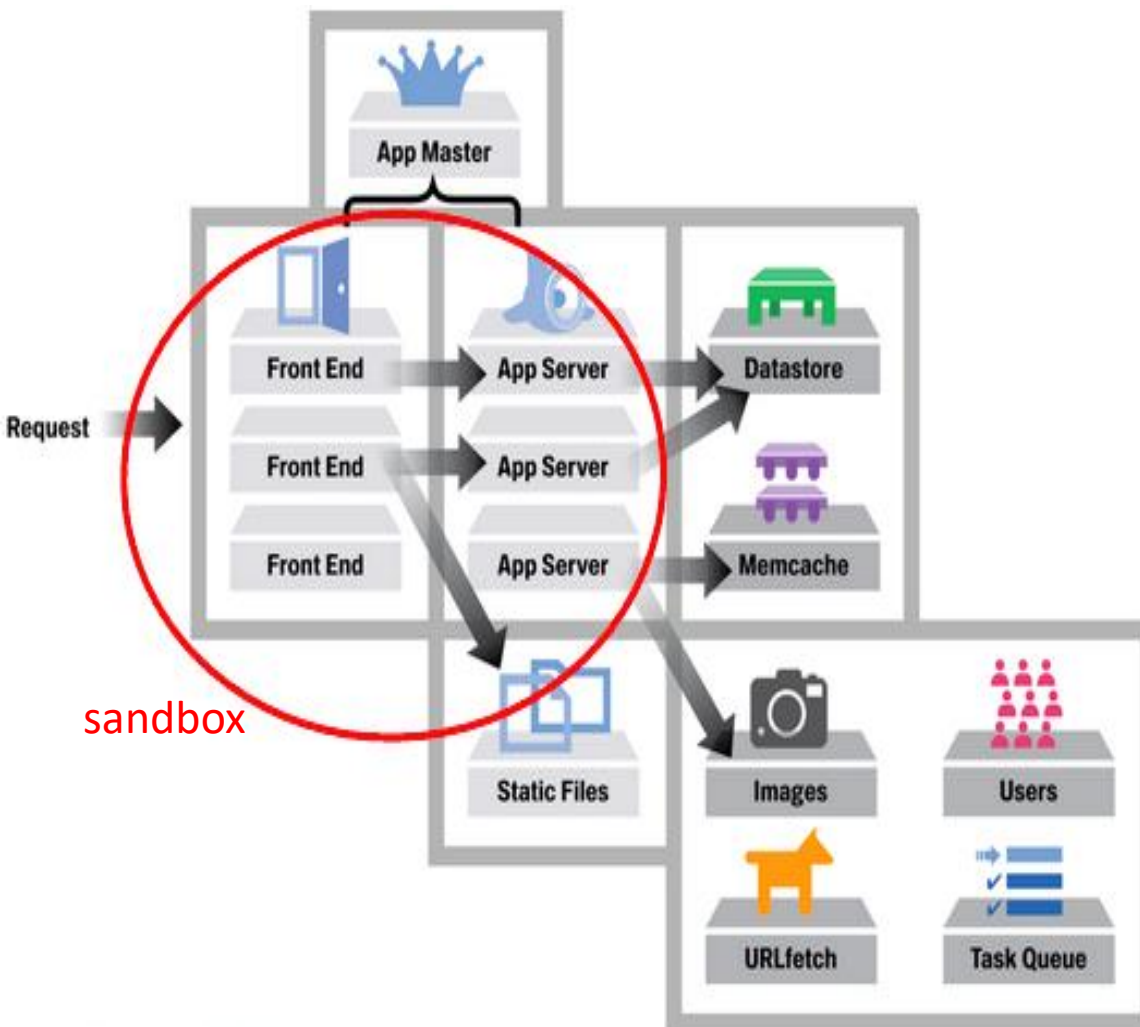


- Auto scale out/in
  - *Workload (resource usage)* *load balance*
  - Keep “optimal” # of VMs at optimal util% across AZs
  - Auto scaling group
  - Policies
    - Launch config
    - Min/max group size
    - AZs
    - conditions → alarms → actions
- Enabled by CloudWatch
- In cooperation with ELB
- ≡ vRealize Orchestrator
- vs. VMware DRS

# Google Compute Engine (GCE)

- IaaS: KVM-based VM
- OS: various Linux, Windows
- Machine types
  - Standard, small
  - High CPU, high memory
- Instance template
- Load balancing: Replica pool
  - Network, HTTP (cross-region)
- Instance group: pool of VMs
  - Size up/down
- **Autoscaler**: managed instance group
  - Network, HTTP (cross-region)
  - CPU util%, Cloud metrics
- Authorizing to/from other Google service: OAuth 2.0
- Resource: global / region / zone
  - Zone: isolated location in a region
  - Region-specific resources can be accessed by resources in the region
  - Isolation of failure, redundancy
- **Disk snapshot: FS only; no VM memory; VM up or down**
- **no VM hot clone**
- Host maintenance policies
  - [default] **VM live migration**
  - Terminate and (optionally) restart
- VM crash: [default] **auto restart**

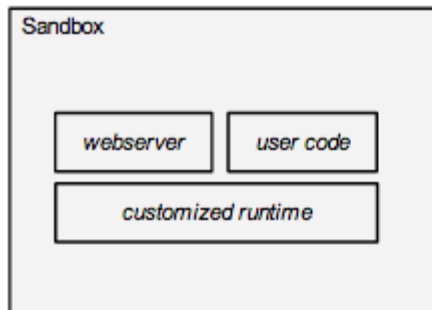

# Google App Engine (GAE) - Standard



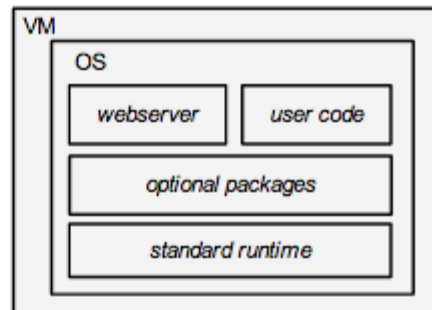
- PaaS: web app
- Sandboxed (container-based) env: Python, **Java 7**, PHP, GO
- Instance **scaling**: auto (req rate, latency), basic, manual
- Storage
  - Datastore: NoSQL
  - Cloud SQL: MySQL
  - Blobstore
  - Cloud Storage
- Communication
  - Task queue: async process
  - Endpoints: client lib/REST
- Scheduled tasks for triggering events



# Google App Engine (GAE) - Flexible



App Engine Hosting Environment

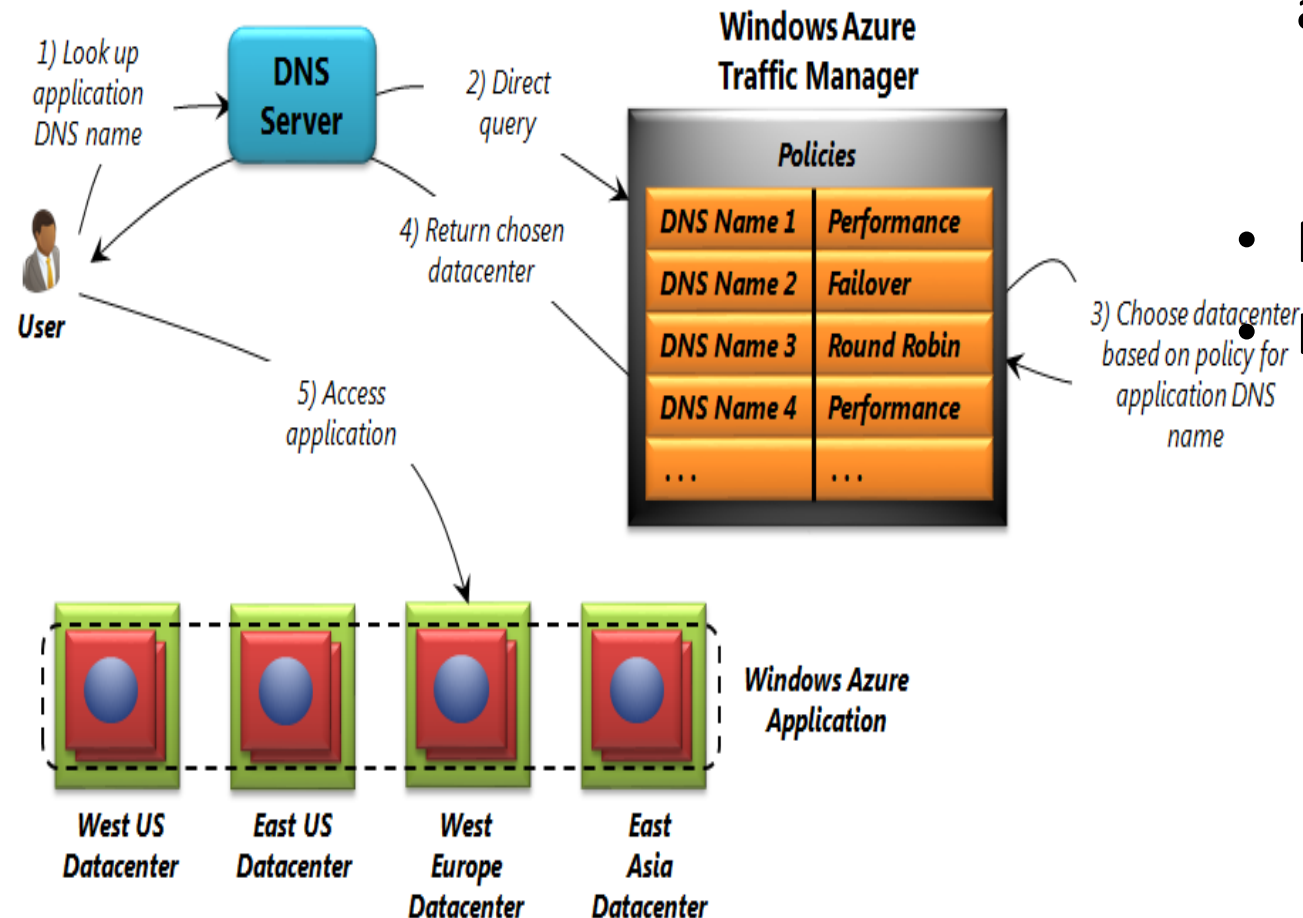


Managed VM Hosting Environment

- Initial name: Managed VMs
- run App Engine apps on configurable GCE (VMs)
  - More cpu / mem options
  - User configurable custom runtimes
  - Java 8, Python, Ruby, GO, Node.js, (Docker-based) custom runtime
  - Apps: mix-and-match standard and flexible env
- Auto-scaling
- Load balancing

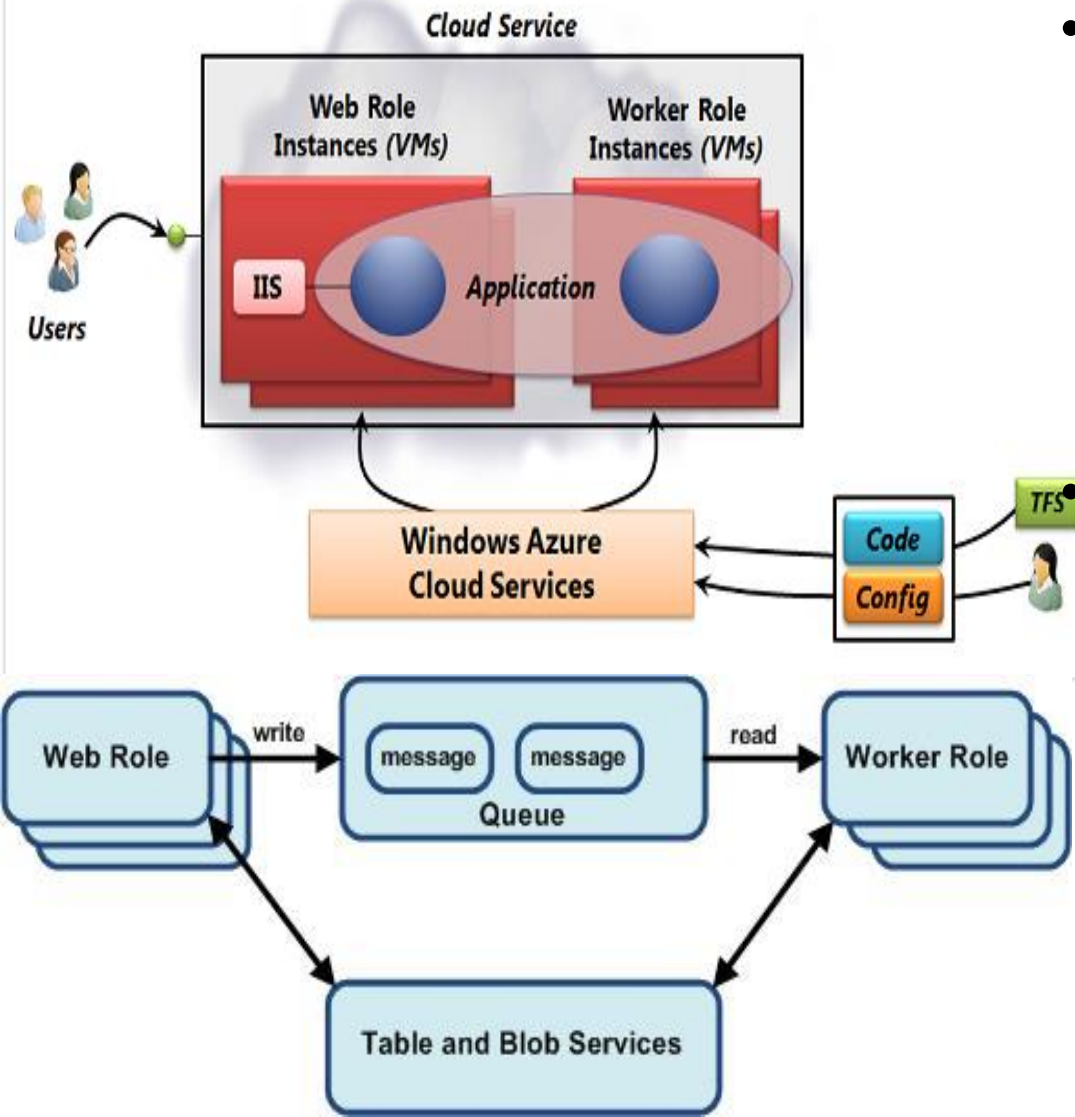
# Azure - Network Services - Traffic Manager

## Traffic Manager



- **Load balance** incoming traffic for performance and HA
  - Distribute traffic over multiple locations
- Monitor health status
- Load balancing options
  - Failover
  - Round robin
  - Performance


# Compute - Cloud Services



- PaaS for apps
  - VM roles: web role, worker role
  - Monitor: failure detection (apps, VM, host) & restart
  - **Auto scale, load balancing**
  - HA & FT: VM pool

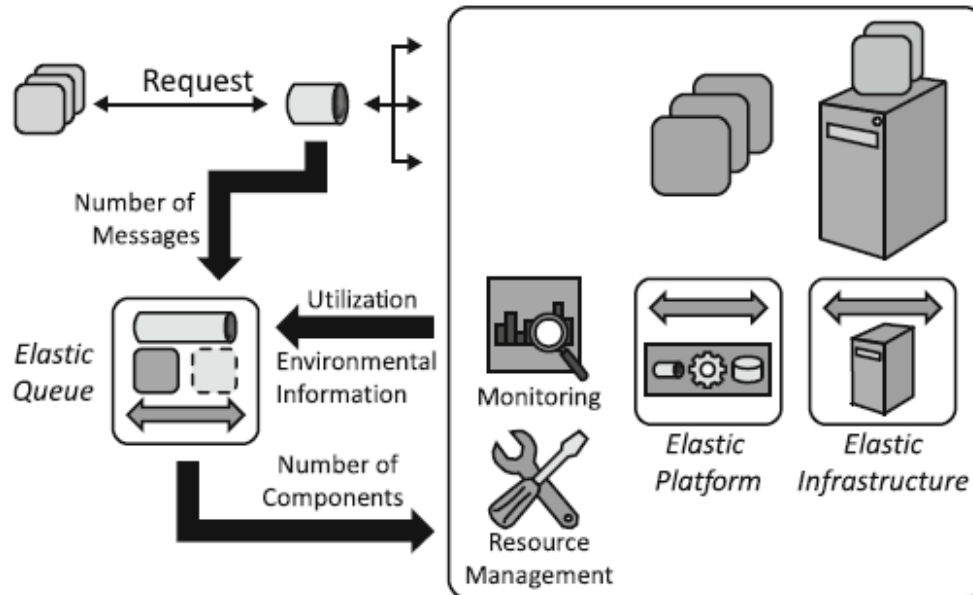
## Cloud service apps

- Do not create VM directly (PaaS)
- User deploy apps and specify # of web/worker roles
- apps state: only in SQL DB, blobs, tables, etc. **Not in file systems of VM (why?)**



# Elastic Queue

- Intent: # of **async accesses via messaging** to an elastically scaled-out app and its util% is used to **adjust # of** required app component **instances**



- AWS CloudWatch

- Solution: the elastic queue monitors the queue (that distributes msg to component instances) and/or util% and determines # of component instances, and performs provisioning and decommissioning on elastic platform or elastic infrastructure
  - Horizontal scale: scale out / scale in
  - Network load balance and workload load balance

# CloudWatch

Name	Volume ID	Capacity	Volume Type	Snapshot	Created	Zone	State	Alarm	Attachment Information
empty	vol-b12f11da	360 GiB	standard	snap-c763b7a6	2011-07-16T13:51:10Z	us-east-1b	in-use	none	i-439b5c25 (adm-1000):sdi4

1 Volume selected

Volume: vol-b12f11da

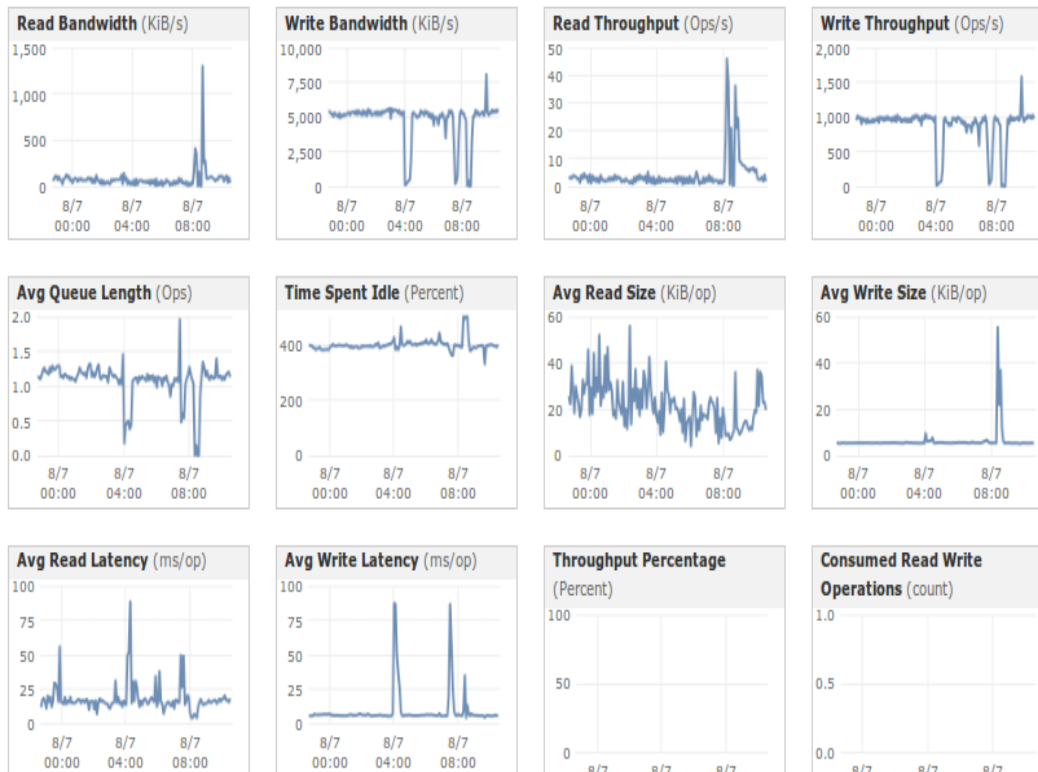
Details Status Checks **Monitoring** Tags

CloudWatch alarms: No alarms configured for vol-b12f11da [View all CloudWatch alarms](#)

Create Alarm

CloudWatch metrics: Times are displayed in UTC.

Time Range: Last 12 Hours




- Mgmt: monitor AWS cloud resources and apps

- Standard metrics
- Custom metrics (from customer's apps and services)

- Policy-based **Auto Scaling**

- Auto scaling group
- Add/remove (EC2) instances **dynamically** based on metrics

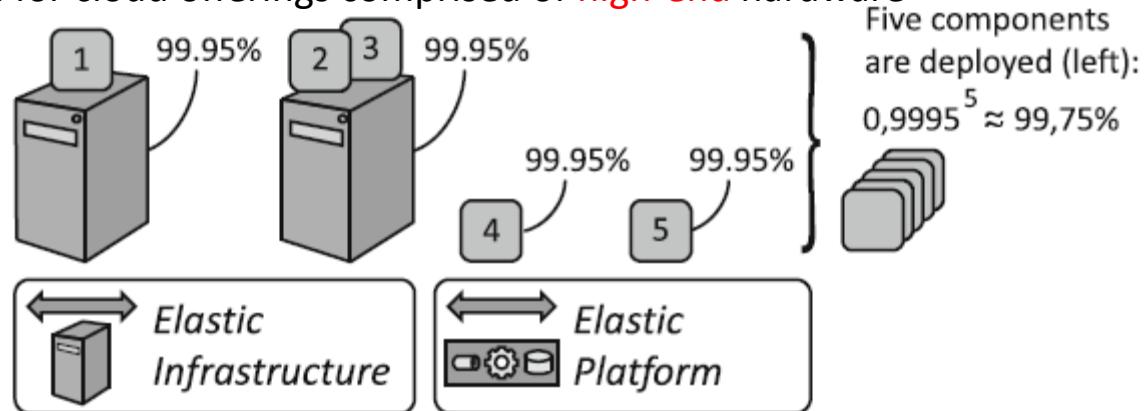
- util%
- queue length
- etc.




# Node-based Availability

- Intent: A cloud provider guarantees the **availability of individual nodes, such as individual virtual servers, middleware components, or hosted app components**

- often used for cloud offerings comprised of **high-end** hardware




- Solution: The provider assures availability for *each individual* app component (elastic platform), or the provided virtual server (elastic infrastructure)
  - How to monitor? Util%, test data/msg, heartbeat, etc.
  - customer estimates overall avail% = (node<sub>1</sub> avail%) \* (node<sub>2</sub> avail%) \* ...
  - If assured avail% is low, customer may incorporate redundancy and failure replacement on the app level (deploying multiple app component instances) <sup>14</sup>



# Environment-based Availability

- Intent: A cloud provider guarantees the (overall) availability of the env hosting individual nodes, such as virtual servers or hosted app components
  - often used for cloud offerings comprised of commodity hardware



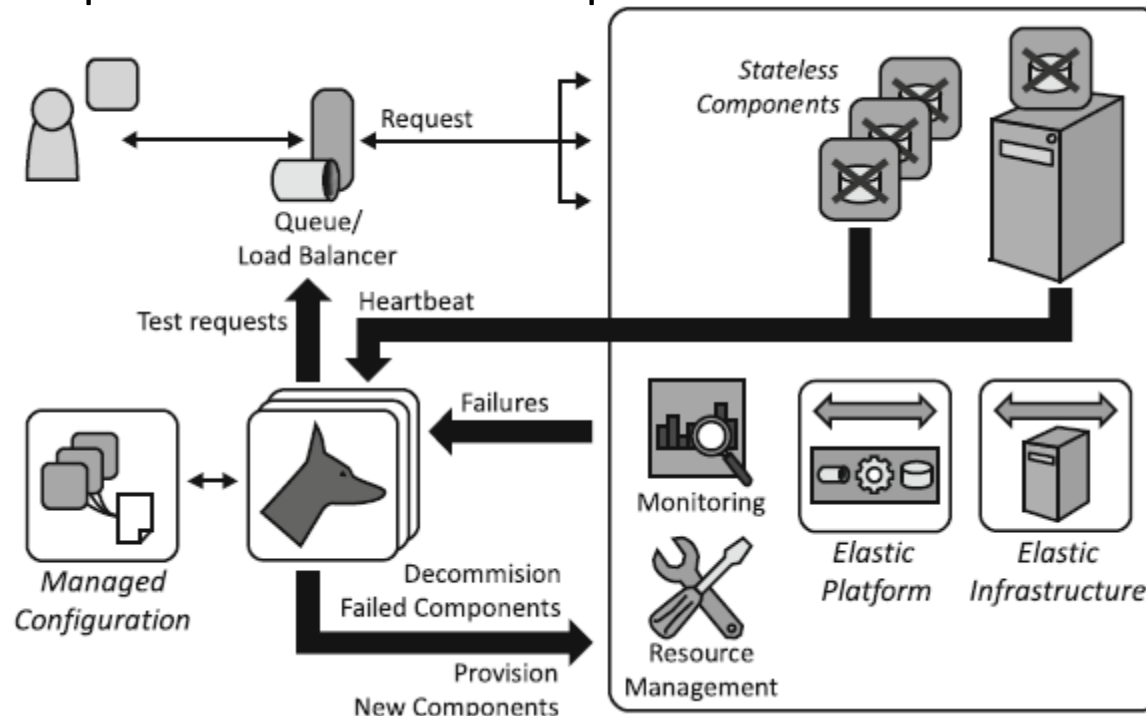
- AWS SLA

- Solution: The provider assures the *overall* availability for the provided env
  - no notion of availability for individual app components or virtual servers
  - customer incorporates failures in app architectures and runtime mgmt
  - AWS SLA: unavailability = {all running instances cannot be reached > 5 min and no replacement instances can be provisioned}
    - Deploy redundant instances to multiple “availability zones”



# Watchdog

- Intent: HA applications **cope with failures by monitoring and replacing** app component instances if the provider-assured availability is insufficient



- AWS CloudWatch
- RightScale
- Scalr

- Solution: Scale-out *stateless* components (virtual servers and/or app components) are monitored by a watchdog component and replaced in case of failures
  - Watchdog itself must be implemented as HA



# CloudWatch

Name	Volume ID	Capacity	Volume Type	Snapshot	Created	Zone	State	Alarm	Attachment Information
empty	vol-b12f11da	360 GiB	standard	snap-c763b7a6	2011-07-16T13:51:10Z	us-east-1b	in-use	none	i-439b5c25 (adm-1000):sdi4

1 Volume selected

Volume: vol-b12f11da

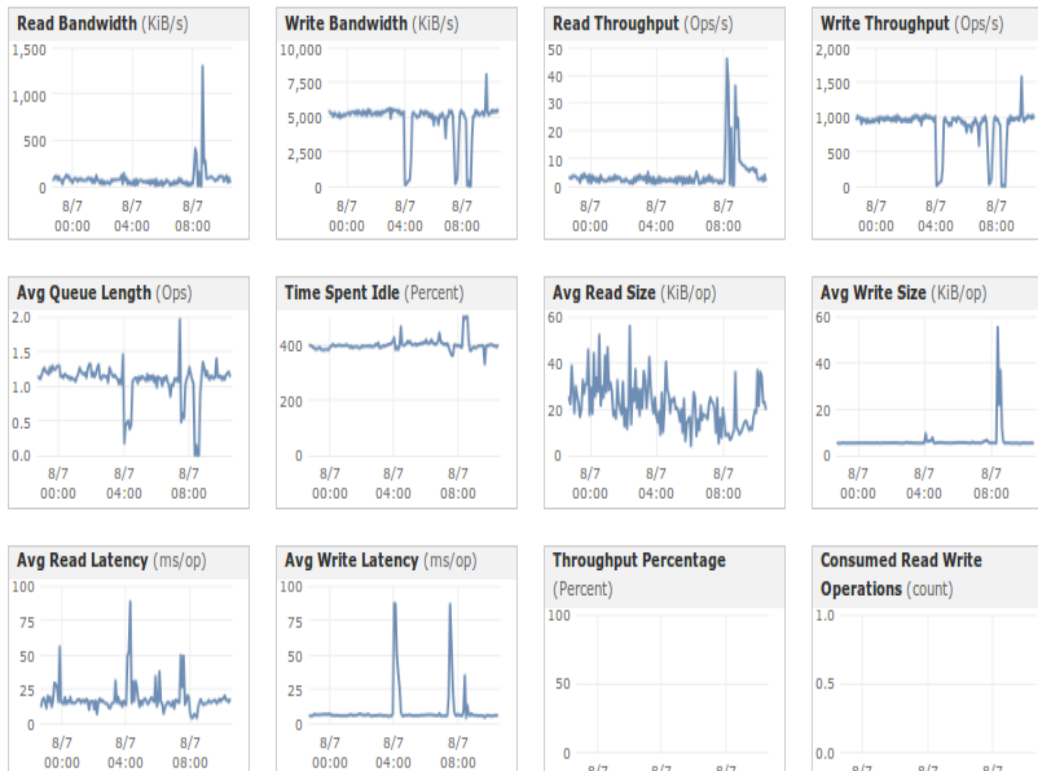
Details Status Checks **Monitoring** Tags

CloudWatch alarms: No alarms configured for vol-b12f11da [View all CloudWatch alarms](#)

Create Alarm

CloudWatch metrics: Times are displayed in UTC.

Time Range: Last 12 Hours



- Mgmt: monitor AWS cloud resources and apps

- Standard metrics
- Custom metrics (from customer's apps and services)


- Policy-based **Auto Scaling**

- Auto scaling group
- Add/remove (EC2) instances **dynamically** based on metrics

- util%
- queue length
- etc.

# Coffee Shop - Elasticity and Resiliency

What shall happen if workload changes or something fails?



# Lessons - Elasticity and Resiliency

- Analyze availability assured by provider (node-based or environment-based)
- In case of low node-based availability and environment-based availability implement a watchdog
- Use msgs and reqs to determine necessary component instances

# Coffee Shop - Summary






















Decomposition

Workload

Data (State)

Component  
Refinement

Elasticity and  
Resiliency



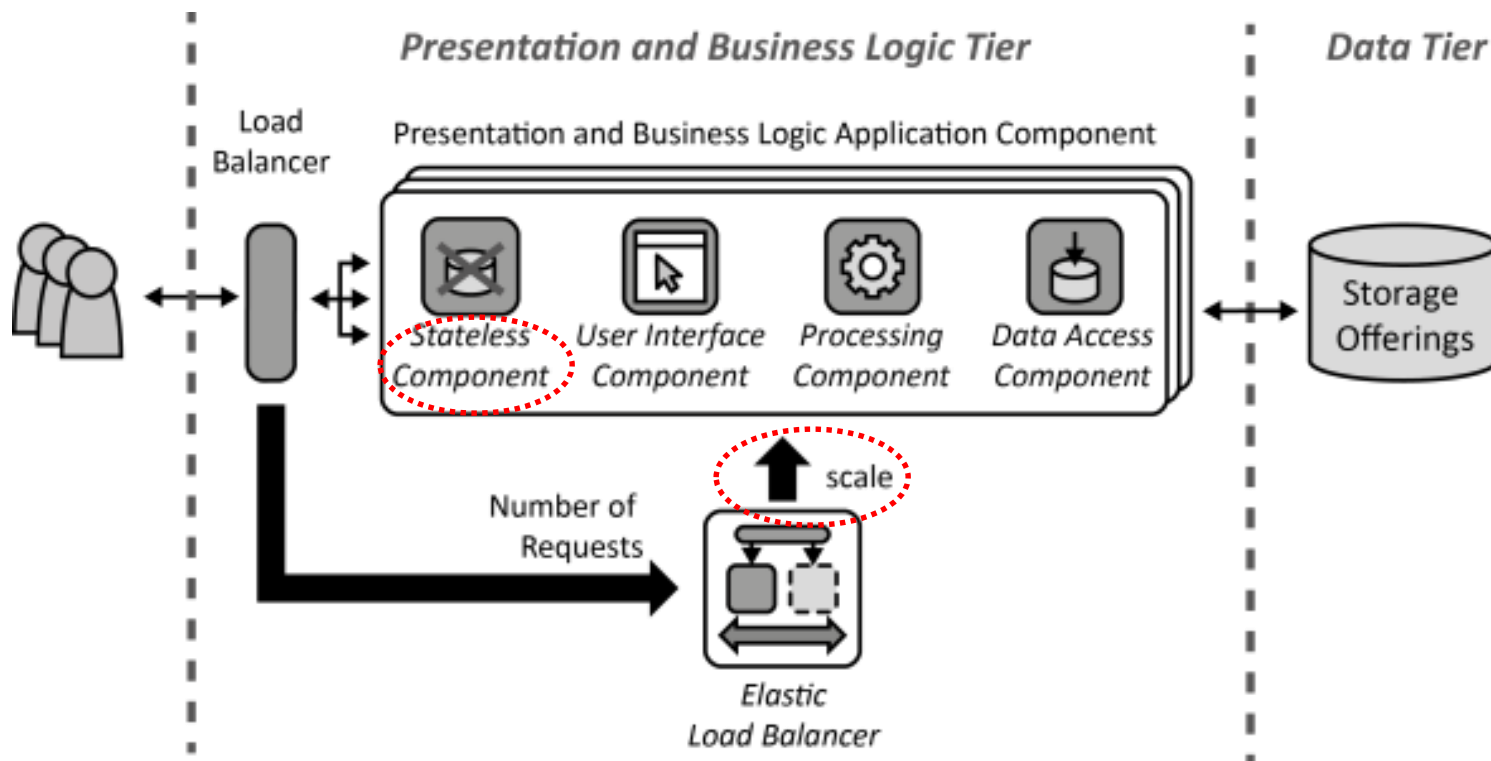
# Composite Cloud Applications

- Native Cloud Applications
  - Two-Tier Cloud Application
  - Three-Tier Cloud Application
  - Content Distribution Network
- Hybrid Cloud Applications
  - Hybrid User Interface
  - Hybrid Processing
  - Hybrid Data
  - Hybrid Backup
  - Hybrid Backend
  - Hybrid Application Functions



# Two-Tier Cloud App

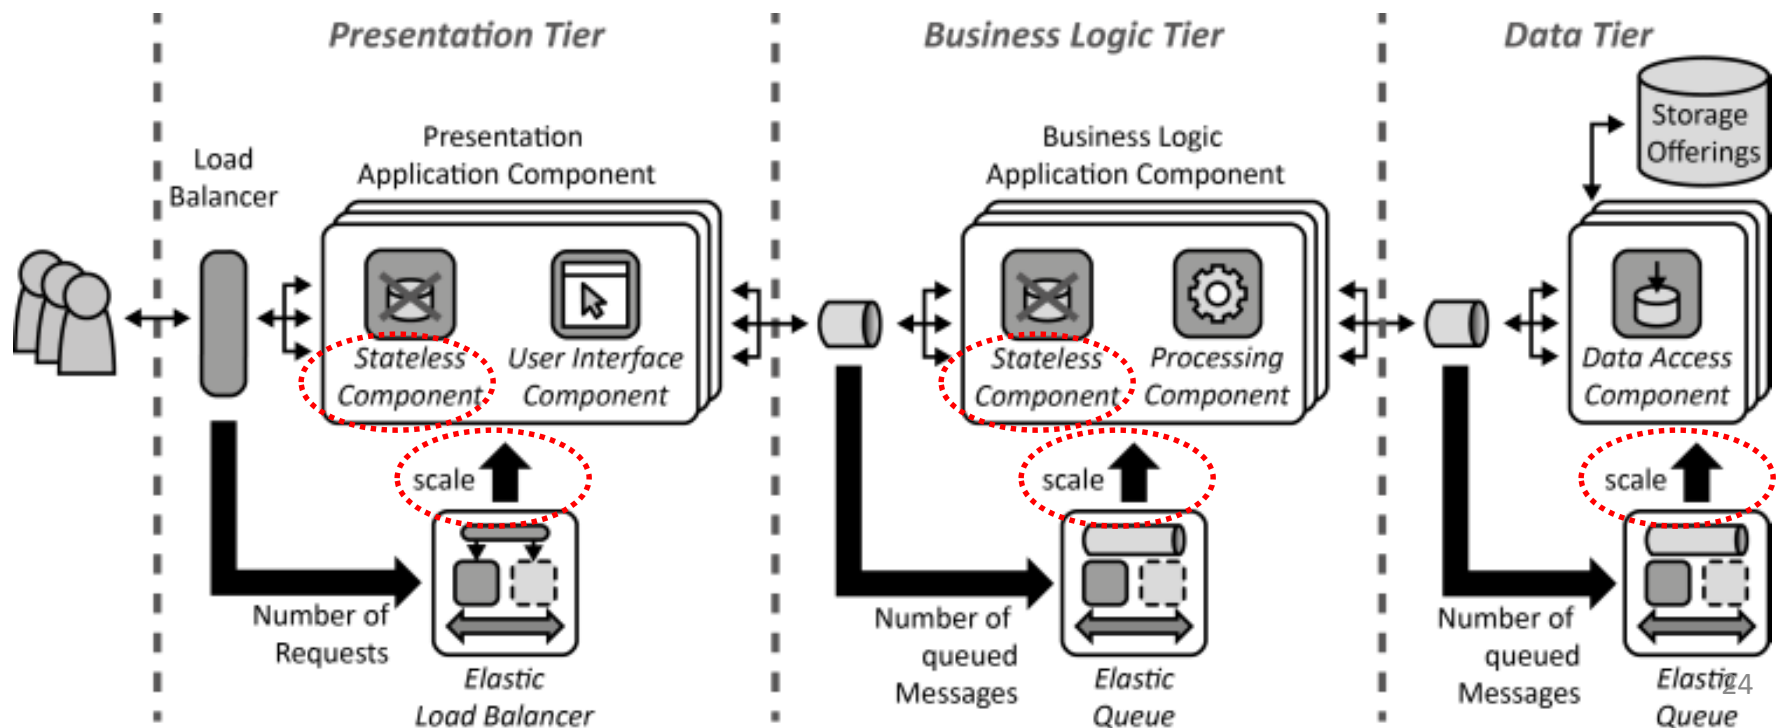
- Intent: Presentation and business logic is **bundled to one stateless tier** that is easy to scale. This tier is separated from the data tier that is harder to scale and often handled by a provider-supplied storage offering





# Three-Tier Cloud App


- Intent: The presentation, business logic, and data handling is realized as **separate tiers** to scale **stateless** presentation and compute-intensive processing independently of the data tier, which is harder to scale and often handled by cloud provider





# Content Distribution Network

- Intent: Apps component **instances** and **data** handled by them are **globally distributed** to meet the access performance required by a global user group




- Solution: Content replicas are established in different physical locations of one or multiple clouds. During distribution of replicas, the topology of distribution networks is considered to ensure locality for all users. Replicas are updated from a central location





# Hybrid User Interface

- Intent: **Varying** workload from a user group interacting **async** w/ an app is handled in an **elastic** env while the remainder of an app resides in a **static** env

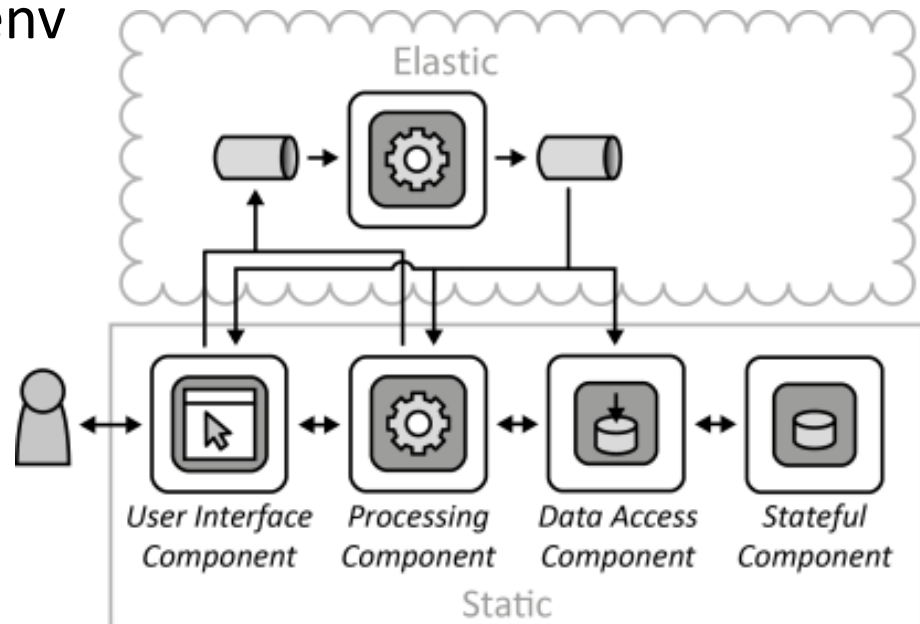


- Solution: The UI Component generating varying workload is hosted in an elastic cloud env. Other app components are hosted in a static env. The UI is integrated w/ the remainder of app in a decoupled fashion using **msg (loose coupling)**



# Hybrid Processing

- Intent: Processing functionality that experiences **varying** workload is hosted in an **elastic** cloud while the remainder of app resides in a **static** env

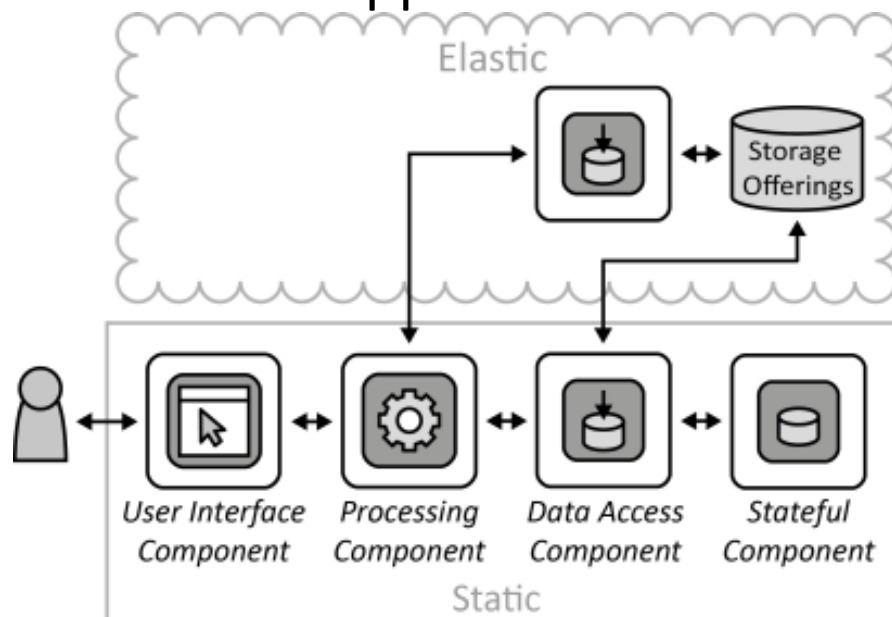


- Solution: The Processing Components experiencing varying workloads are provisioned in an elastic cloud. Loose Coupling is ensured by exchanging info b/w the hosting envs **async** via **msgs**



# Hybrid Data

- Intent: Data of **varying** size is hosted in an **elastic** cloud while the remainder of an app resides in a **static** env

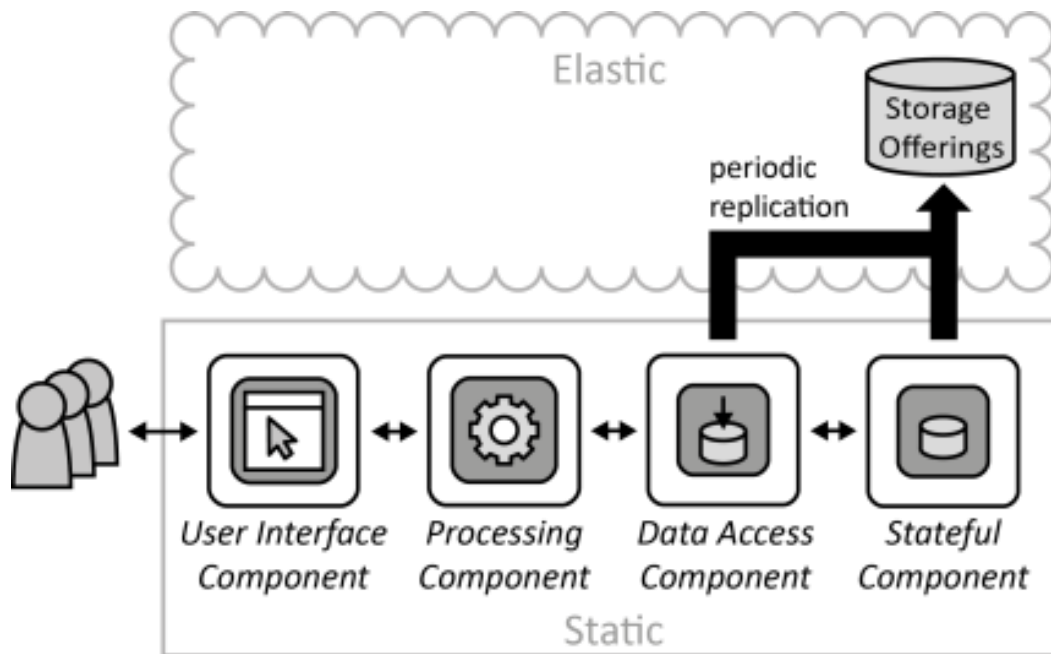


- Solution: Data whose varying size makes it unsuitable for hosting in a static env is handled by Storage Offerings in an elastic cloud. Data is accessed by Data Access Components that are either hosted in the static env or in the elastic env



# Hybrid Backup

- Intent: Data is periodically extracted from an app to be archived in an elastic cloud for **disaster recovery** purposes




- Solution: A Distributed Application is hosted in a local static env of the company. Data handled by **Stateful** Components is periodically extracted and replicated to a cloud storage offering



# Hybrid Backend

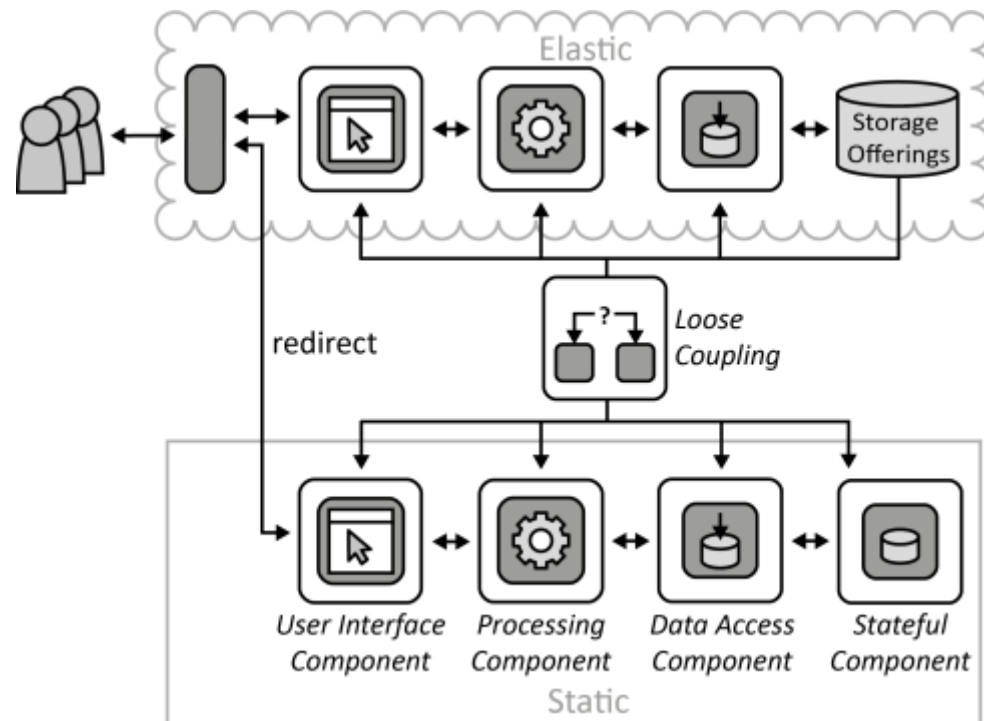
- Intent: Backend functionality comprised of data intensive processing and data storage is experiencing **varying** workloads and is hosted in an **elastic** cloud while the rest of an app is hosted in a **static** data center



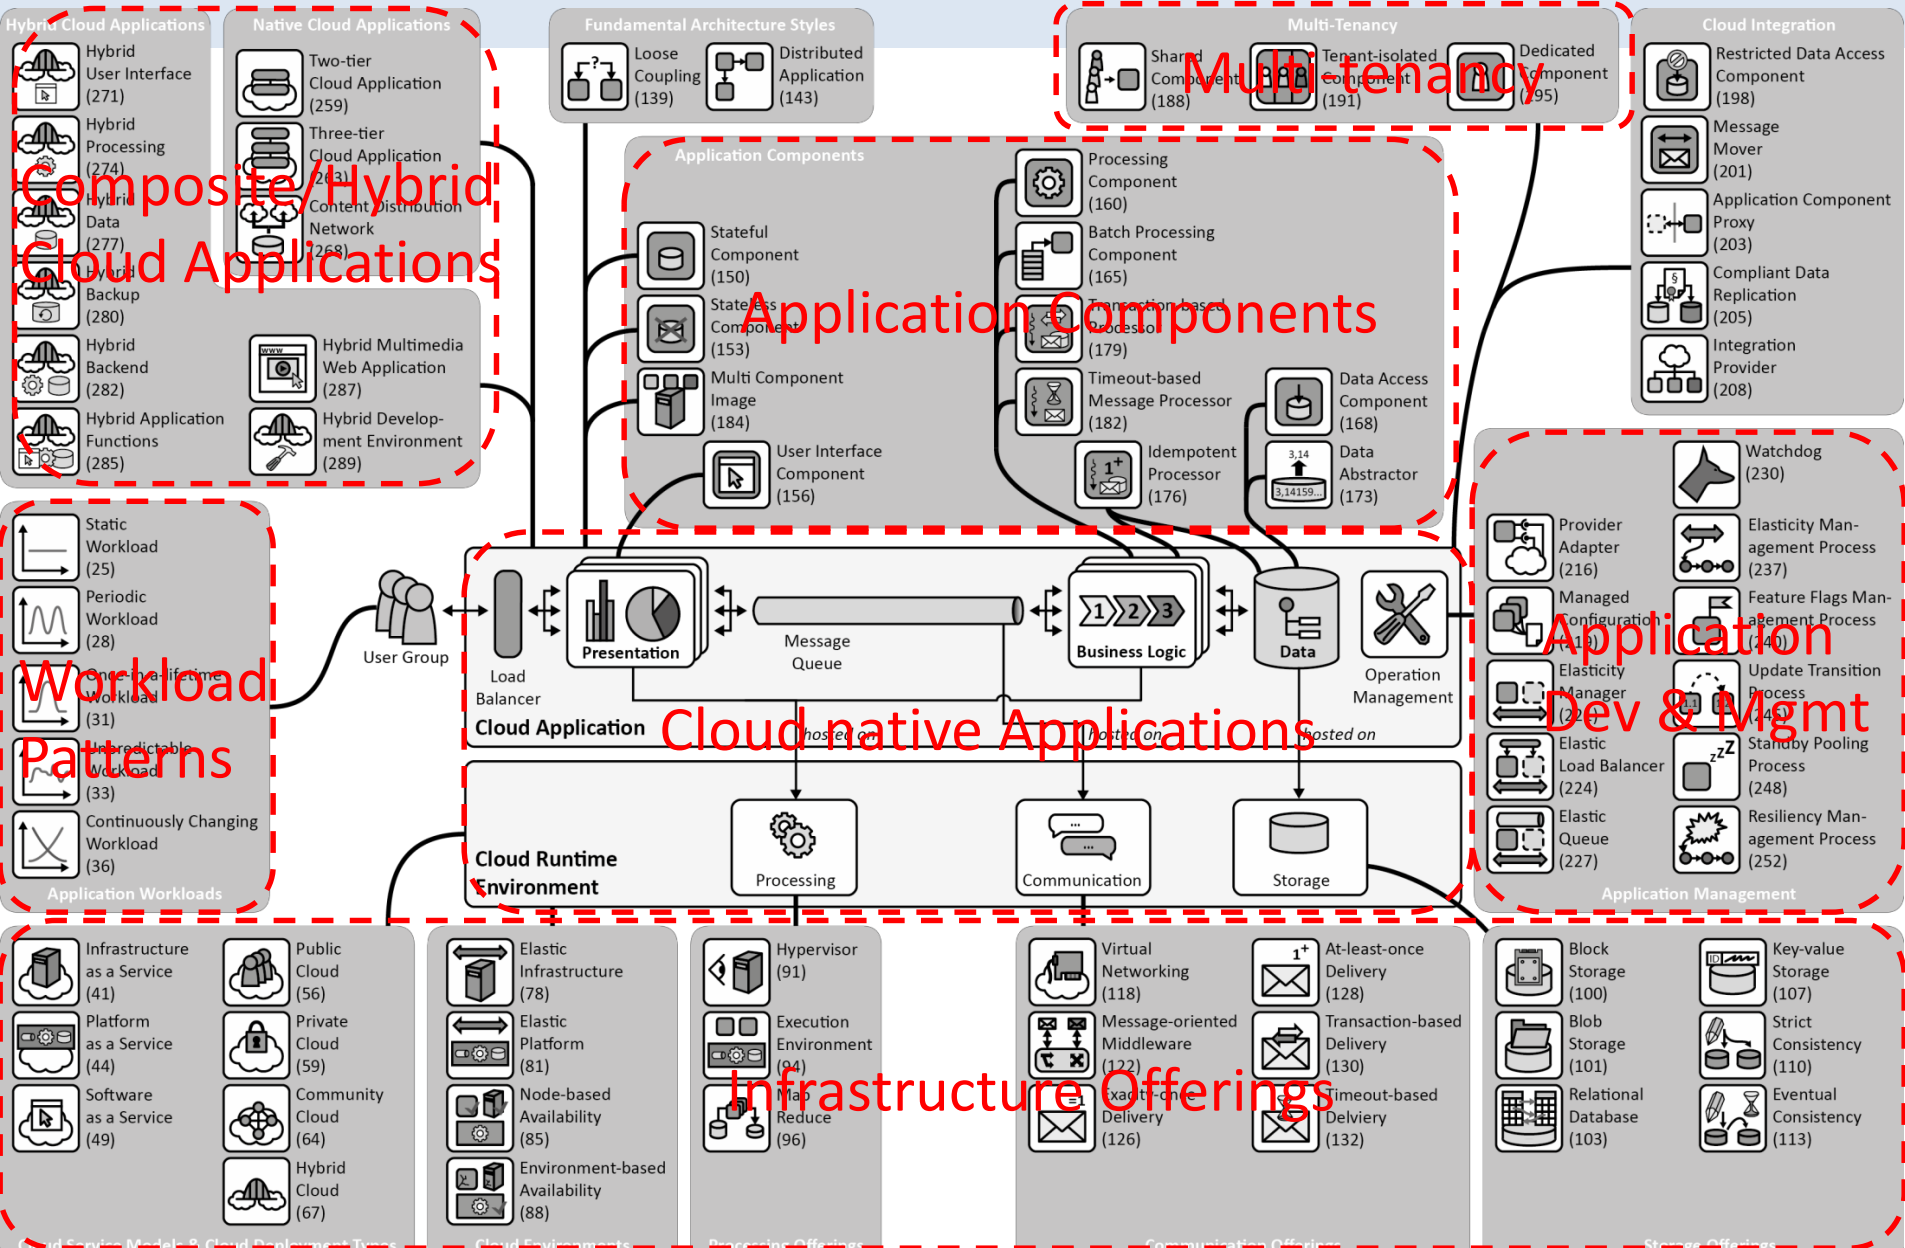


# Hybrid App Functions

- Intent: Some app functionality provided by UIs, processing, and data handling is experiencing **varying** workload and is hosted in an **elastic** cloud while **other** app functionality of the same type is hosted in a **static** env



# <http://www.cloudcomputingpatterns.org>



# References

- <http://www.cloudcomputingpatterns.org/>
- Christoph Fehling, Frank Leymann, Ralph Retter, Walter Schupeck, and Peter Arbitter, *Cloud Computing Patterns: Fundamentals to Design, Build, and Manage Cloud Applications*. Springer; 2014
  - <http://www.springer.com/978-3-7091-1567-1>
- The coffee shop example is adapted from
  - <https://indico.scc.kit.edu/indico/event/26/session/1/contribution/12/material/slides/0.pdf>
  - <http://www.sei.cmu.edu/library/assets/presentations/retter-saturn2013.pdf>