

# CMPE 282 Cloud Services

## ***Hadoop***

Instructor: Kong Li

# Content

- MapReduce
- Apache Hadoop
- HDFS
  - Name space
  - Data access
  - HA – Active-standby namenode
  - HA - Data Replication
- Daemons
- YARN
- MapReduce Job
- Fault Tolerance
- Speculative Execution
- User application
- Spark
- Storm



# Big Data

- Characteristics
  - Volume: MB, GB, TB, PB, etc
  - Variety: different forms or types
  - Velocity: batch, near realtime, realtime
- Search for actionable insights
  - Regardless of structured, semi-structured, or unstructured data
  - Q: How to analyze structured, semi-structured, and unstructured data?
- Evolution: Batch ➔ real time ➔ prediction
- Tools
  - Generic: NoSQL, SQL, search
  - Batch: MapReduce, Hive, Pig, etc.
  - Real time / streaming: Spark (streaming), Storm, etc
  - Machine learning: Mahout, Spark ML, etc
- Q: how to use the right tool for the job?
  - <http://www.slideshare.net/AmazonWebServices/aws-november-webinar-series-architectural-patterns-best-practices-for-big-data-on-aws>

# MapReduce

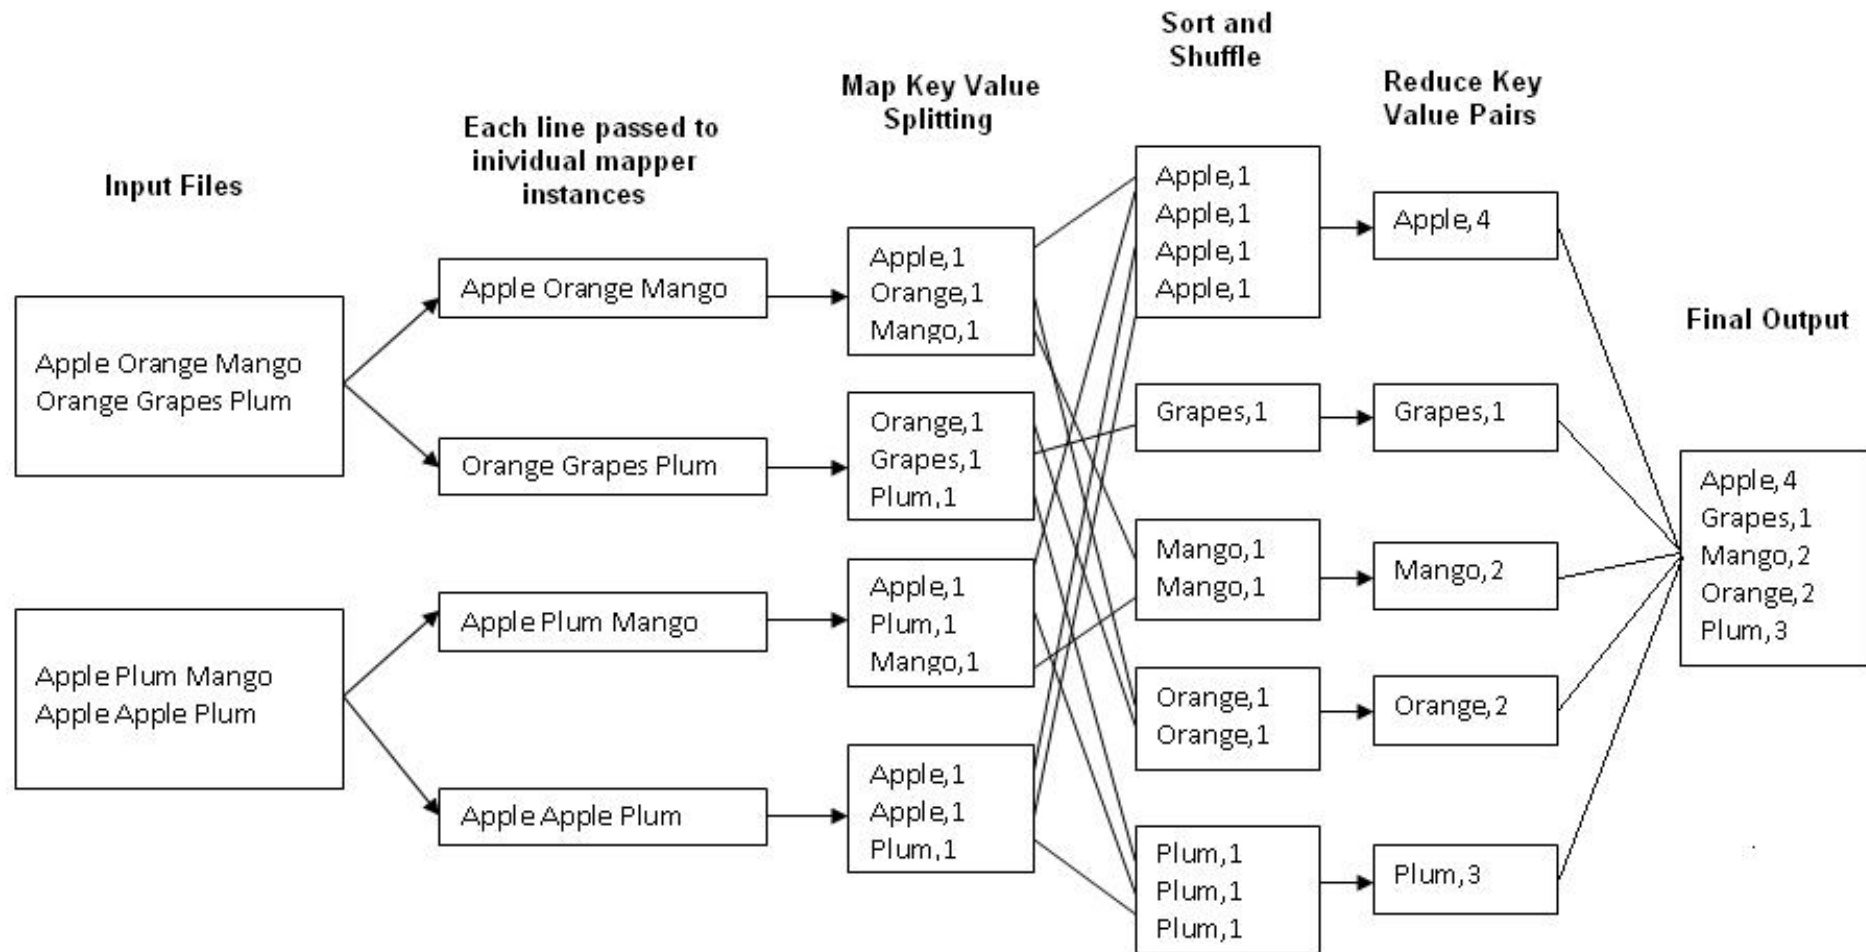
- New programming paradigm
  - Purpose: mine large datasets
    - Structured, semi-structured, unstructured
  - Implementation: computation cluster
- Two phases/stages
  - Map: applies to all the members of the dataset and returns a list of results, executed in parallel
  - Reduce: collates and resolves the results from one or more mapping operations, executed in parallel

Google calls it:	Hadoop equivalent:
MapReduce	Hadoop
GFS	HDFS
Bigtable	Hbase
Chubby	Zookeeper

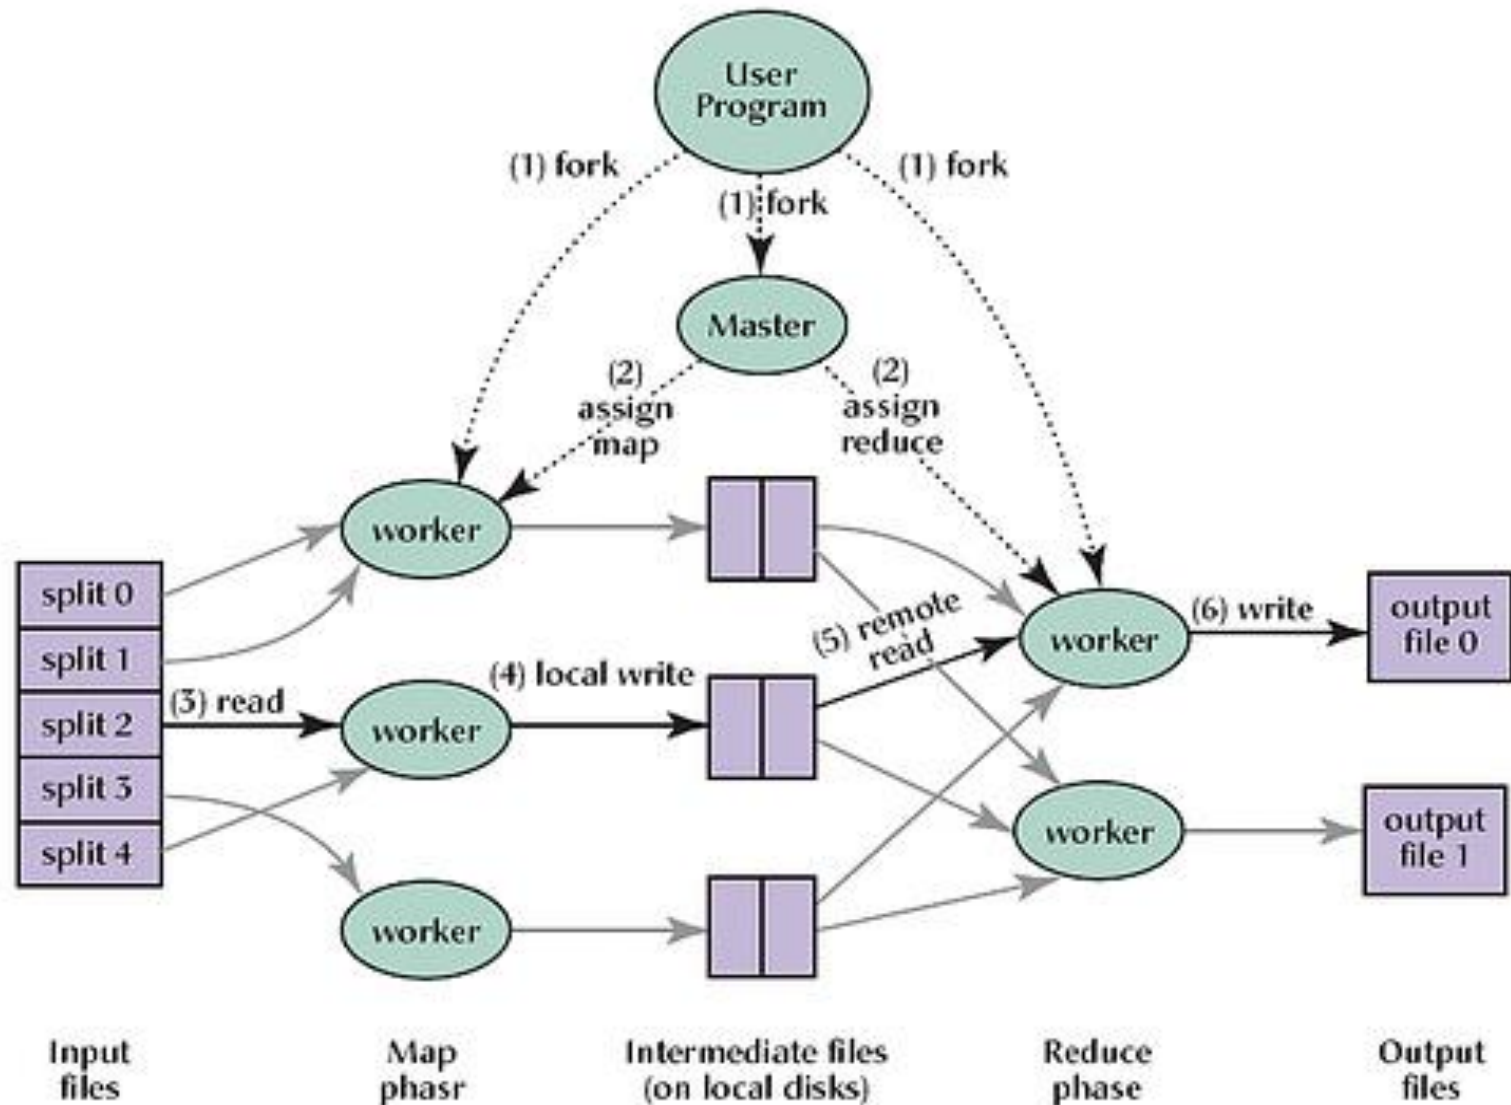
# MapReduce (cont'd)

- Details
  - Large datasets are split into subsets called splits
  - Map: **parallelized** operation: splits → intermediate results
    - $\text{map}(k1, v1)$  applied to each split → **list of  $(k2, v2)$**
  - Shuffle/sort intermediate results into new splits
    - Group **list of  $(k2, v2)$** 's by  $k2$  → **list of  $(k2, \text{list}(v2))$**
  - Reduce: **parallelized** operation: intermediate results → final result
    - $\text{reduce}(k2, \text{list}(v2)) \rightarrow v3$
  - **All values are processed *independently***
  - **final results: same as sequential execution on the entire dataset**  
(assuming deterministic map/reduce)
- Separate app (business logic) from multi-processing logic
  - MapReduce framework: dispatching, scalability, locking, and logic flow
  - App: the business logic w/o worrying about infrastructure or scalability issues

# MapReduce: Word Count

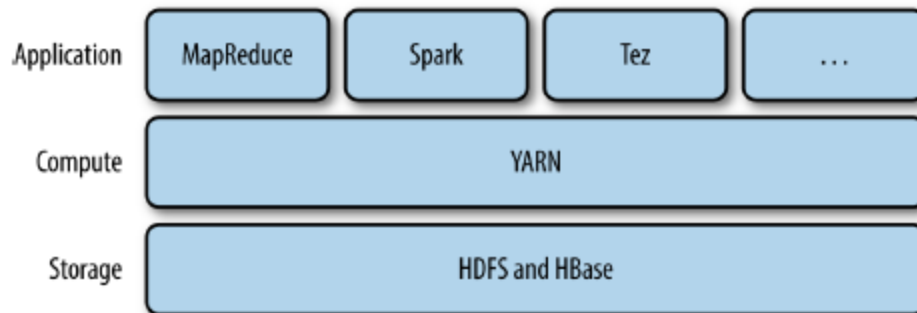


# MapReduce (cont'd)



# Hadoop

- Distributed computing framework
- Large scale: HDFS + YARN + MapReduce
  - HDFS: scalable storage layer
  - **YARN**: cluster compute layer, cluster resource mgmt
  - MapReduce app: map/reduce tasks on cluster of machines
    - New programming model: computation on top of set of key/value pair
    - **Batch** processing: A job runs until a given set of data has been processed
- For semi-structured, unstructured, as well as structured data
  - Can use RDBMS (w/ SQL) and NoSQL together
- **Built-in resiliency, fault tolerance**
- Batch → streaming/real time?
  - Apache Spark
  - Apache Storm





# Apache Hadoop Ecosystem



## Apache Hadoop Ecosystem



**Ambari**

Provisioning, Managing and Monitoring Hadoop Clusters



**Scoop**

Data Exchange



**Zookeeper**

Coordination



**Oozie**

Workflow



**Pig**

Scripting



**Mahout**

Machine Learning

**R Connectors**

Statistics



**Hive**

SQL Query



**Hbase**

Columnar Store



**YARN Map Reduce v2**

Distributed Processing Framework

**HDFS**

Hadoop Distributed File System



# Apache Hadoop Ecosystem (cont'd)

- **HDFS**: scalable distributed file system
- **YARN**: framework for job scheduling and cluster resource mgmt
- **MapReduce**: YARN-based system for parallel processing of large data sets
- **HBase**: scalable, distributed NoSQL DB on HDFS for large tables
- **Hive**: data warehouse infrastructure providing data summarization and ad hoc querying
- **Cassandra**: scalable, multi-master, HA, NoSQL DB
- **Pig**: high-level data-flow language and execution framework for parallel computation
- **ZooKeeper**: a coordination service for distributed apps
- **Oozie**: workflow scheduler system for MR jobs using DAGs
- **Ambari**: provisioning, managing, and monitoring Apache Hadoop clusters
- **Mahout**: scalable machine learning and data mining library
- **Spark**: distributed framework for in-memory analytics on large data sets; co-exist w/ Hadoop - SQL, machine learning, *stream* processing, and graph computation
- **Storm**: distributed real-time computation framework for event stream processing
  - real-time analytics, online machine learning, continuous computation, ETL

# HDFS


- A distributed file system: makes tradeoffs that are good for MapReduce
  - Single virtual file system spread over many machines
  - Traditional hierarchical file organization: directory, file
- Goals:
  - Simple Coherency Model: **write-once-read-many** access model for files
  - Large data sets
  - **Moving computation is cheaper than moving data**
- Good for:
  - **Very large read-only or append-only files** (individual file size GB or TB)
  - **Sequential access patterns**
- Not so good for:
  - Storing lots of small files
  - Low-latency access
  - Multiple writers
  - Writing to arbitrary offsets in the file
- **Optimized for sequential read and local accesses**

# HDFS: namespace

- Name space: separated from FS of OS
  - Files are stored as sets of (large) blocks on top of FS of OS
    - Blocks are replicated for durability and availability (discussed later)
  - Files in HDFS are not visible in the normal FS
    - `hadoop fs -ls`
  - Only the blocks and the block metadata are visible in FS of OS
  - Also support 3<sup>rd</sup> party FS e.g., CloudStore and Amazon S3

```
[ahae@carbon ~]$ ls -la /tmp/hadoop-ahae/dfs/data/current/
total 209588
drwxrwxr-x 2 ahae ahae      4096 2013-10-08 15:46 .
drwxrwxr-x 5 ahae ahae      4096 2013-10-08 15:39 ..
-rw-rw-r-- 1 ahae ahae 11568995 2013-10-08 15:44 blk_-3562426239750716067
-rw-rw-r-- 1 ahae ahae   90391 2013-10-08 15:44 blk_-3562426239750716067_1020.meta
-rw-rw-r-- 1 ahae ahae      4 2013-10-08 15:40 blk_5467088600876920840
-rw-rw-r-- 1 ahae ahae     11 2013-10-08 15:40 blk_5467088600876920840_1019.meta
-rw-rw-r-- 1 ahae ahae 67108864 2013-10-08 15:44 blk_7080460240917416109
-rw-rw-r-- 1 ahae ahae   524295 2013-10-08 15:44 blk_7080460240917416109_1020.meta
-rw-rw-r-- 1 ahae ahae 67108864 2013-10-08 15:44 blk_-8388309644856805769
-rw-rw-r-- 1 ahae ahae   524295 2013-10-08 15:44 blk_-8388309644856805769_1020.meta
-rw-rw-r-- 1 ahae ahae 67108864 2013-10-08 15:44 blk_-9220415087134372383
-rw-rw-r-- 1 ahae ahae   524295 2013-10-08 15:44 blk_-9220415087134372383_1020.meta
-rw-rw-r-- 1 ahae ahae     158 2013-10-08 15:40 VERSION
```

# HDFS: Data Access

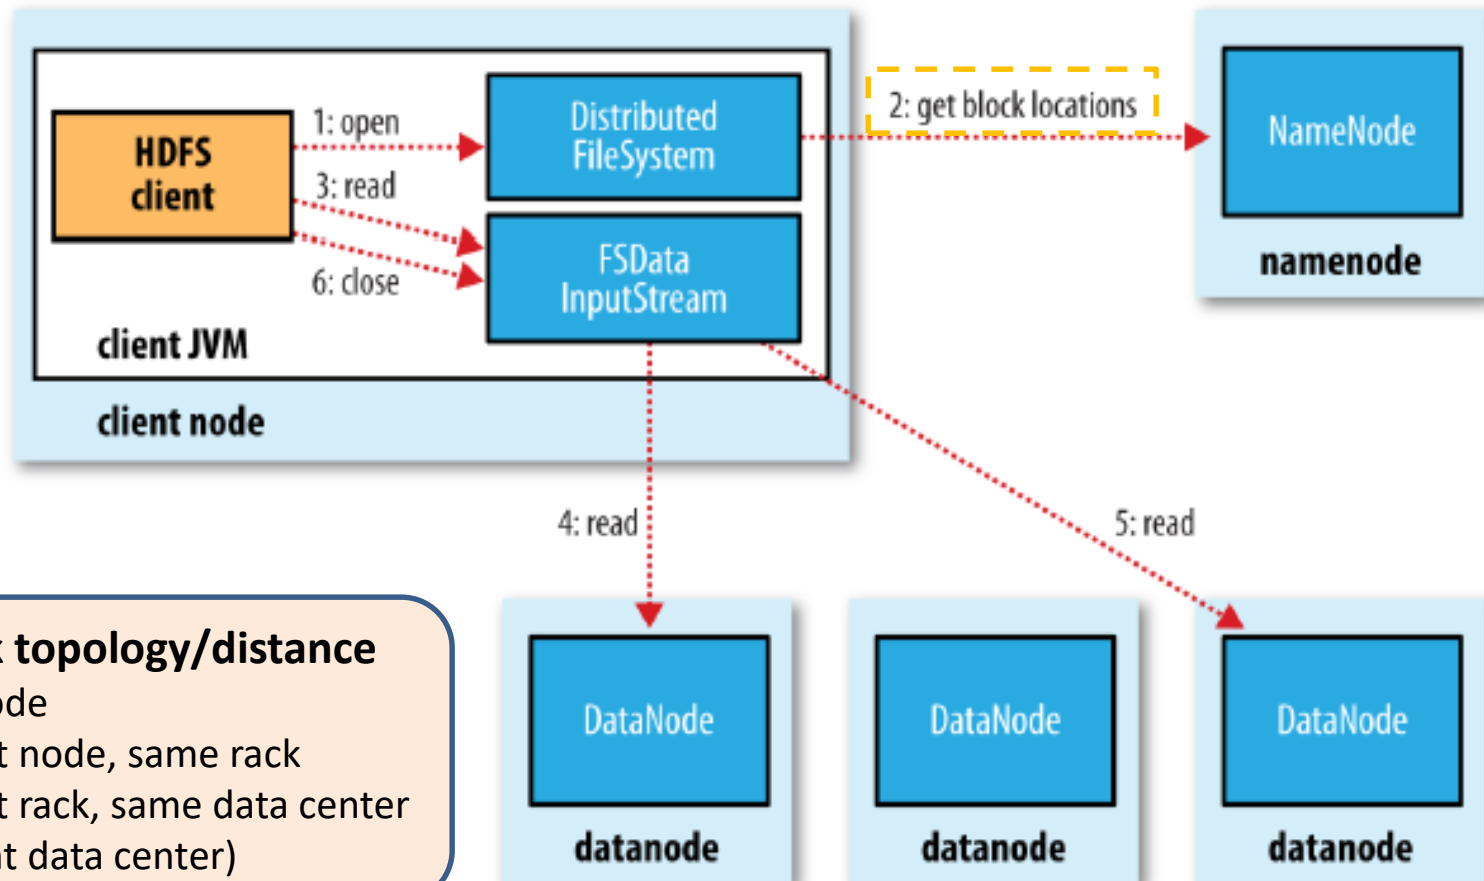


## Secondary namenode

- Not a *real* namenode
- Not for HA
- periodically merge namespace image w/ edit log to prevent the edit log from becoming too large

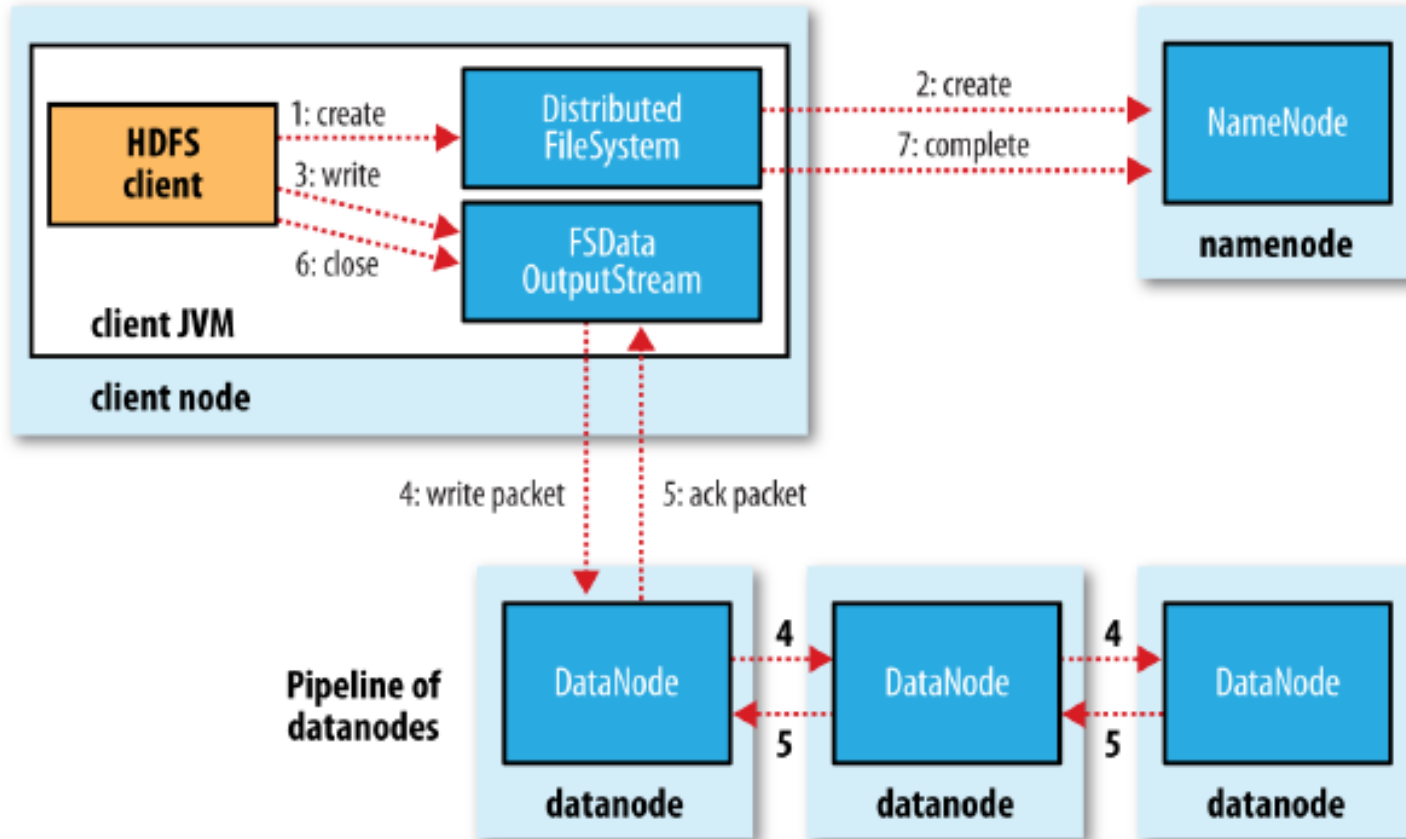
- File == multiple blocks
  - By default, block size = 128 MB
  - Blocks of a given file spread across multiple datanodes
  - `hdfs fsck / -files -blocks`
- Namenode (master)
  - manages FS namespaces: FS tree, metadata of all files/directories
    - persist info in namespace image, edit log
  - Caches datanodes containing which blocks for files
- Datanode (worker)
  - Read/write blocks: requests from clients and namenode
  - Actual data transfer is directly between client & datanode
  - Periodically report blocks to namenode

# HDFS: File Read



- **Step 2**
  - Response: datanodes are **sorted** according to their proximity to the client
  - Client reads local datanode if given blocks are stored locally, or from **closest** datanode

# HDFS: File Write



(Assume replication level = 3)

Rack-aware replica placement (later)

# HDFS: Data Access from Java


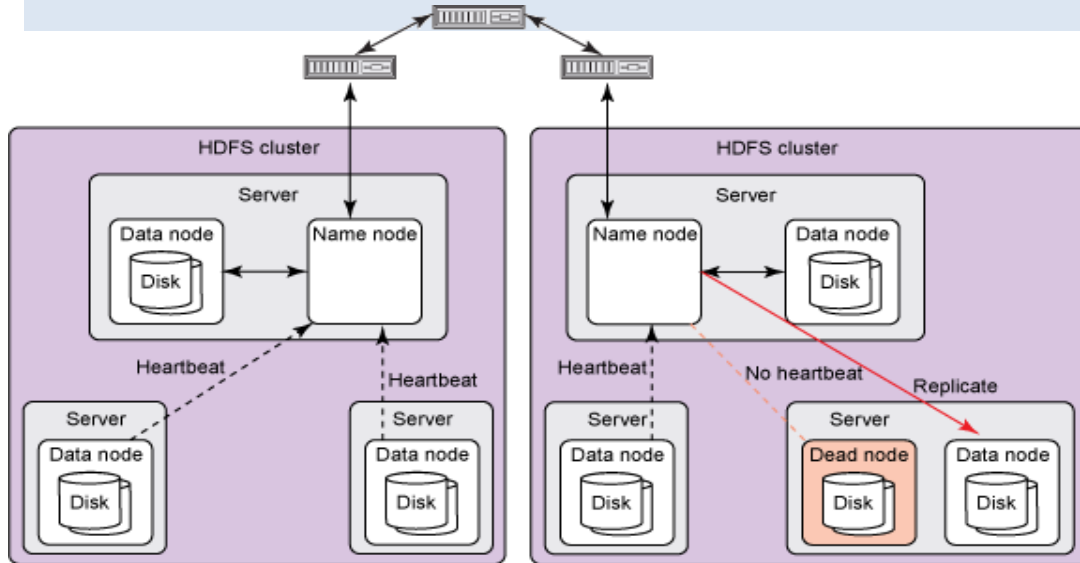
- Programs can read/write HDFS files directly
  - Not needed in MapReduce; I/O is handled by the framework
- Files: represented as URIs
  - Example: `hdfs://localhost/foo/bar/example.txt`
- Access is via the `FileSystem` API
  - To get access to the file: `FileSystem.get()`
  - For reading, call `open()` -- returns `InputStream`
  - For writing, call `create()` -- returns `OutputStream`



# HDFS: HA - Active-Standby NN

- Namenode – single point of failure
- HDFS HA
  - **Active-standby namenodes**: since Hadoop 2
  - Data Replication
- Active-standby namenodes
  - Two namenodes in an active-standby configuration
  - Namenodes use highly available shared storage to share the edit log
  - **Block mappings stored in namenode's memory (not on disk)** → datanodes send block reports to both namenodes
  - Clients configured to handle namenode failover transparently
  - The secondary namenode's role is subsumed by the standby namenode
  - **Failover controlled by ZooKeeper**

# HDFS: HA - Data Replication



- **Namenode**
  - receives periodical heartbeats from datanodes
  - makes all decisions regarding replication of blocks
- **Rack-aware replica placement policy: fault tolerance, locality/performance**
  - 1<sup>st</sup> replica: same node as the client
  - 2<sup>nd</sup> replica: a off-rack node
  - 3<sup>rd</sup> replica: the same rack as the second, but on a different node
  - Others: random nodes in the cluster, but avoid placing too many replicas on the same rack

Tradeoffs

# Daemons

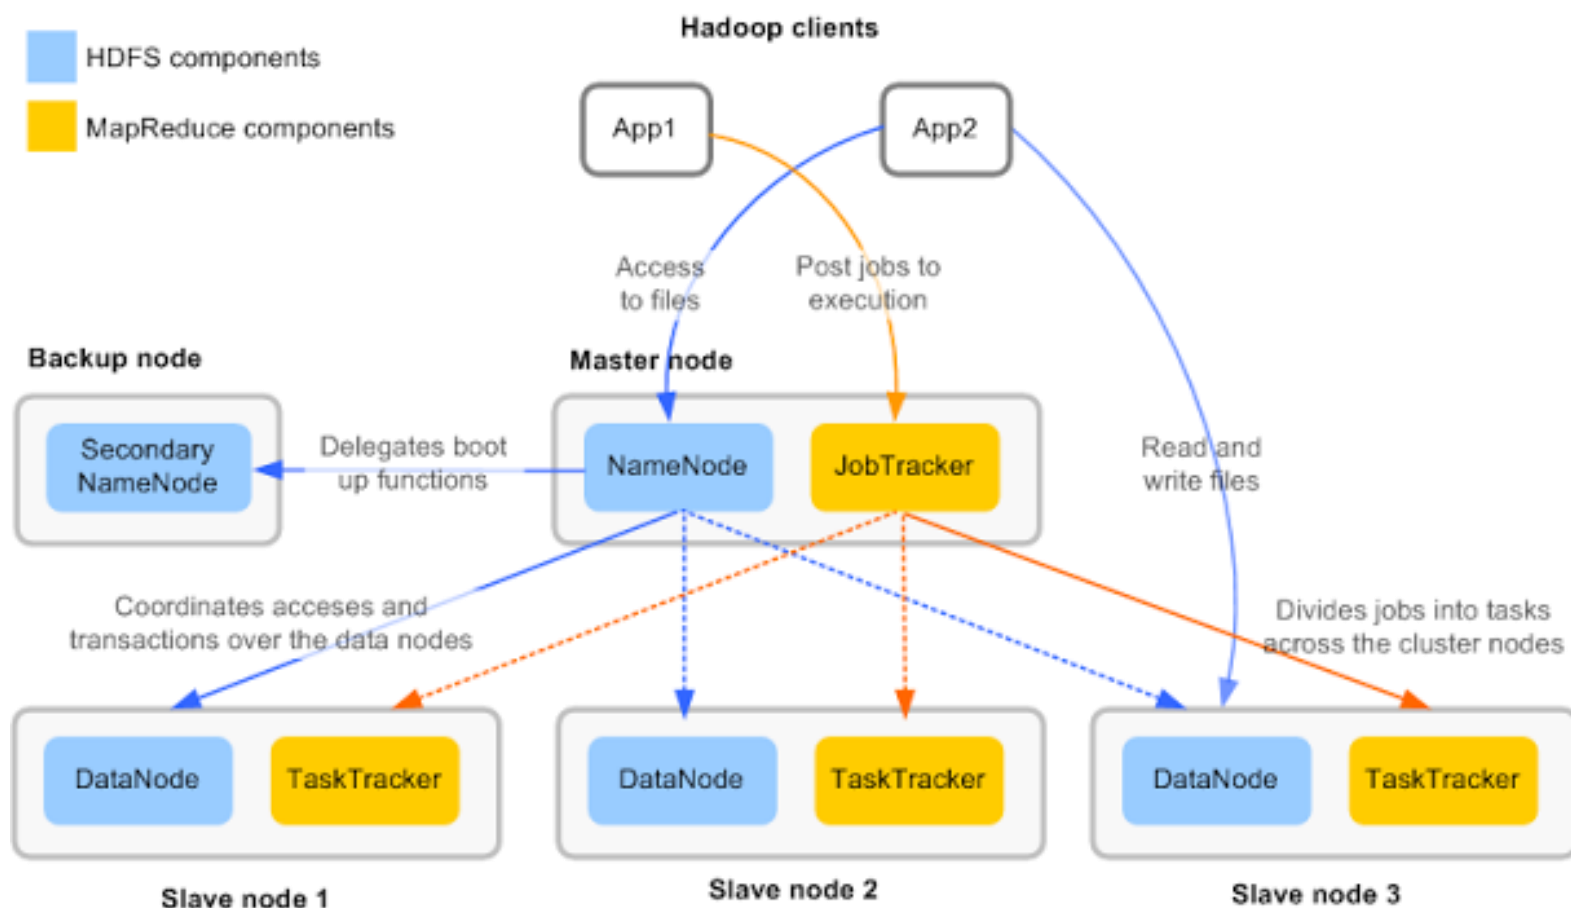
- HDFS
  - Namenode: manages FS namespaces
  - Datanode: stores HDFS blocks
  - v1: Secondary namenode (**not for HA**); v2: standby namenode
- MapReduce
  - JobTracker: one per cluster
    - Coordinates all jobs; schedules tasks to TaskTrackers
    - Monitors progress of TaskTrackers; restarting failed or slow tasks (discussed later)
  - TaskTracker: one per node
    - Runs map tasks and reduce tasks
    - Sends progress report to jobTracker
- YARN or MapReduce 2
  - JobTracker → **global resource manager, per-app application master, timeline server**
  - TaskTracker → **node manager**
  - Multi-tenancy: MapReduce app, Spark app, etc.
- A server may host multiple daemons

# Hadoop Cluster

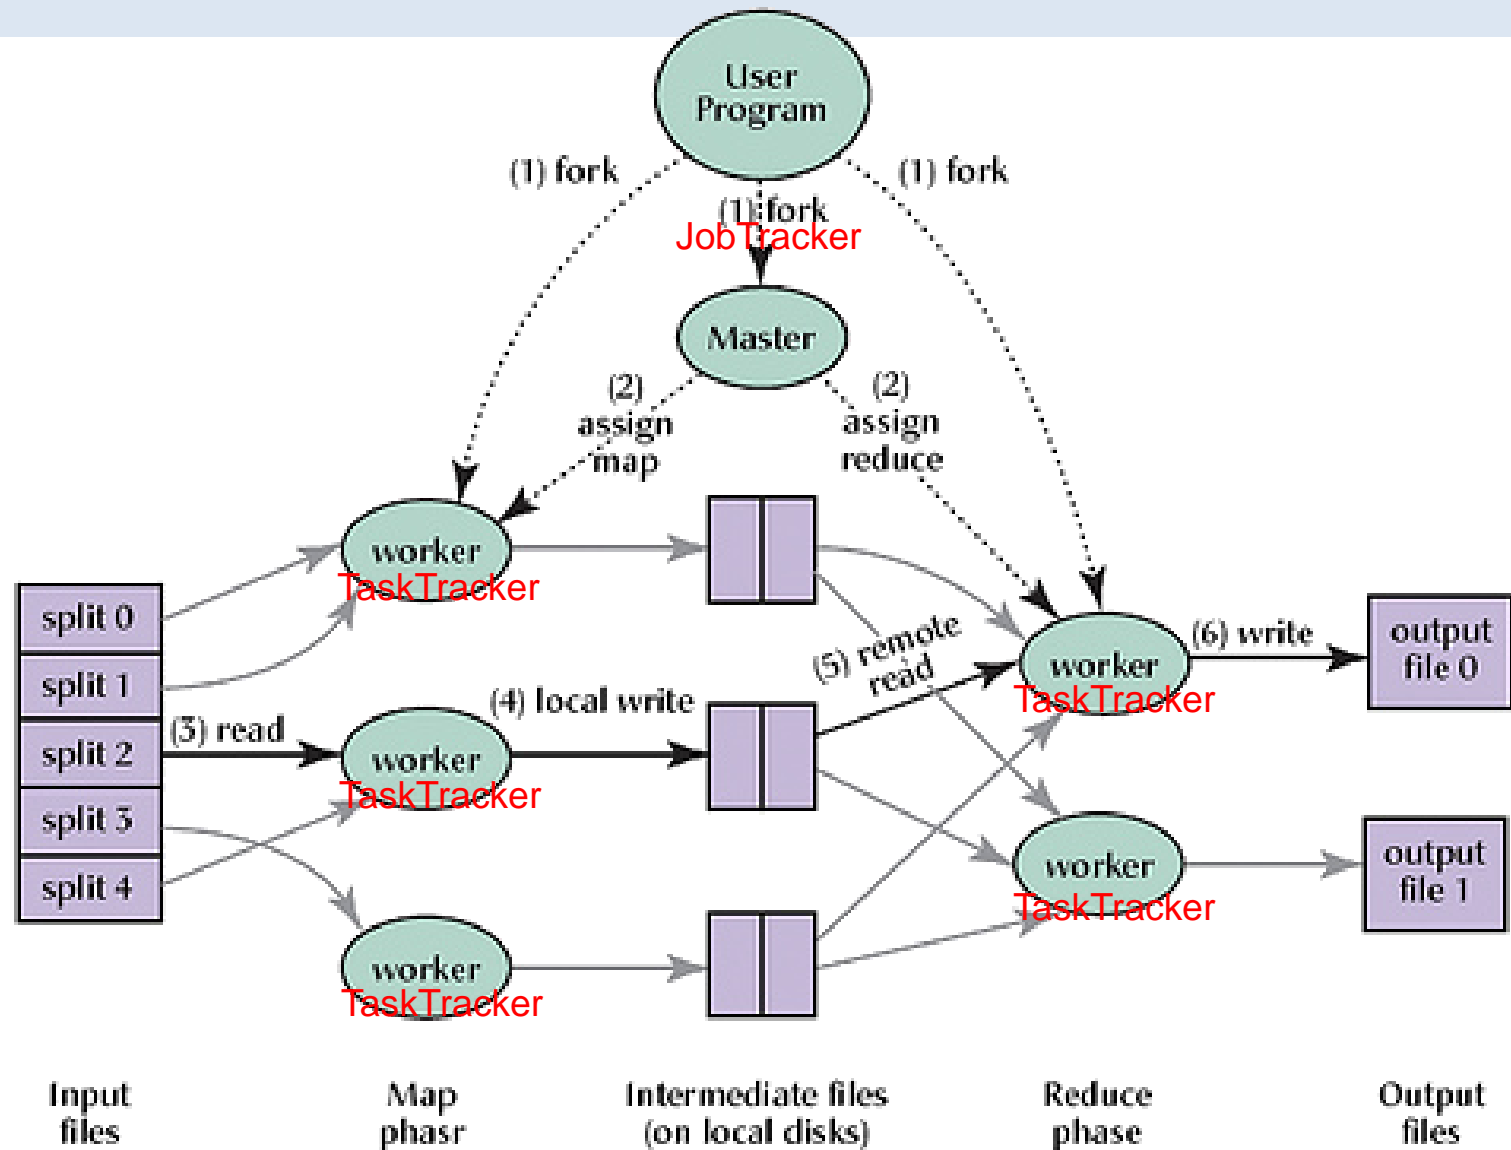


For more information visit me at  
[www.hadooper.blogspot.com](http://www.hadooper.blogspot.com)

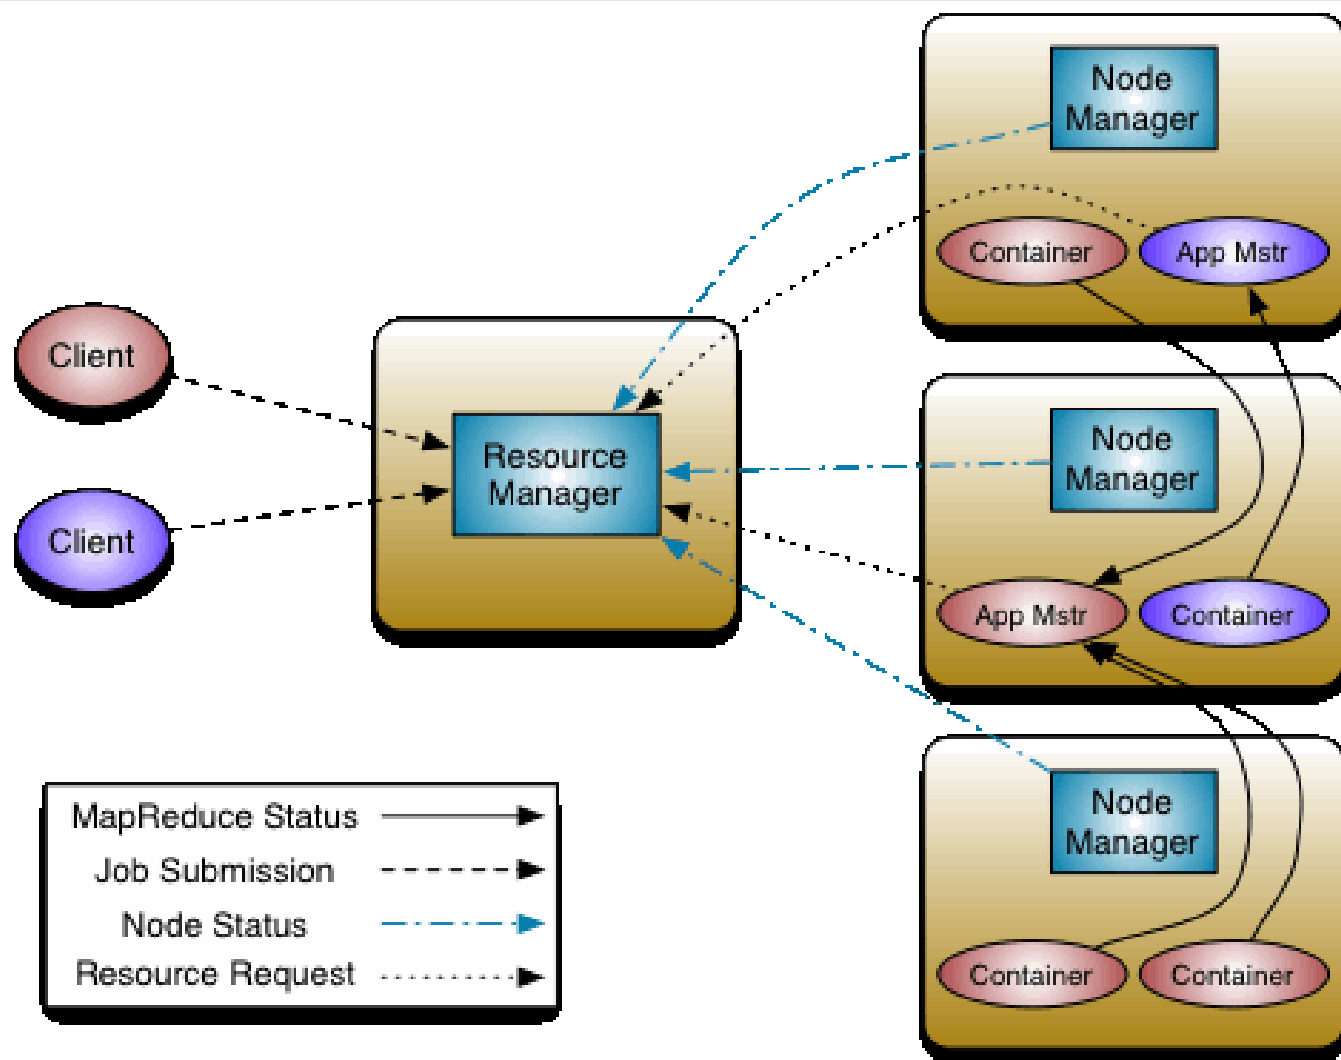
## Hadoop base platform brief



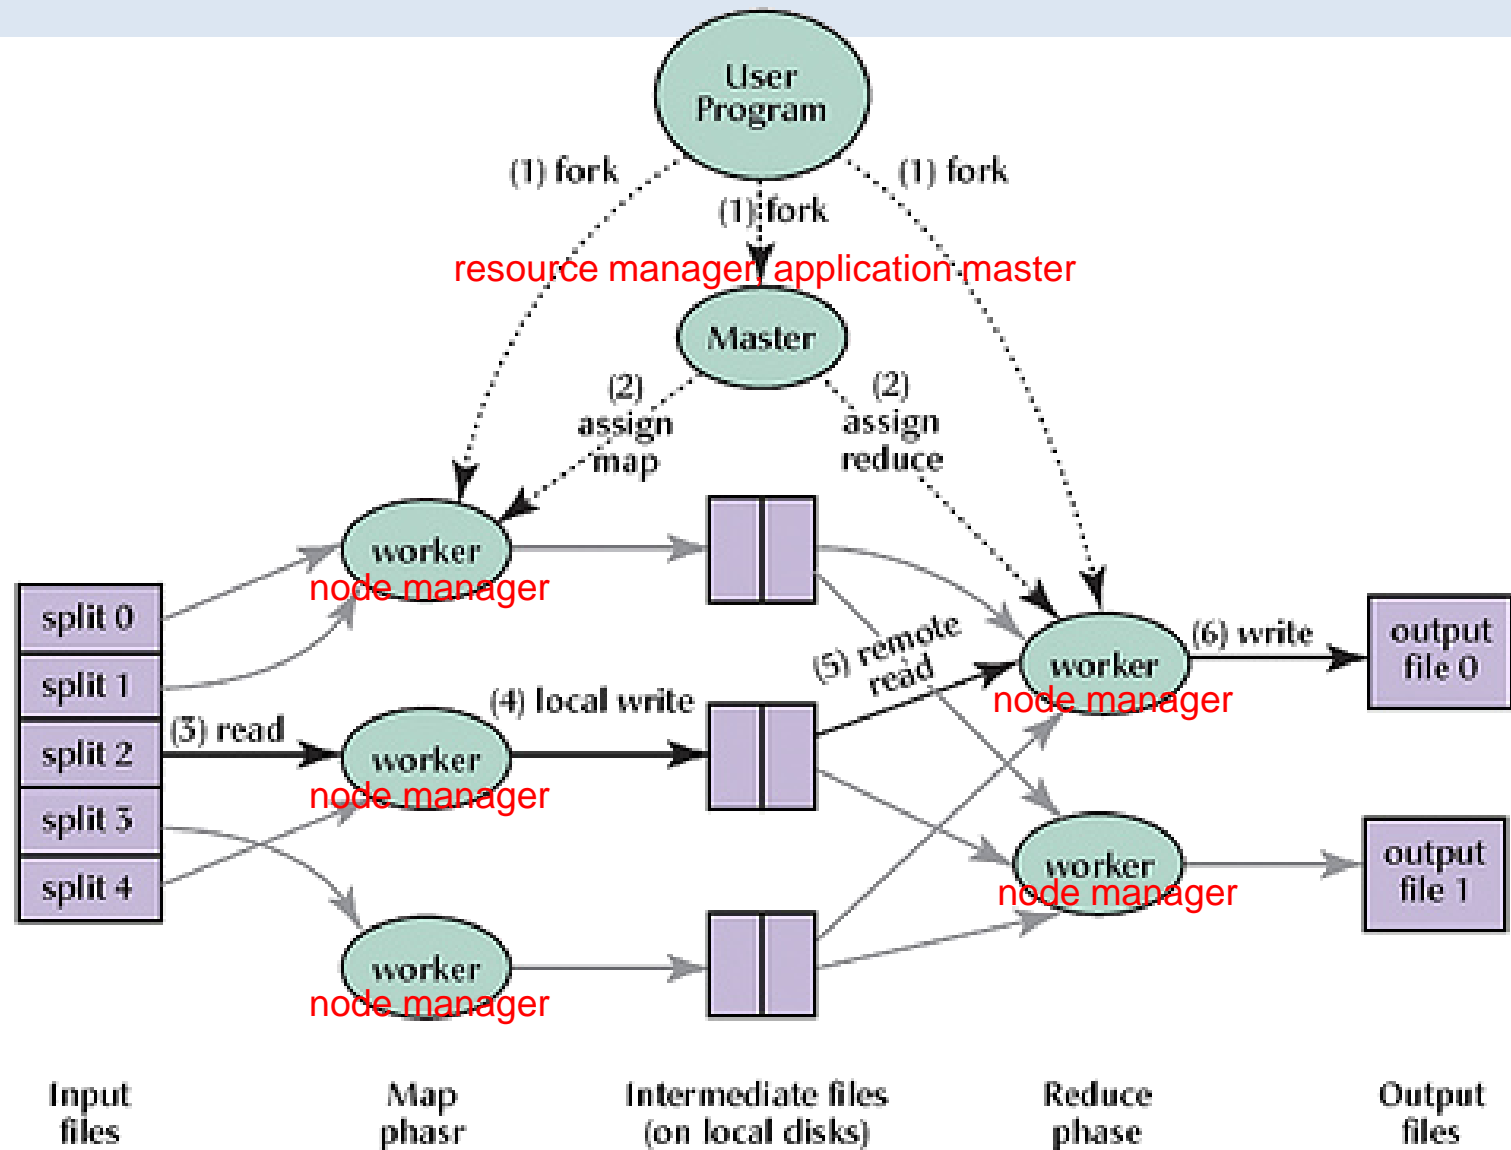
# MapReduce



# Hadoop Cluster (YARN)

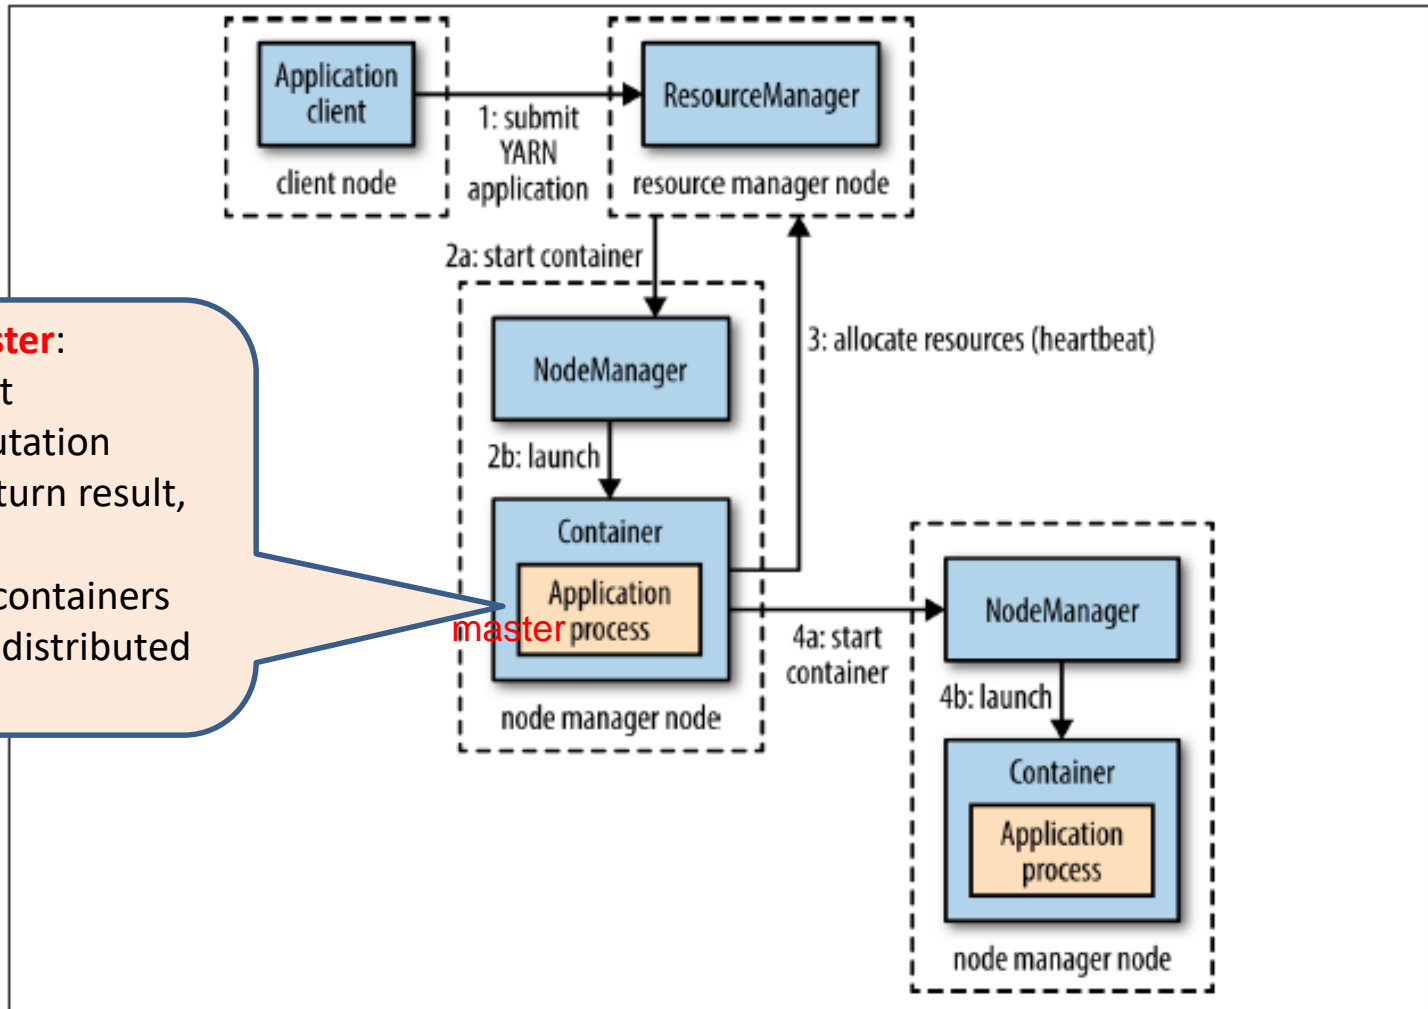


# MapReduce2 / YARN



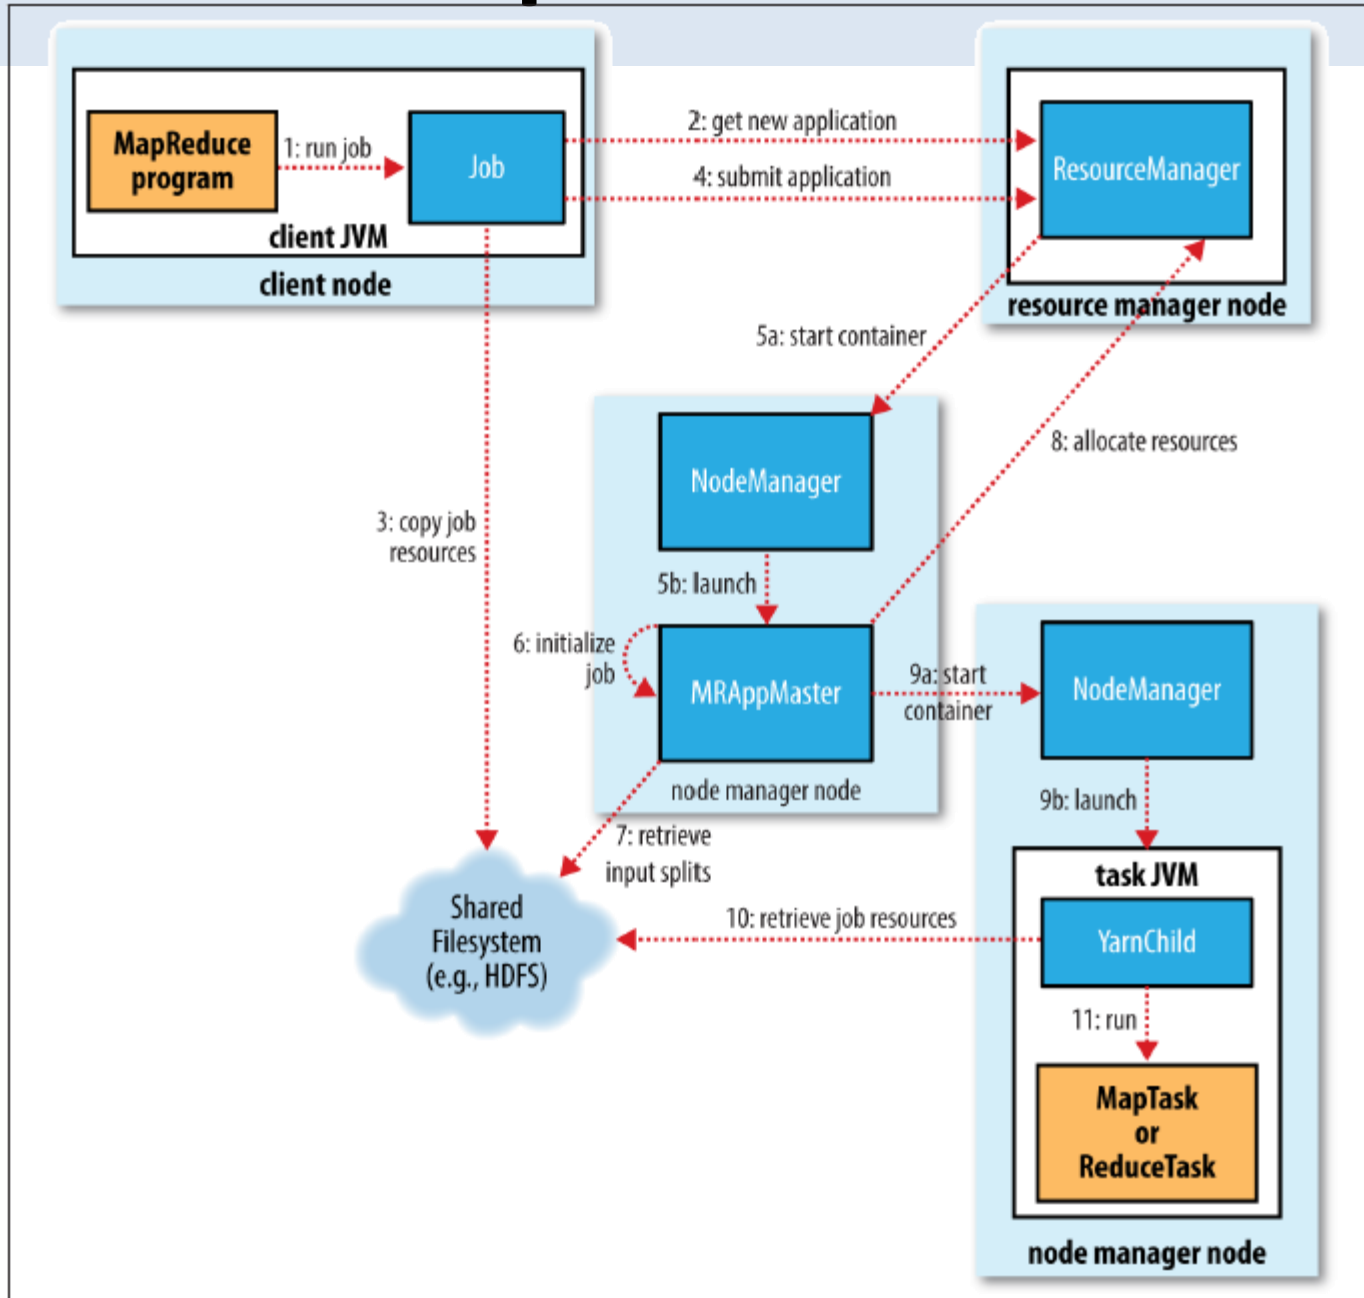
# YARN: Application Run

- For any YARN application: include MapReduce, Spark, etc.






# YARN: MapReduce Job Run




# YARN: Schedulers


## i. FIFO Scheduler



## ii. Capacity Scheduler




## iii. Fair Scheduler



# MapReduce Job

- Input data: divided into multiple *splits*
  - By default, 128 MB (HDFS block size)
- Map task:  $(K1, V1) \rightarrow \text{list}(K2, V2)$ 
  - Input: HDFS
    - Data locality: data-local, rack-local, off-rack
  - Output: local disk (why?)
  - One map task per split
    - Run map function for each *record in the split*
- Reduce task:  $(K2, \text{list}(V2)) \rightarrow \text{list}(K3, V3)$ 
  - Input: *shuffle-and-sorted map output's intermediate key-value pairs by key*
    - Each reduce task can be fed by many map tasks
  - Output: HDFS
  - # of reduce tasks: specified independently
    - **Nothing** to do with input data size
    - In driver by default, `job.setNumReduceTasks(1);`
    - OK to have 0 reduce task

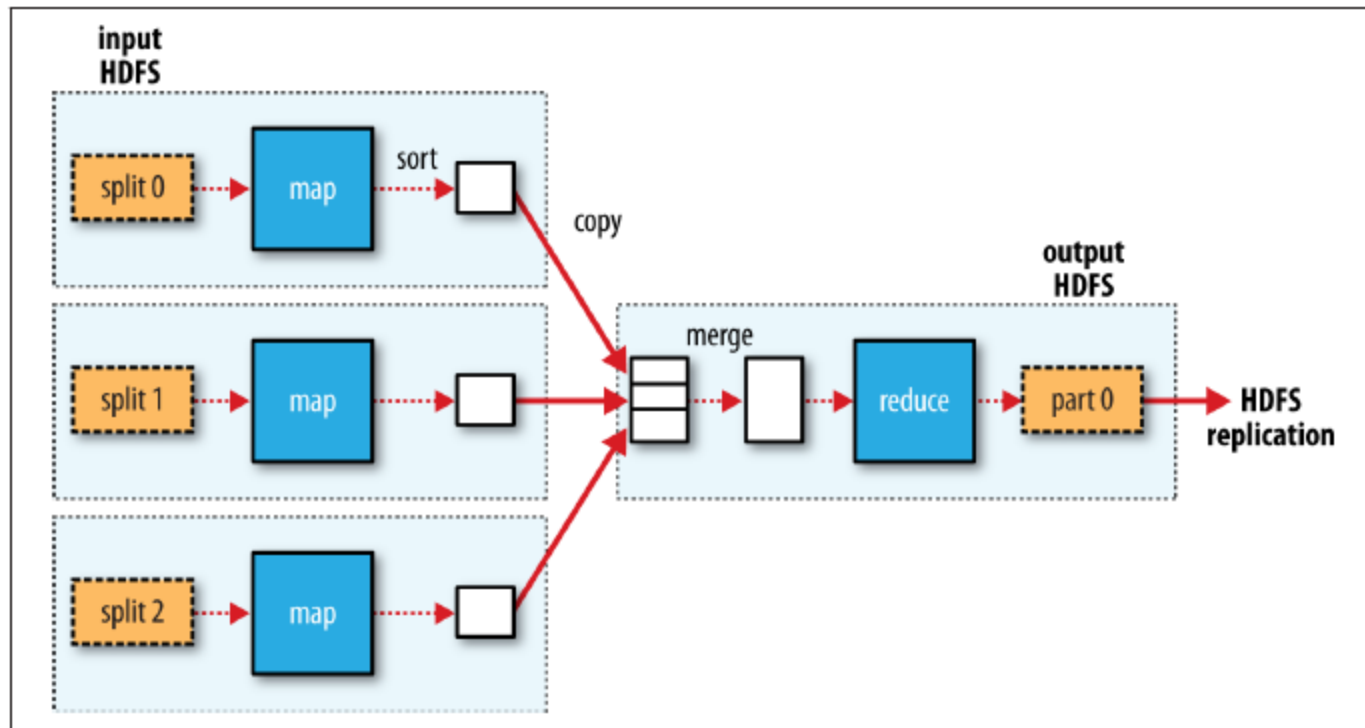


## Trade-offs:

More splits  $\rightarrow$  more map tasks,  
more parallelism, higher overhead  
Less splits  $\rightarrow$  less map tasks, more  
sequentiality, less overhead

# Data Flow: Combiner Function

- Combiner function:  $(K2, \text{list}(V2)) \rightarrow \text{list}(K2, V2)$ 
  - Optional, part of map phase
  - Often the combiner and reduce functions are the same
  - Why? **minimize the data transferred** between map and reduce tasks
  - Same MapReduce output even when no combiner function**



# Data Flow: Combiner Function (cont'd)

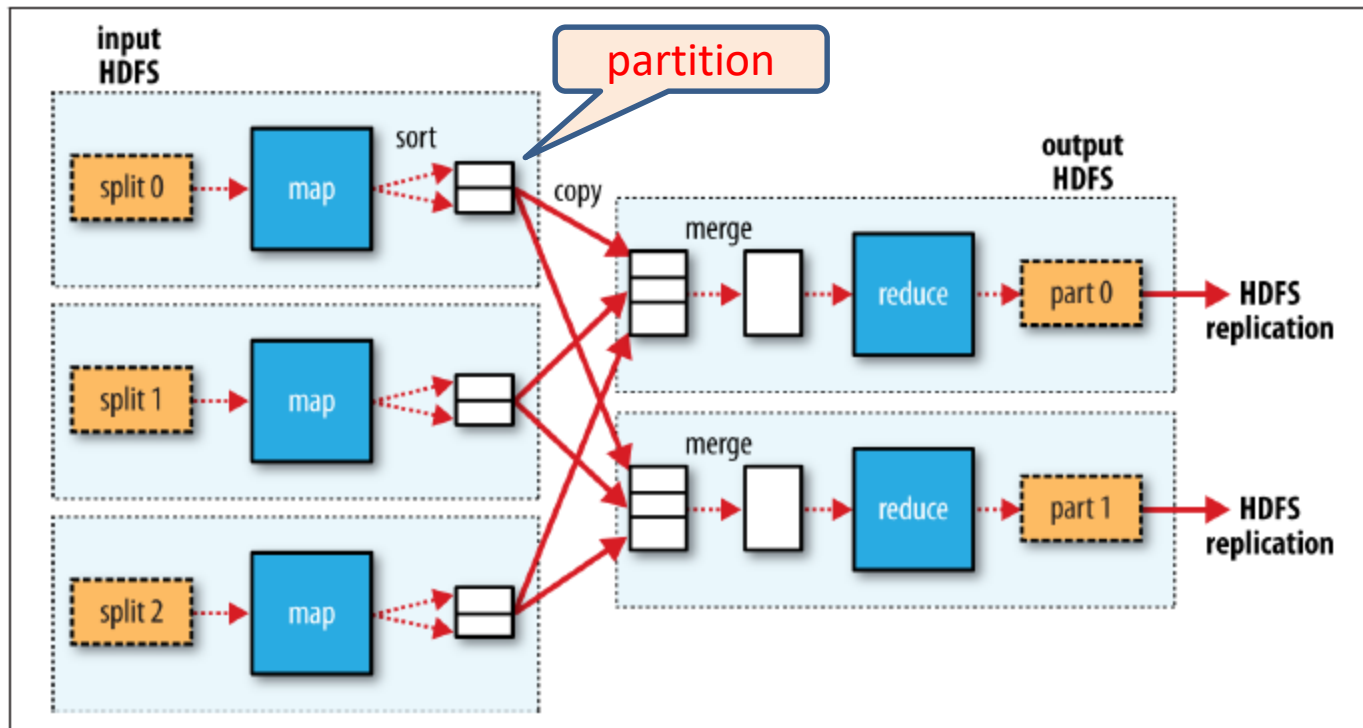
- Example: Find highest temperature in 1950s

	Without Combiner Function	With Combiner Function
Map output (two map tasks)	<ul style="list-style-type: none"><li>• (1950, 0), (1950, 20), (1950, 10)</li><li>• (1950, 25), (1950, 15)</li></ul>	<ul style="list-style-type: none"><li>• (1950, 20)</li><li>• (1950, 25)</li></ul>
Reduce input (one reduce task)	(1950, [0, 20, 10, 25, 15])	(1950, [20, 25])
Reduce output	25	25

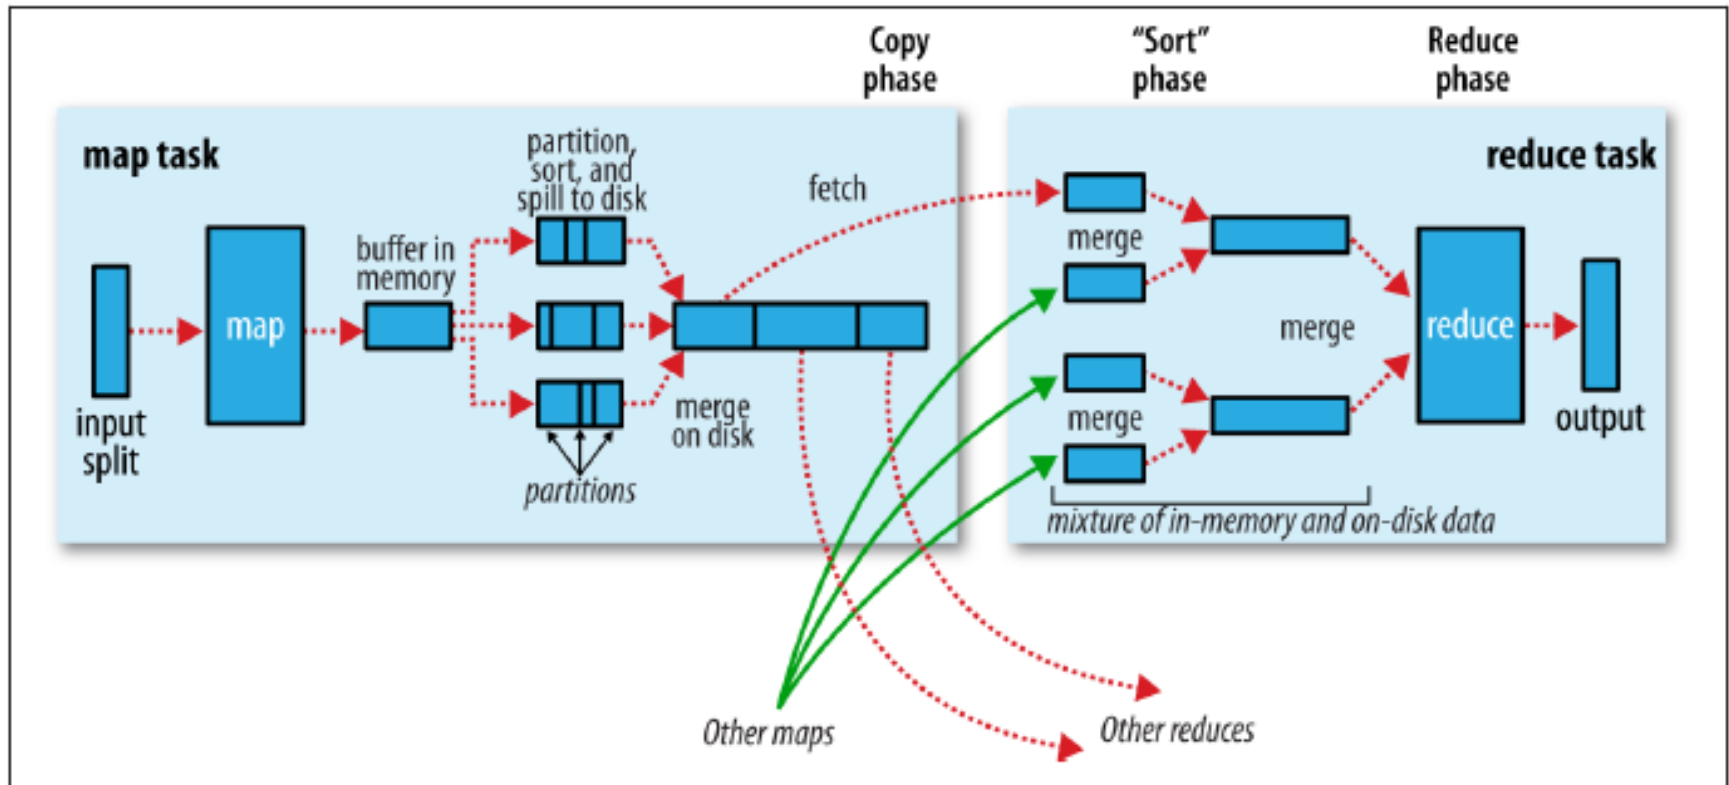
- $\max(0, 20, 10, 25, 15) = \max(\max(0, 20, 10), \max(25, 15)) = \max(20, 25) = 25$
- Note: **not** all functions can be done in this way
  - Ex:  $\text{mean}(0, 20, 10, 25, 15) = 14$   
But  $\text{mean}(\text{mean}(0, 20, 10), \text{mean}(25, 15)) = \text{mean}(10, 20) = 15$
  - Workaround?

# Data Flow: Partition Function

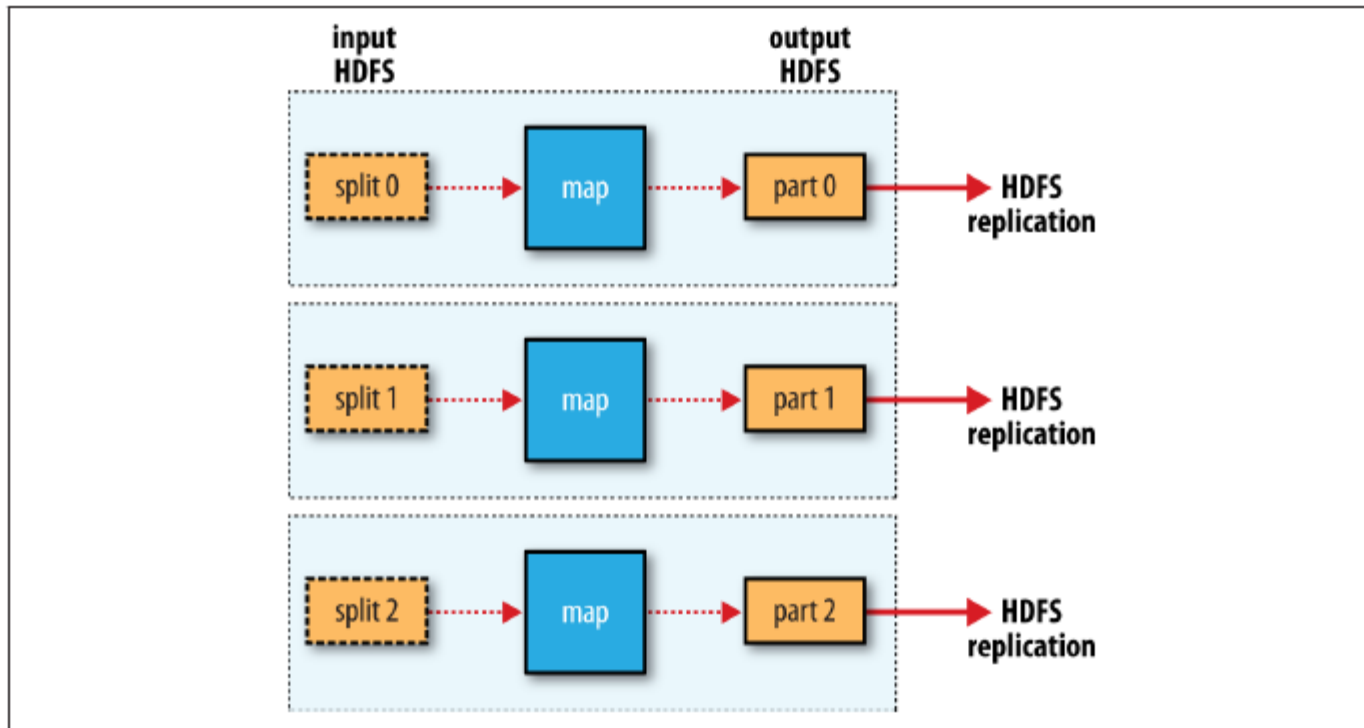
- Partition function:  $(K2, V2) \rightarrow \text{integer}$ 
  - For  $> 1$  reduce tasks
  - Optional – can overwrite the default one
  - Partition a single map task output to multiple reduce tasks
  - Records of any given key are in a single partition



# Data Flow: Shuffle-and-Sort

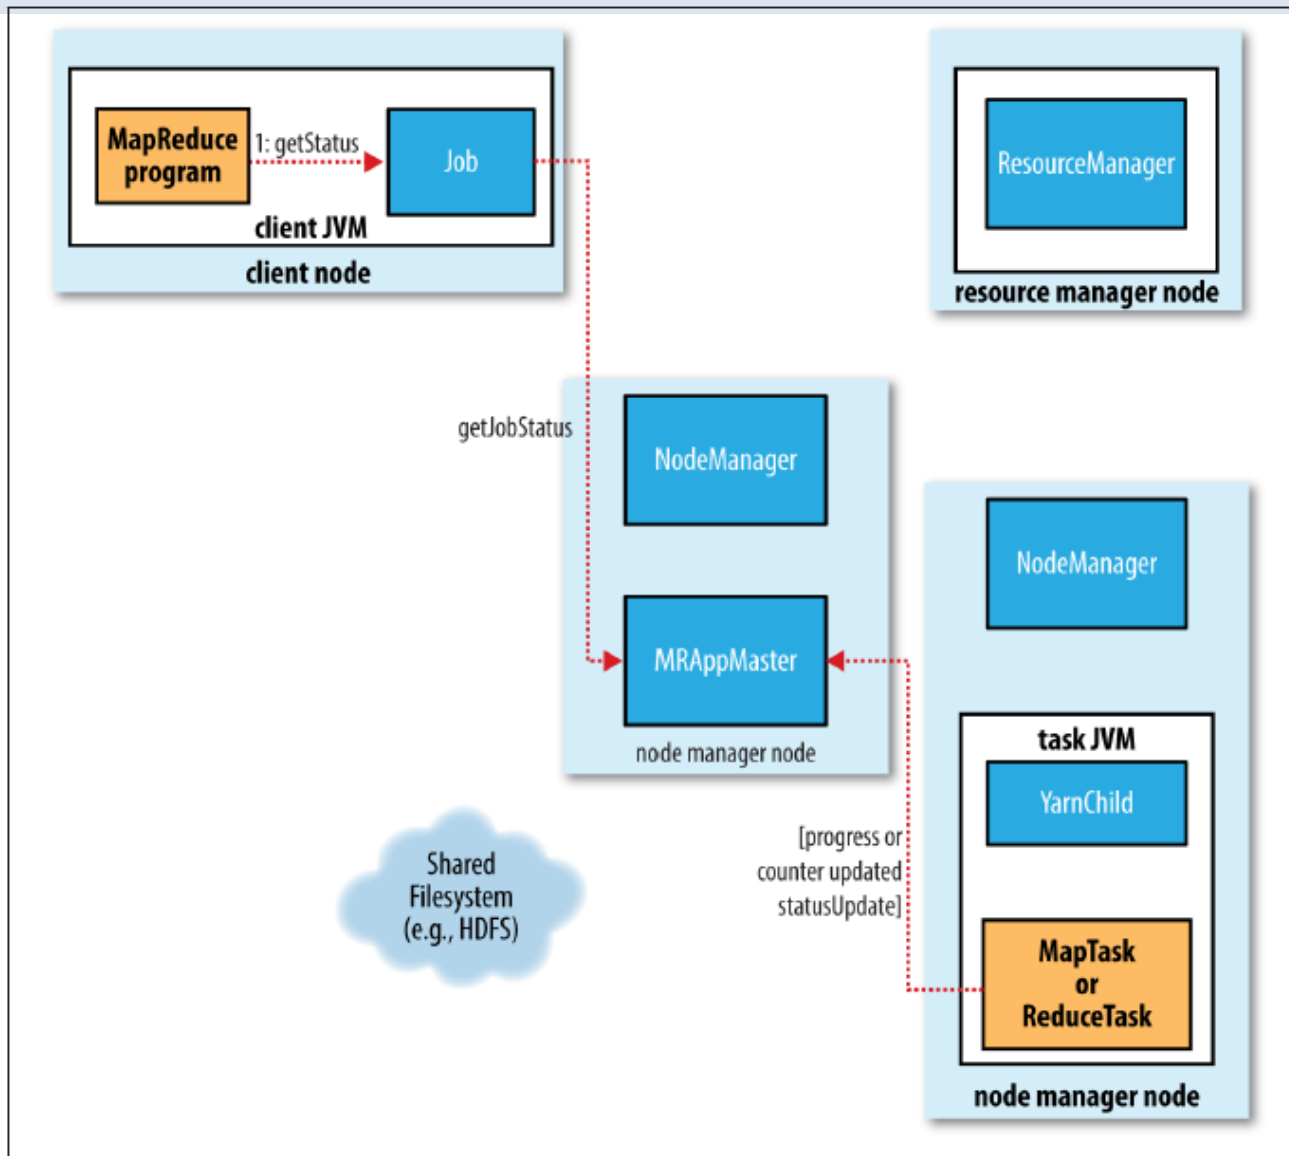


# Data Flow: no reduce task





# MapReduce Job Status Update



# FT – MR Task Failure

- Problem: what if a MapReduce task fails during a job
- Solution:
  - JobTracker (application master) notices that runtime exception, JVM exit, or timeout of progress update from the node's TaskTracker (node manager)
  - JobTracker (application master) re-schedules the failed node's task
    - Map: re-execute *all* map tasks assigned to the node
    - Reduce: re-execute *all* reduce tasks assigned to the node
      - Good enough?
      - May also need to re-execute *all* map tasks on the failed node
        - » Failed node may also have completed map tasks, and other nodes may not have finished copying out the results

# FT – Application Master Failure

- Problem: what if an application master fails
- Solution:
  - Application master periodically heartbeats resource manager
  - When timeout on heartbeat, resource manager starts a new instance of the failed master in a new container
    - If it is MapReduce master failure, use the job history (from timeline server) to recover the state of tasks; no need to rerun the entire app

# FT – Node Manager Failure

- Problem: what if an node manager fails
- Solution:
  - Node manager periodically heartbeats resource manager
  - When timeout on heartbeat, resource manager
    - Removes the failed node manager from the pool of nodes to schedule containers on
    - Recover any task or application master running on the failed node manager
      - For incomplete MapReduce job, application master re-runs map tasks that were completed successfully on the failed node manager (why?)

# FT – Resource Manager Failure

- Problem: what if a resource manager fails
- Solution:
  - Run two resource managers in active-standby configuration
  - Standby resource manager
    - recovers info of all apps from HA state store (backed by ZooKeeper or HDFS)
    - restarts all apps in the cluster
  - Clients and node managers must be configured to handle resource manager failover

# Speculative Execution

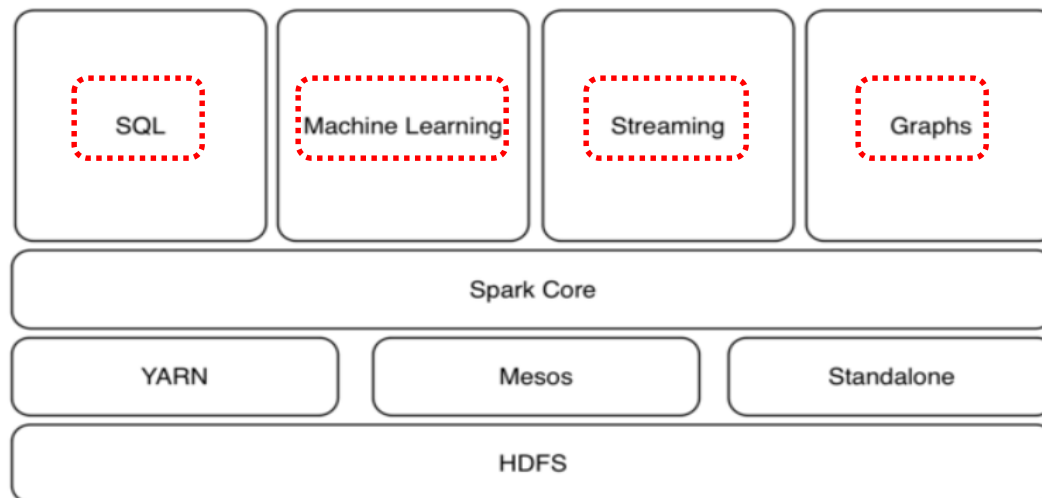
- Problem: a slow task would delay the entire job
- Tasks are executed *in isolation* from one another → the same input split can be processed by multiple nodes *in parallel*
- Solution:
  - When a job is almost done, schedule *redundant* copies of remaining tasks across several *idle* nodes
  - The one finishes first becomes the definitive copy; the others' tasks are killed (results are discarded)
- Note
  - Is an optimization, not required to make it more reliable
  - Can independently apply to map tasks and reduce tasks
  - Different from “running duplicated tasks at the same time”
    - Waste cluster resource

# MapReduce User Applications

- User provides the codes (Java, scripts, etc) for mapper, reducer, and a driver
- The rest is handled by the framework
- What problems are MapReduce good for?
  - Good for single-pass algorithm w/ filter/aggregate stages
  - Complicated one may need multi-map/reduce stages
- Are there problems one cannot solve efficiently with MapReduce?
- Are there problems it can't solve at all?

# Apache Spark


- **In-memory** analytics on large data sets (vs. read/write files in MapReduce)
- Good for
  - interactive analysis - series of ad hoc exploratory queries by a user on dataset
  - Iterative algorithms – functions applied to data set repeatedly
- Alternative **data** processing engine: Directed Acyclic Graph (DAG) engine
  - Data in RAM: better performance than Hadoop MapReduce
  - Does *not* replace Hadoop: standalone, or inside Hadoop cluster (a YARN app)
  - Does *not* have own storage / file system: can use Hadoop (HDFS, HBase, etc), or other (cloud) data platforms






# Apache Spark (cont'd)

- **Resilient Distributed Datasets (RDD)**
  - **partitioned** in multiple nodes; **in-memory *between* jobs**
  - **Immutable (read-only)**: can generate new one easily; remember dependencies
  - **Lazy transformation**: no results computed until an action requires the results
  - **recovery capability**: re-apply same transformation to RDD



- Cluster mgr: local, standalone, Mesos, YARN
- Executor: runs tasks



- API: Scala, Java, Python

# Apache Spark (cont'd)

**\$ spark-shell**

```
> val lines = sc.textFile("input/ncdc/micro-tab/sample.txt")
> val records = lines.map(_.split("\t"))
> val filtered = records.filter(rec => (rec(1) != "9999" && rec(2).matches("[01459]")))
> val tuples = filtered.map(rec => (rec(0).toInt, rec(1).toInt))
> val maxTemps = tuples.reduceByKey((a, b) => Math.max(a, b))
> maxTemps.foreach(println(_))
> maxTemps.saveAsTextFile("output")
```

**\$ hadoop fs -cat output/\***


**\$ spark-shell**

**>:paste**

```
sc.textFile("input/ncdc/micro-tab/sample.txt")
  .map(_.split("\t"))
  .filter(rec => (rec(1) != "9999" && rec(2).matches("[01459]")))
  .map(rec => (rec(0).toInt, rec(1).toInt))
  .reduceByKey((a, b) => Math.max(a, b))
  .saveAsTextFile("output2")
```

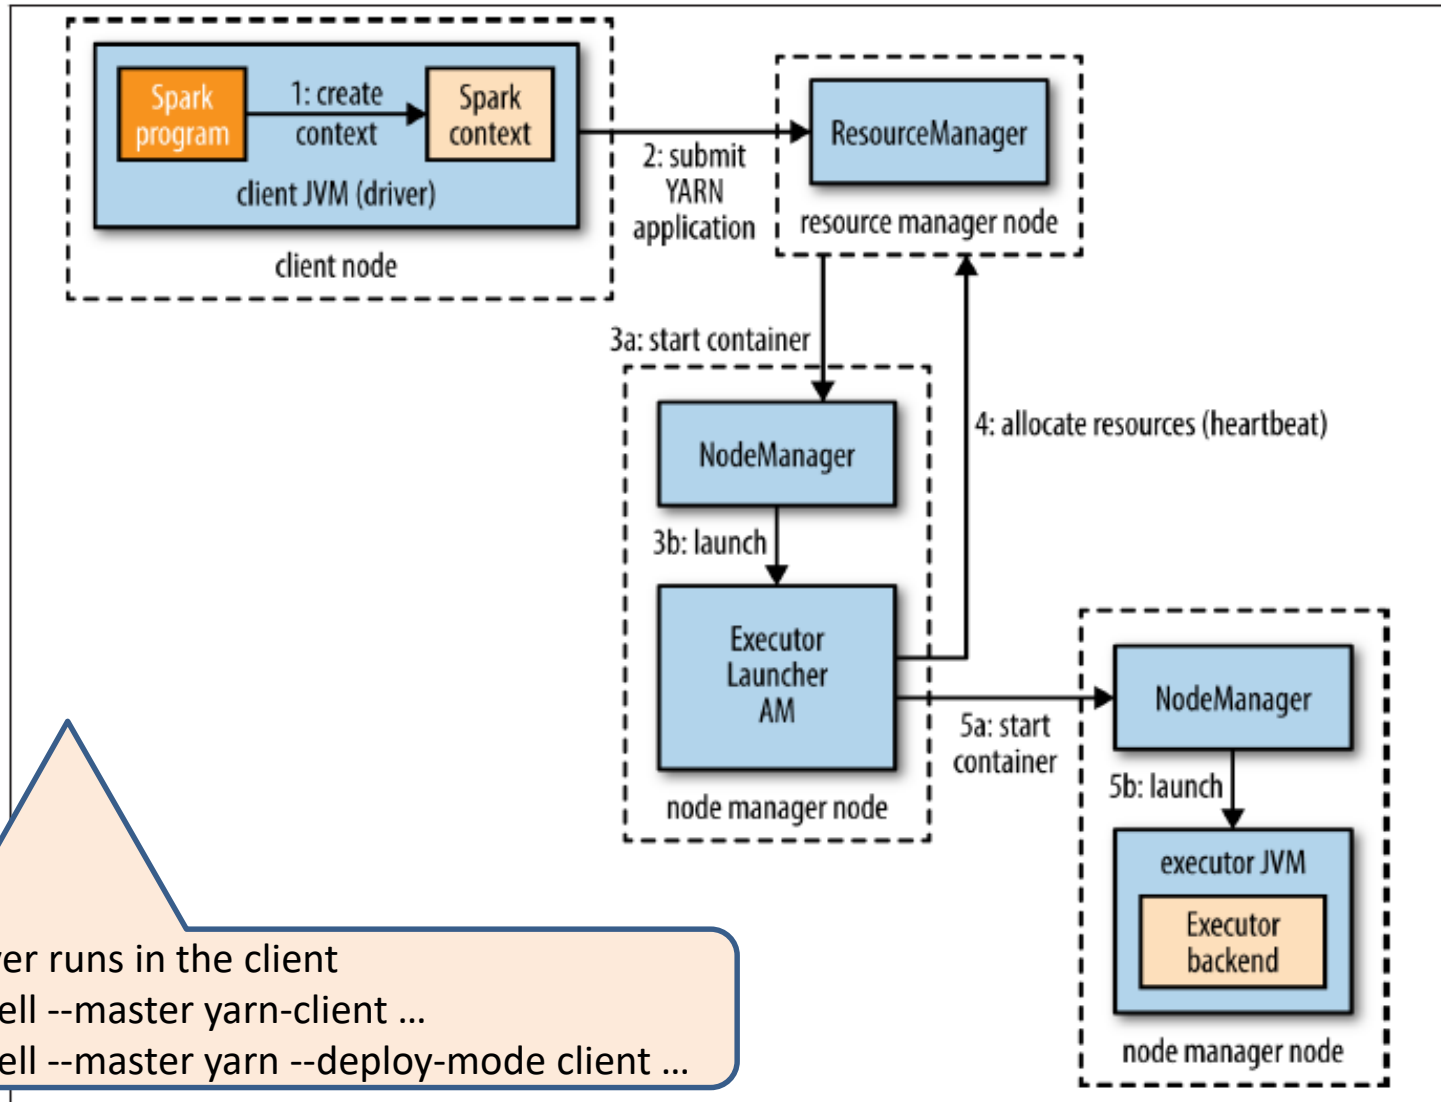
**^D (i.e., control-D)**

**\$ hadoop fs -cat output2/\***

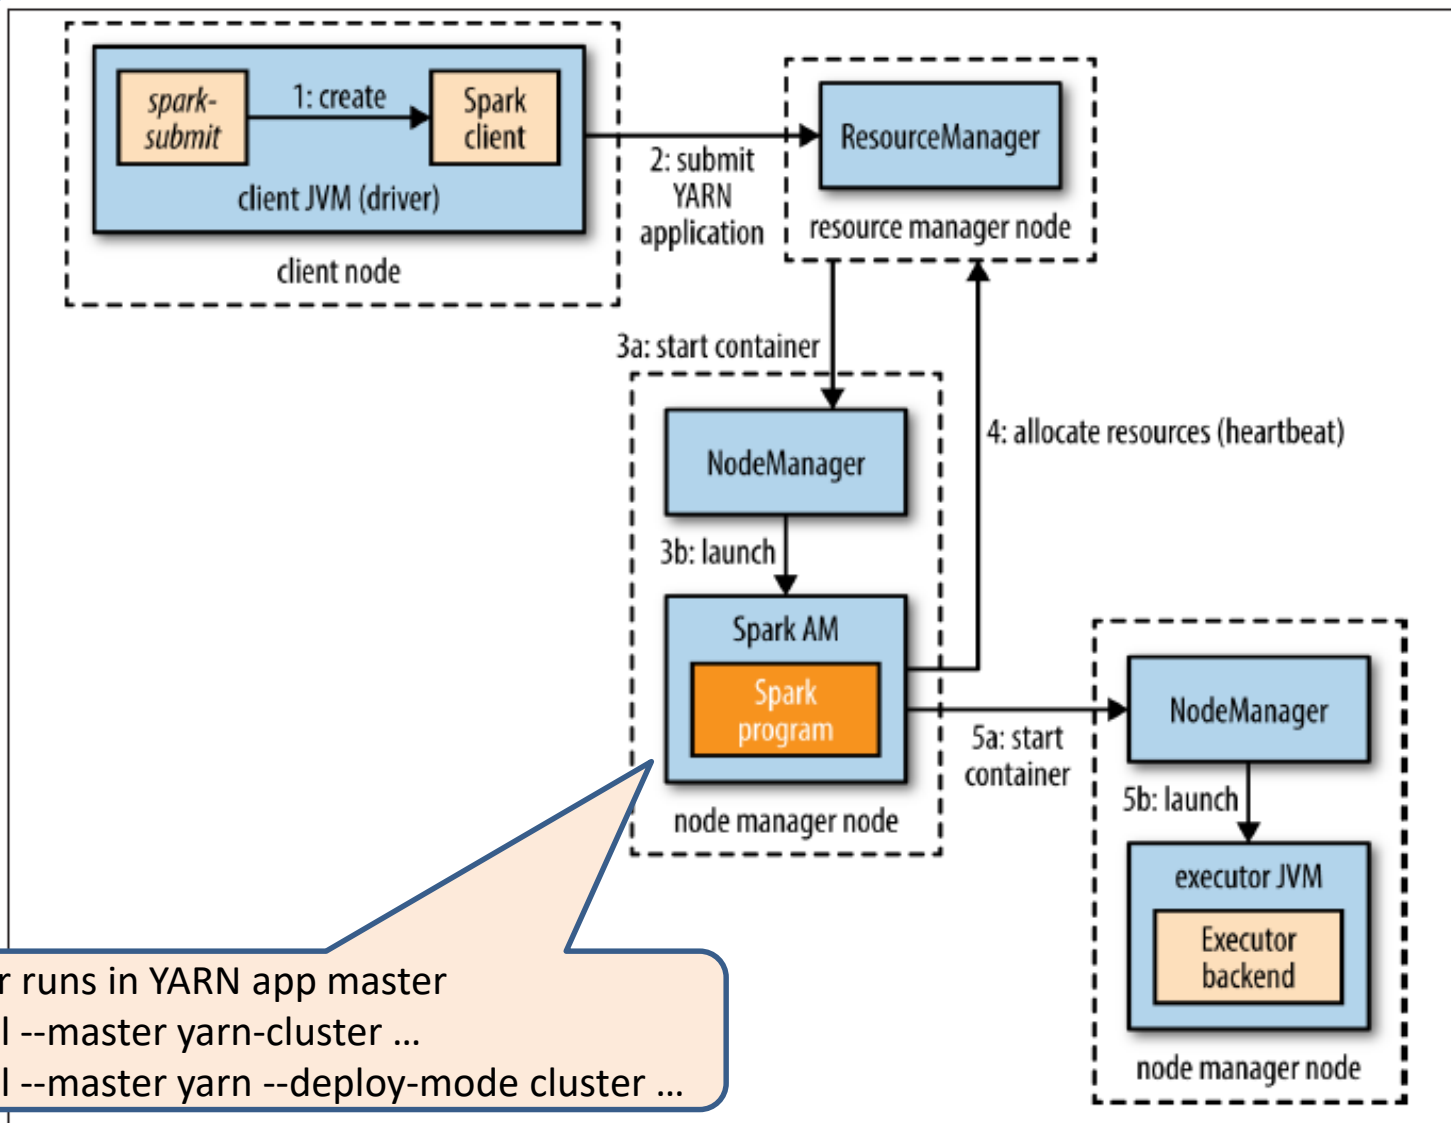


```
$ spark-submit --class MaxTemperature --master local \
spark-examples.jar input/ncdc/micro-tab/sample.txt output3
```

# Apache Spark: YARN Client Mode



# Apache Spark: YARN Cluster Mode



# Apache Storm

- Distributed computation framework for **event stream processing**

- Scalable
- **Real-time computations on continuous streams of data**
- Incremental computation (e.g., rolling avg, etc.)
- running forever until they are explicitly killed
- Fault tolerance: restart failed worker (perhaps on different node)


Hadoop & Spark: process a batch at a time  
Storm: process an event (microbatch) a time

- **No** dependency on MapReduce/Hadoop; it may

- Adapter: interop w/ HDFS

- Storm *topology*: collection of wired spouts and bolts


- *Spout*: filter/process input data stream to output *tuples*
- *Bolt*: process input *tuples*, can be fed to other bolts
- Instances of spout and bolt distributed on multi nodes in cluster



# MapReduce Services

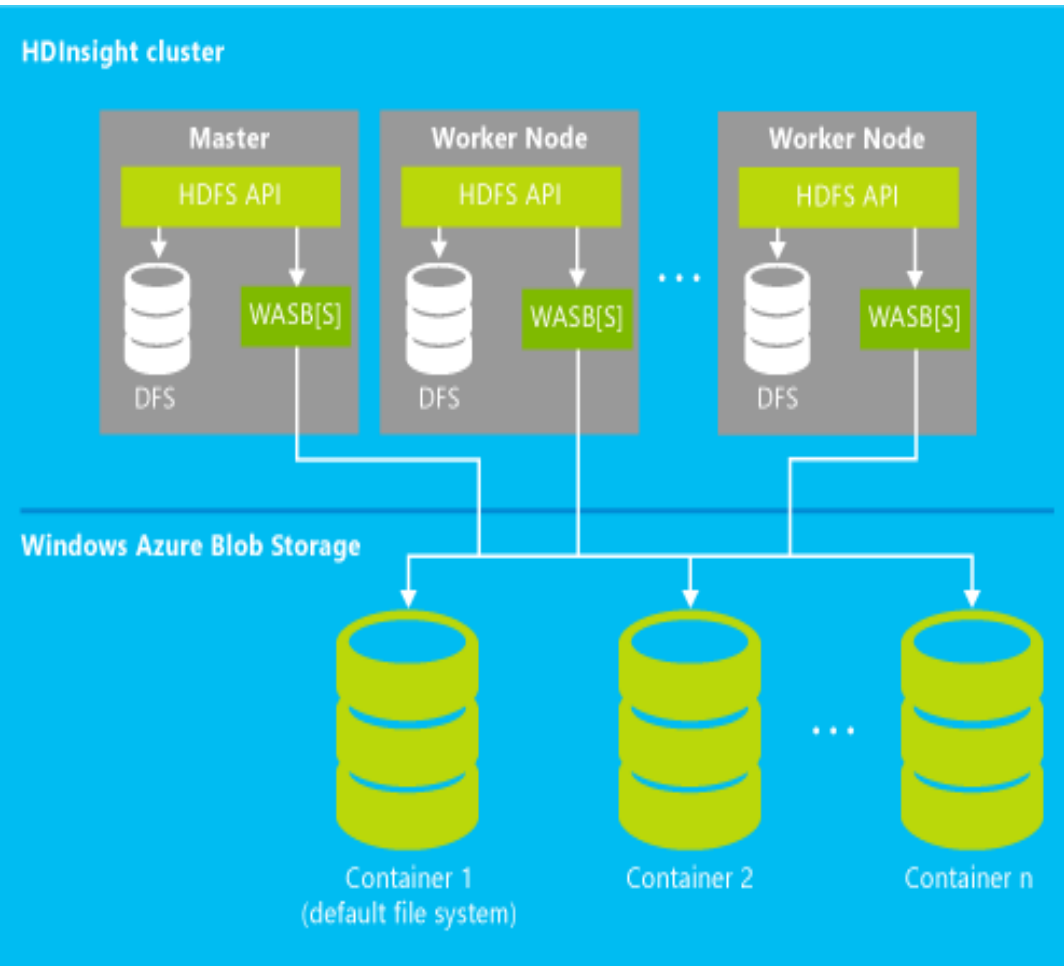
- AWS EMR: <http://aws.amazon.com/elasticmapreduce/>
- Azure HDInsight: <http://azure.microsoft.com/en-us/services/hdinsight/>
- Google Cloud DataProc: <https://cloud.google.com/dataproc/>
- Etc.
- Hadoop sandbox (VM on VMware player or VirtualBox, Docker container image): free
  - CDH: [http://www.cloudera.com/content/www/en-us/downloads/quickstart\\_vms.html](http://www.cloudera.com/content/www/en-us/downloads/quickstart_vms.html)
  - HDP: <https://hortonworks.com/downloads/#data-platform>
  - <https://www.mapr.com/products/mapr-sandbox-hadoop/download>

# AWS - Elastic MapReduce (EMR)



- Use Hadoop to distribute data and process across a cluster of EC2 instances
  - Integrated w/ S3 and CloudWatch
- Scaling up / down
  - Deploy multiple clusters
  - Resize a running cluster
- Apache Spark on EMR

# Azure - HDInsight



- Hadoop-based service
  - MapReduce
  - Various Hadoop related tools
- HA: HDInsight cluster
- query on-premises and cloud-based Hadoop clusters
  - Local to node:  
`hdfs://<namenode>/<path>`
  - Azure Blob:  
`wasb[s]://<containername>@<accountname>.blob.core.windows.net/<path>`
- Include
  - Apache Spark
  - Apache Hbase: NoSQL
  - Apache Storm: real-time stream processing



# References

- Tom White, *Hadoop: The Definite Guide*, 4/E
  - [http://bit.ly/hadoop\\_tdg\\_4e](http://bit.ly/hadoop_tdg_4e)
  - <https://github.com/tomwhite/hadoop-book/>
- Hadoop, HDFS
  - <http://hadoop.apache.org>
  - <http://hadooper.blogspot.com>
  - <http://hadoop.apache.org/docs/current/hadoop-project-dist/hadoop-hdfs/HdfsDesign.html>
  - [www.ibm.com/developerworks/library/wa-introhdfs/](http://www.ibm.com/developerworks/library/wa-introhdfs/)
- Hadoop sandbox (VM on VMware player or VirtualBox, docker container image): free
  - CDH: [http://www.cloudera.com/downloads/quickstart\\_vms.html](http://www.cloudera.com/downloads/quickstart_vms.html)
  - HDP: <https://hortonworks.com/downloads/#data-platform>
  - <https://www.mapr.com/products/mapr-sandbox-hadoop/download>
- Spark: <http://spark.apache.org/>
- Storm: <http://storm.apache.org/>