

# CMPE 282 Cloud Services

## ***MapReduce Lab***


Instructor: Kong Li

# Content

- HDFS Cmd Line Interface
- MapReduce Tutorial
- MapReduce Considerations
- Common Mistakes
- HW

# MapReduce: Refresher


- Required
  - map:  $(K1, V1) \rightarrow \text{list}(K2, V2)$
  - reduce:  $(K2, \text{list}(V2)) \rightarrow \text{list}(K3, V3)$
- Optional
  - Combiner F:  $(K2, \text{list}(V2)) \rightarrow \text{list}(K2, V2)$ 
    - Part of map phase
    - Often the combiner and reduce functions are the same
  - Partition F:  $(K2, V2) \rightarrow \text{integer}$



- HDFS: job input & output
- Driver program
- JobTracker, Task Tracker
- Resource manager, node manager

# MapReduce Job

- Input data: divided into multiple *splits*
  - By default, 128 MB (HDFS block size)
- Map task:  $(K1, V1) \rightarrow \text{list}(K2, V2)$ 
  - Input: HDFS
    - Data locality: data-local, rack-local, off-rack
  - Output: local disk (why?)
  - One map task per split
    - Run map function for each *record in the split*
- Reduce task:  $(K2, \text{list}(V2)) \rightarrow \text{list}(K3, V3)$ 
  - Input: *shuffle-and-sorted map output's intermediate key-value pairs by key*
    - Each reduce task can be fed by many map tasks
  - Output: HDFS
  - # of reduce tasks: specified independently
    - **Nothing** to do with input data size
    - In driver by default, `job.setNumReduceTasks(1);`
    - OK to have 0 reduce task



## Trade-offs:

More splits  $\rightarrow$  more map tasks,  
more parallelism, higher overhead  
Less splits  $\rightarrow$  less map tasks, more  
sequentiality, less overhead

# HDFS Cmd Line

- Copy file from local disk to HDFS

```
$ hadoop fs -copyFromLocal input/docs/quangle.txt \  
hdfs://localhost/user/tom/quangle.txt
```

\$ **hadoop -help**

```
$ hadoop fs -copyFromLocal input/docs/quangle.txt /user/tom/quangle.txt
```

```
$ hadoop fs -copyFromLocal input/docs/quangle.txt quangle.txt
```

- Copy file from HDFS to local disk

```
$ hadoop fs -copyToLocal quangle.txt quangle.copy.txt
```

- Mkdir on HDFS

```
$ hadoop fs -mkdir -p books
```

- List files on HDFS

```
$ hadoop fs -ls .
```

Found 2 items

```
drwxr-xr-x - tom supergroup 0 2014-10-04 13:22 books
```

```
-rw-r--r-- 1 tom supergroup 119 2014-10-04 13:21 quangle.txt
```

- Parallel copying

```
$ hadoop distcp file1 file2
```

# WordCount

```
import java.io.IOException;
import java.util.StringTokenizer;
import org.apache.hadoop.conf.Configuration;
import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Job;
import org.apache.hadoop.mapreduce.Mapper;
import org.apache.hadoop.mapreduce.Reducer;
import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;

public class WordCount {
    public static class TokenizerMapper
        extends Mapper<Object, Text, Text, IntWritable> {
        .....

    public void map(Object key, Text value, Context context)
        throws IOException, InterruptedException {
        StringTokenizer itr = new StringTokenizer(value.toString());
        while (itr.hasMoreTokens()) {
            context.write(new Text(itr.nextToken()), new IntWritable(1));
            .....
        }
    }
}
```

```
public static class IntSumReducer
    extends Reducer<Text, IntWritable, Text, IntWritable> {
    .....

    public void reduce(Text key, Iterable<IntWritable> values,
        Context context)
        throws IOException, InterruptedException {
        int sum = 0;
        for (IntWritable val : values) {
            sum += val.get();
        }
        context.write(key, new IntWritable(sum));
        .....
    }
}

public static void main(String[] args) throws Exception {
    Configuration conf = new Configuration();
    Job job = Job.getInstance(conf, "word count");
    job.setJarByClass(WordCount.class);
    job.setMapperClass(TokenizerMapper.class);
    job.setCombinerClass(IntSumReducer.class);
    job.setReducerClass(IntSumReducer.class);
    job.setOutputKeyClass(Text.class);
    job.setOutputValueClass(IntWritable.class);
    FileInputFormat.addInputPath(job, new Path(args[0]));
    FileOutputFormat.setOutputPath(job, new Path(args[1]));
    System.exit(job.waitForCompletion(true) ? 0 : 1);
}
}
```

# WordCount - map

```
public static class TokenizerMapper
```

keyin, valuein, keyout, valueout

```
    extends Mapper<Object, Text, Text, IntWritable>{
```

Classes must match keyin, valuein

```
    public void map(Object key, Text value, Context context)
```

```
        throws IOException, InterruptedException {
```

```
        StringTokenizer itr = new StringTokenizer(value.toString());
```

```
        while (itr.hasMoreTokens()) {
```

Classes must match  
keyout, valueout

```
            context.write(new Text(itr.nextToken()), new IntWritable(1));
```

```
        }
```

```
    }
```

```
}
```

**Input file01:**

Hello World Bye World

**Input file02:**

Hello Hadoop Goodbye Hadoop

**map 1 output:**

<Hello, 1>

<World, 1>

<Bye, 1>

<World, 1>

**map 2 output:**

<Hello, 1>

<Hadoop, 1>

<Goodbye, 1>

<Hadoop, 1>

# WordCount – map + combiner

```
public static class TokenizerMapper
```

keyin, valuein, keyout, valueout

```
extends Mapper<Object, Text, Text, IntWritable>{
```

Classes must match keyin, valuein

```
public void map(Object key, Text value, Context context)
```

```
throws IOException, InterruptedException {
```

```
StringTokenizer itr = new StringTokenizer(value.toString());
```

```
while (itr.hasMoreTokens()) {
```

Classes must match  
keyout/valueout

```
context.write(new Text(itr.nextToken()), new IntWritable(1));
```

```
}
```

```
}
```

```
}
```

**Input file01:**

Hello World Bye World

**Input file02:**

Hello Hadoop Goodbye Hadoop

**map+combiner 1**

**output:**

<Bye, 1>

<Hello, 1>

<World, 2>

**map+combiner 2**

**output:**

<Goodbye, 1>

<Hadoop, 2>

<Hello, 1>

8

Combiner's keyin/valuein must match map's keyout/valueout  
Combiner's keyout/valueout must match reduce's keyin/valuein



# WordCount - reduce

```
public static class IntSumReducer
    extends Reducer<Text, IntWritable, Text, IntWritable> {

    public void reduce(Text key, Iterable<IntWritable> values,
        Context context)
        throws IOException, InterruptedException {

        int sum = 0;
        for (IntWritable val : values) {
            sum += val.get();
        }
        context.write(key, new IntWritable(sum));
    }
}
```

keyin, valuein, keyout, valueout  
keyin/valuein must match  
map's keyout/valueout

Classes must match keyin, valuein

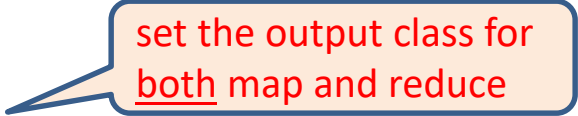
Classes must match  
keyout/valueout

## reduce output:

<Bye, 1>  
<Goodbye, 1>  
<Hadoop, 2>  
<Hello, 2>  
<World, 2>

# WordCount - driver

```
public static void main(String[] args) throws Exception {  
    Configuration conf = new Configuration();  
    Job job = Job.getInstance(conf, "word count");  
    job.setJarByClass(WordCount.class);  
  
    job.setMapperClass(TokenizerMapper.class);  
    job.setCombinerClass(IntSumReducer.class);  
    job.setReducerClass(IntSumReducer.class);  
  
    job.setOutputKeyClass(Text.class);  
    job.setOutputValueClass(IntWritable.class);  
  
    FileInputFormat.addInputPath(job, new Path(args[0]));  
    FileOutputFormat.setOutputPath(job, new Path(args[1]));  
    System.exit(job.waitForCompletion(true) ? 0 : 1);  
}
```



set the output class for  
both map and reduce

# WordCount - run

- Assume env variables JAVA\_HOME and PATH are set

```
$ export HADOOP_CLASSPATH=${JAVA_HOME}/lib/tools.jar
```

- Compile WordCount.java and create a jar

```
$ hadoop com.sun.tools.javac.Main WordCount.java
```

```
$ jar cf wc.jar WordCount*.class
```

- Assume /user/joe/wordcount/input as input, .../output as output

```
$ hadoop fs -ls /user/joe/wordcount/input/
```

```
/user/joe/wordcount/input/file01 /user/joe/wordcount/input/file02
```

```
$ hadoop fs -cat /user/joe/wordcount/input/file01
```

```
Hello World Bye World
```

```
$ hadoop fs -cat /user/joe/wordcount/input/file02
```

```
Hello Hadoop Goodbye Hadoop
```

- Run the app

```
$ hadoop jar wc.jar WordCount /user/joe/wordcount/input /user/joe/wordcount/output
```

- Output

```
$ hadoop fs -cat /user/joe/wordcount/output/part-r-00000
```

```
Bye 1
```

```
Goodbye 1
```

```
Hadoop 2
```

```
Hello 2
```

```
World 2
```

Left cmd lines can be replaced with

```
$ javac -classpath `hadoop classpath` -d . WordCount.java
```

```
$ jar cf wc.jar WordCount*.class
```

```
$ export HADOOP_CLASSPATH=wc.jar
```

```
...
```

```
$ hadoop WordCount /user/joe/wordcount/input \
/user/joe/wordcount/output
```

Map, reduce, and driver can be three separate classes – see WordCount2

# Common Mistakes to Avoid

- Mapper and reducer should be **stateless**

- **No static variables** - after `map` + `reduce` return, they should remember nothing about the processed data!
- Why: **No guarantees about which key-value pairs will be processed by which workers!**
- There are exceptions – see Top Ten pattern

```
HashMap h = new HashMap();  
map(key, value) {  
    if (h.contains(key)) {  
        h.add(key, value);  
        emit(key, "X");  
    }  
}
```

- Don't do your own I/O

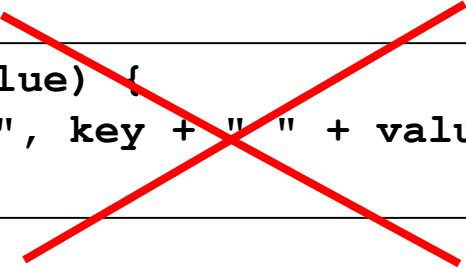
- Don't try to read from, or write to, files in the file system
- The MapReduce framework does **all** the I/O for you:

```
map(key, value) {  
    File foo =  
        new File("xyz.txt");  
    while (true) {  
        s = foo.readLine();  
        ...  
    }  
}
```

- All the incoming data will be fed as arguments to `map` and `reduce`
- **Any data your functions produce should be output via `context.write`**

# Common Mistakes to Avoid (cont'd)

```
map(key, value) {  
    emit("FOO", key + " " + value);  
}
```



```
reduce(key, value[]) {  
    /* do some computation on  
    all the values */  
}
```

- Mapper should not map too much data to the same key
  - Avoid mapping *everything* to the same key
    - Otherwise the reduce worker will be overwhelmed
  - OK if some reduce workers have more work than others
    - Ex: In WordCount, the reduce worker that works on the key 'and' has a lot more work than the reduce worker that works on 'syzygy'

# MapReduce: Considerations

- Map: execution order **not** deterministic; processing time unpredictable
- Reduce tasks cannot start before **all** Maps have finished
- **Not** suitable for continuous input streams
- Spike in network util% after Map / before Reduce phase
- Number & size of key/value pairs: Obj creation & serialization overhead
- Aggregate partial results when possible → Use combiner
- How many map tasks?
  - Smaller splits + many mappers vs larger splits + fewer mappers
    - Tradeoffs: Resource consumption, parallelism
  - Split (size) can be controlled by `InputFormat.getSplits()`, or job property `mapreduce.input.fileinputformat.split.maxsize`
- How many reduce tasks? usually determined by the algorithm
- **Locality**: Master tries to do work on nodes that have replicas of the data
- Fault tolerance: JobTracker re-executes the failed node's task(s)
- **Speculative execution**: master can deal with stragglers (slow machines) by re-executing their tasks somewhere else

# Designing MapReduce Algorithms

- **Key decision:** What should be done by `map`, and what by `reduce`?
  - `map` can do something to each **individual** key-value pair, but it can't look at **other** key-value pairs
    - Ex: Filtering out key-value pairs we don't need
  - `map` can emit **more than one** intermediate key-value pair for **each** incoming key-value pair
    - Ex: Incoming data is text line, `map` produces (word,1) for each word
  - Output value from `map` is a class which can have **several** properties
    - Ex: Map output can be (key, {min, max})
  - `reduce` can aggregate data; it can look at multiple values, as long as `map` has mapped them to the **same** (intermediate) key
    - Ex: Count the number of words, add up the total cost, ...
- Need to get the **intermediate format** right
  - If `reduce` needs to look at several values together, `map` must emit them using the **same** key
- Multiple MapReduce jobs can be **chained** together

# Minimal MapReduce driver, with defaults explicitly set

```
public class MinimalMapReduceWithDefaults extends Configured implements Tool {
    @Override
    public int run(String[] args) throws Exception {
        Job job = JobBuilder.parseInputAndOutput(this, getConf(), args);
        if (job == null) {
            return -1;
        }
        job.setInputFormatClass(TextInputFormat.class);
        job.setMapperClass(Mapper.class);
        job.setMapOutputKeyClass(LongWritable.class);
        job.setMapOutputValueClass(Text.class);
        job.setPartitionerClass(HashPartitioner.class);
        job.setNumReduceTasks(1);
        job.setReducerClass(Reducer.class);
        job.setOutputKeyClass(LongWritable.class);
        job.setOutputValueClass(Text.class);
        job.setOutputFormatClass(TextOutputFormat.class);
        return job.waitForCompletion(true) ? 0 : 1;
    }

    public static void main(String[] args) throws Exception {
        int exitCode = ToolRunner.run(new MinimalMapReduceWithDefaults(), args);
        System.exit(exitCode);
    }
}
```

each line of each input file as a separate record

`setOutputKeyClass()` and `setOutputValueClass()` set the output class for both map and reduce. If they are different, call `setMapOutputKeyClass()` and `setMapOutputValueClass()`

(key, value) pairs as "key \tvalue" on individual lines of a text file



# CMPE vCenter Server Lab Rules

- You have permission only on
  - Resource Pool: **CMPE282 SEC1**
  - VM folder: **CMPE LABS/CMPE282 SEC1/workspace**
  - Datastore: **Classroom**
- Naming convention for any newly-created VM/template/vApp
  - Naming convention: **<YourName>-<os><version>-<L3SID>** , e.g., john-ub1404-123
  - If necessary (avoid more naming collision), append **-1**, **-2**, etc at the end
- VM creation rules:
  - You can create VM only based on template, ISO file, or OVF/OVA **provided by the instructor**
    - Any created VM must connect to a specified network **without** internet access, unless allowed by instructor
  - You are **not** allowed to create VMs based on **your** uploaded ISO file or OVF/OVA
- Connection to vCenter Server and VM console
  - Web client: supported browser + several plug-ins (client support + remote console)
  - VM console (Web client or VI client): require SJSU [VPN](#)
  - ssh
- It is a shared environment
  - Be responsible - never disrupt other users
  - Clean up (**power off** or **delete** VM) as soon as you finish – penalty if you fail to do so
  - **Any malicious action will face discipline**

Your own area: a folder  
**YourName-L3SID**  
under workspace

Login acct:  
**vsphere.local\cmpe282\_sec1\_student**

# HW

- See Canvas for details

# References

- <http://hadoop.apache.org/docs/current/hadoop-mapreduce-client/hadoop-mapreduce-client-core/MapReduceTutorial.html>
- <https://github.com/tomwhite/hadoop-book/archive/master.zip>
- <https://github.com/adamjshook/mapreducepatterns/archive/master.zip>
  - Get pom.xml from <https://pragmaticintegrator.wordpress.com/2014/09/09/running-mapreduce-design-patterns-on-cloudera-cdh5/>
- Cloudera Quickstart doc  
<https://www.cloudera.com/documentation/enterprise/5-10-x/topics/quickstart.html>