

CMPE 282 Cloud Services

***MapReduce Design Patterns –
Data Organization***

Instructor: Kong Li

Content

- What and Why
- MapReduce refresher
- Summarization Patterns
- Filtering Patterns
- Data Organization Patterns
- Join Patterns

Data Organization Patterns

What: Reorganize, restructure

Why: I want to change the way my data is organized

- **Structured to hierarchical**
- **Partitioning**
- **Binning**
- Total order sorting
- Shuffling

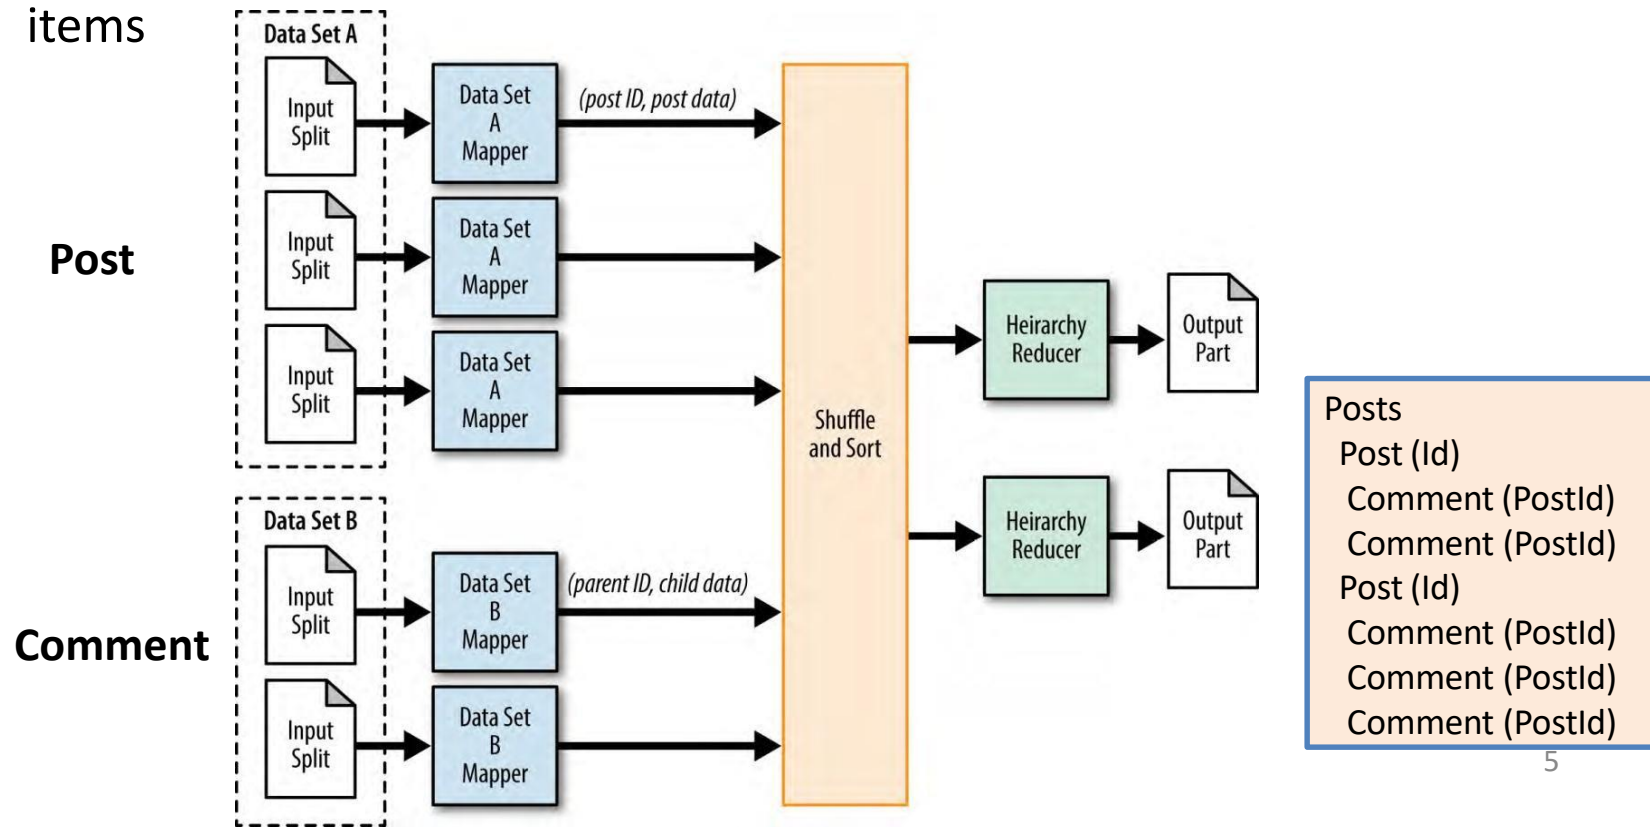
Structured to Hierarchical 1/5

- Intent
 - Transform row-based data to a hierarchical format (JSON or XML)
- Motivation
 - Migrating data from an RDBMS to Hadoop
 - Reformatting data into a more conducive structure
- Applicability
 - You have data sources that are linked by FKs
 - Your data is structured and row-based

Posts
Post
Comment
Comment
Post
Comment
Comment
Comment

Structured to Hierarchical 2/5

- Structure:
 - Mapper: load data and parse records into one cohesive format
 - Combiner: little use
 - Reducer: build the hierarchical data structure from the list of data items



Structured to Hierarchical 3/5

- Known uses
 - Pre-joining data
 - Preparing data for HBase or MongoDB
- Performance analysis
 - How much data is being sent from mappers to reducers?
 - **Almost all** data is moved across network
 - The memory footprint of the obj that the reducer builds
 - For a post that has a million comments?

Structured to Hierarchical 4/5

- PostCommentBuildingDriver.java: Given a list of posts and comments, create corresponding XML hierarchy
- In-1: Posts.xml
- In-2: Comments.xml

```
public static void main(String[] args) throws Exception {
    Configuration conf = new Configuration();
    Job job = new Job(conf, "PostCommentHierarchy");
    job.setJarByClass(PostCommentBuildingDriver.class);

    MultipleInputs.addInputPath(job, new Path(args[0]),
        TextInputFormat.class, PostMapper.class);

    MultipleInputs.addInputPath(job, new Path(args[1]),
        TextInputFormat.class, CommentMapper.class);

    job.setReducerClass(UserJoinReducer.class);

    job.setOutputFormatClass(TextOutputFormat.class);
    TextOutputFormat.setOutputPath(job, new Path(args[2]));

    job.setOutputKeyClass(Text.class);
    job.setOutputValueClass(Text.class);

    System.exit(job.waitForCompletion(true) ? 0 : 2);
}
```

```
public static class PostMapper extends Mapper<Object, Text, Text, Text> {

    private Text outkey = new Text();
    private Text outvalue = new Text();

    public void map(Object key, Text value, Context context)
        throws IOException, InterruptedException {

        Map<String, String> parsed = MRDPUtils.transformXmlToMap(value
            .toString());

        // The foreign join key is the post ID
        outkey.set(parsed.get("Id"));

        // Flag this record for the reducer and then output
        outvalue.set("P" + value.toString());
        context.write(outkey, outvalue);
    }
}

public static class CommentMapper extends Mapper<Object, Text, Text, Text> {

    private Text outkey = new Text();
    private Text outvalue = new Text();

    public void map(Object key, Text value, Context context)
        throws IOException, InterruptedException {

        Map<String, String> parsed = MRDPUtils.transformXmlToMap(value
            .toString());

        // The foreign join key is the post ID
        outkey.set(parsed.get("PostId"));

        // Flag this record for the reducer and then output
        outvalue.set("C" + value.toString());
        context.write(outkey, outvalue);
    }
}
```

Structured to Hierarchical 5/5

```
public static class PostCommentHierarchyReducer extends
    Reducer<Text, Text, Text, NullWritable> {

    private ArrayList<String> comments = new ArrayList<String>();
    private DocumentBuilderFactory dbf = DocumentBuilderFactory.newInstance();
    private String post = null;

    public void reduce(Text key, Iterable<Text> values, Context context)
        throws IOException, InterruptedException {
        // Reset variables
        post = null;
        comments.clear();

        // For each input value
        for (Text t : values) {
            // If this is the post record, store it, minus the flag
            if (t.charAt(0) == 'P') {
                post = t.toString().substring(1, t.toString().length())
                    .trim();
            } else {
                // Else, it is a comment record. Add it to the list, minus
                // the flag
                comments.add(t.toString()
                    .substring(1, t.toString().length()).trim());
            }
        }
        // If there are no comments, the comments list will simply be empty.

        // If post is not null, combine post with its comments.
        if (post != null) {
            // nest the comments underneath the post element
            String postWithCommentChildren = nestElements(post, comments);

            // write out the XML
            context.write(new Text(postWithCommentChildren),
                NullWritable.get());
        }
    }
}
```

```
private String nestElements(String post, List<String> comments) {
    // Create the new document to build the XML
    DocumentBuilder bldr = dbf.newDocumentBuilder();
    Document doc = bldr.newDocument();

    // Copy parent node to document
    Element postEl = getXmlElementFromString(post);
    Element toAddPostEl = doc.createElement("post");

    // Copy the attributes of the original post element to the new one
    copyAttributesToElement(postEl.getAttributes(), toAddPostEl);

    // For each comment, copy it to the "post" node
    for (String commentXml : comments) {
        Element commentEl = getXmlElementFromString(commentXml);
        Element toAddCommentEl = doc.createElement("comments");

        // Copy the attributes of the original comment element to
        // the new one
        copyAttributesToElement(commentEl.getAttributes(),
            toAddCommentEl);

        // Add the copied comment to the post element
        toAddPostEl.appendChild(toAddCommentEl);
    }

    // Add the post element to the document
    doc.appendChild(toAddPostEl);

    // Transform the document into a String of XML and return
    return transformDocumentToString(doc);
}

private Element getXmlElementFromString(String xml) {
    // Create a new document builder
    DocumentBuilder bldr = dbf.newDocumentBuilder();

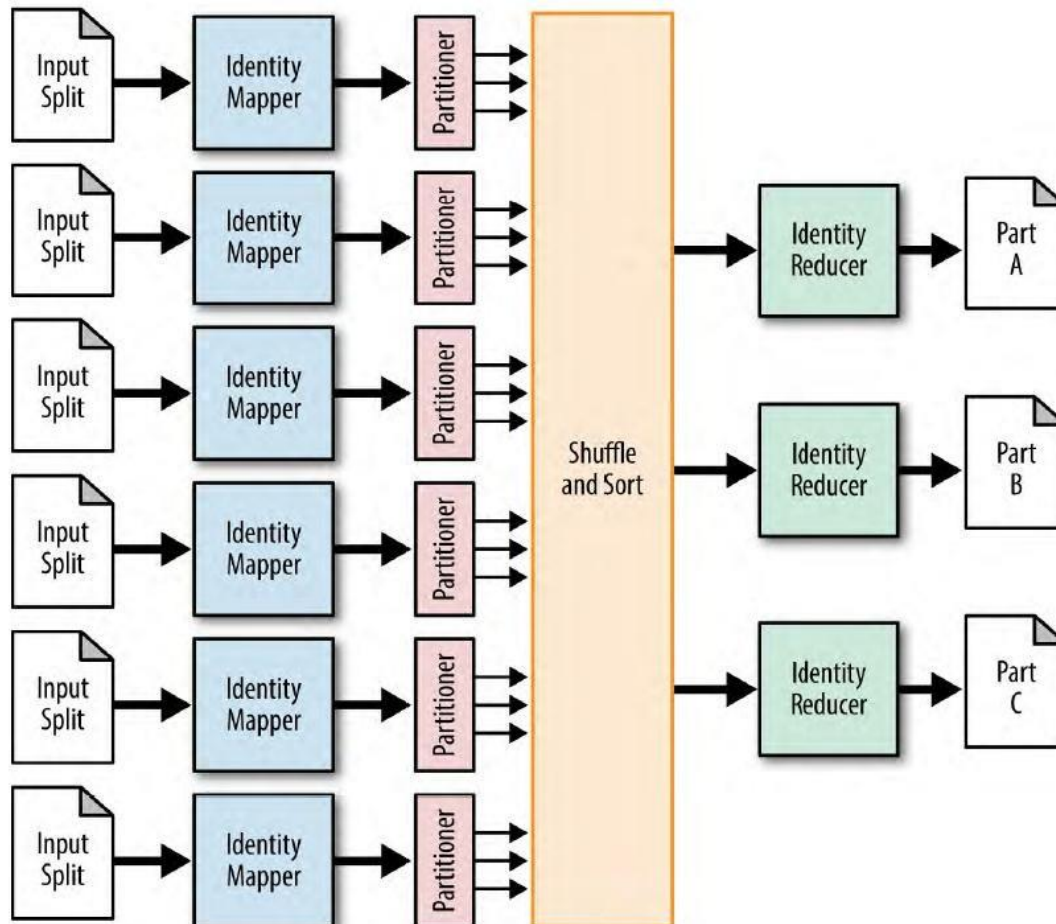
    return bldr.parse(new InputSource(new StringReader(xml)))
        .getDocumentElement();
}
```


Partitioning 1/6

- Intent
 - Move records into categories; doesn't care the order of records
 - Take similar records in a data set and partition them into distinct, smaller data sets
- Motivation
 - If you want to look at a particular set of data, the data items are normally spread out across the entire data set → requires an **entire scan** of all of the data
- Applicability
 - Knowing how many partitions **ahead of time**
 - Ex: by day of the week → 7 partitions


Partitioning 2/6

- Structure
 - exploits the fact that the partitioner partitions data
 - **Custom partitioner**: determine the partition for each record



Partitioning 3/6

- Known uses
 - **Partition pruning** by continuous value (e.g., date)
 - analyzes FROM and WHERE clauses in SQL statements to eliminate unneeded partitions when building the partition access list
 - Partition pruning by category
 - Country, phone area code, language
 - Sharding
 - **Shard**: a single server or replica set holding a part of sharded collection
- Performance analysis
 - The resulting partitions may not have similar # of records
 - **Heavy workload for Reducers with high % of records**
 - Workarounds:
 - Split very large partitions into several smaller partitions, even if just randomly
 - Assign multiple reducers to one partition and then randomly assign records into each to spread it out a bit better
 - Heavy traffic in shuffle-exchange network



Partitioning 4/6

- PartitionedUsers.java
: Given a set of user info, partition the records based on the year of last access date, one partition per year (2008 - 2011)
- In: Users.xml
- Fix: extend to 2016

```
public static class LastAccessDateMapper extends
    Mapper<Object, Text, IntWritable, Text> {

    // This object will format the creation date string into a Date object
    private final static SimpleDateFormat frmt = new SimpleDateFormat(
        "yyyy-MM-dd'T'HH:mm:ss.SSS");

    private IntWritable outkey = new IntWritable();

    protected void map(Object key, Text value, Context context)
        throws IOException, InterruptedException {

        Map<String, String> parsed = MRDPUtils.transformXmlToMap(value
            .toString());

        // Grab the last access date
        String strDate = parsed.get("LastAccessDate");

        // Parse the string into a Calendar object
        Calendar cal = Calendar.getInstance();
        cal.setTime(frmt.parse(strDate));
        outkey.set(cal.get(Calendar.YEAR));

        // Write out the year with the input value
        context.write(outkey, value);
    }
}
```

Partitioning 5/6

- Driver:

```
...  
// Set custom partitioner and min last access date  
job.setPartitionerClass(LastAccessDatePartitioner.class);  
LastAccessDatePartitioner.setMinLastAccessDate(job, 2008);  
  
// Last access dates span between 2008-2011, or 4 years  
job.setNumReduceTasks(4);  
...
```

- Reducer:

```
public static class ValueReducer extends  
    Reducer<IntWritable, Text, Text, NullWritable> {  
  
    protected void reduce(IntWritable key, Iterable<Text> values,  
        Context context) throws IOException, InterruptedException {  
        for (Text t : values) {  
            context.write(t, NullWritable.get());  
        }  
    }  
}
```

Partitioning 6/6

```
public static class LastAccessDatePartitioner extends
    Partitioner<IntWritable, Text> implements Configurable {

    private static final String MIN_LAST_ACCESS_DATE_YEAR =
        "min.last.access.date.year";

    private Configuration conf = null;
    private int minLastAccessDateYear = 0;

    public int getPartition(IntWritable key, Text value, int numPartitions) {
        return key.get() - minLastAccessDateYear;
    }

    public Configuration getConf() {
        return conf;
    }

    public void setConf(Configuration conf) {
        this.conf = conf;
        minLastAccessDateYear = conf.getInt(MIN_LAST_ACCESS_DATE_YEAR, 0);
    }

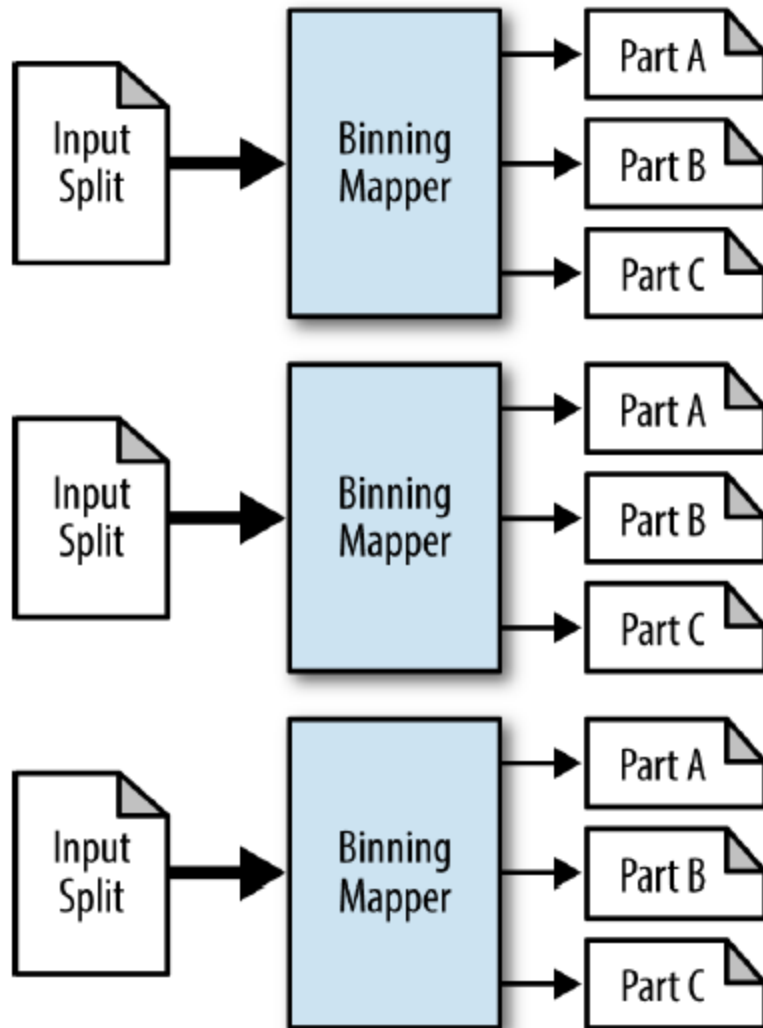
    public static void setMinLastAccessDate(Job job,
        int minLastAccessDateYear) {
        job.getConfiguration().setInt(MIN_LAST_ACCESS_DATE_YEAR,
            minLastAccessDateYear);
    }
}
```

2008 → 0
2009 → 1
2010 → 2
2011 → 3

Binning 1/4

- Intent
 - For each record in the data set, file each one into **one or more** categories
- Motivation
 - Binning: similar to partitioning and often can be used to solve the same problem
 - **Binning splits data up in the map phase instead of in the partitioner**
 - Possibility of handling **dynamic** # of categories
 - Each mapper has one file per possible output bin
 - 1000 Bins x 1000 Mappers = 1000,000 files

Binning 2/4



- Structure
 - Mapper: if the record meets the criteria, it is sent to that bin
 - No combiner, partitioner, or reducer
- Performance analysis
 - Each mapper outputs one **small** file per bin
 - map-only jobs scalability & performance
 - No sort, shuffle, or reduce to be performed
 - Most of the processing done on data that is **local**

Binning 3/4

- Binning.java: Given a set of StackOverflow posts, bin the posts into four bins based on the **tags** hadoop, pig, hive, and hbase. Also, create a separate bin for posts mentioning hadoop in the text or title

- In: Posts.xml

```
...  
// Configure the MultipleOutputs by adding an output called "bins"  
// With the proper output format and mapper key/value pairs  
MultipleOutputs.addNamedOutput(job, "bins", TextOutputFormat.class,  
                                Text.class, NullWritable.class);  
  
// Enable the counters for the job  
// If there are a significant number of different named outputs, this  
// should be disabled  
MultipleOutputs.setCountersEnabled(job, true);  
  
// Map-only job  
job.setNumReduceTasks(0);  
...
```

- MultipleOutputs
 - writing to **additional** outputs other than the job default output
 - Each additional output, or **named output**, may be configured with its own OutputFormat, with its own key class and with its own value class
 - void write(String namedOutput, K key, V value, String baseOutputPath)
Write key and value to baseOutputPath using the namedOutput

Binning 4/4

```
public static class BinningMapper extends
    Mapper<Object, Text, Text, NullWritable> {

    private MultipleOutputs<Text, NullWritable> mos = null;

    protected void setup(Context context) {
        // Create a new MultipleOutputs using the context object
        mos = new MultipleOutputs(context);
    }

    protected void map(Object key, Text value, Context context)
        throws IOException, InterruptedException {

        Map<String, String> parsed = MRDPUtils.transformXmlToMap(value
            .toString());

        String rawtags = parsed.get("Tags");

        // Tags are delimited by ><. i.e. <tag1><tag2><tag3>
        String[] tagTokens = StringEscapeUtils.unescapeHtml(rawtags).split(
            "><");

        // For each tag
        for (String tag : tagTokens) {
            // Remove any > or < from the token
            String groomed = tag.replaceAll(">|<", "").toLowerCase();

            // If this tag is one of the following, write to the named bin
            if (groomed.equalsIgnoreCase("hadoop")) {

                mos.write("bins", value, NullWritable.get(), "hadoop-tag");
            }
            if (groomed.equalsIgnoreCase("pig")) {
                mos.write("bins", value, NullWritable.get(), "pig-tag");
            }
            if (groomed.equalsIgnoreCase("hive")) {
                mos.write("bins", value, NullWritable.get(), "hive-tag");
            }
            if (groomed.equalsIgnoreCase("hbase")) {
                mos.write("bins", value, NullWritable.get(), "hbase-tag");
            }
        }

        // Get the body of the post
        String post = parsed.get("Body");

        // If the post contains the word "hadoop", write it to its own bin
        if (post.toLowerCase().contains("hadoop")) {
            mos.write("bins", value, NullWritable.get(), "hadoop-post");
        }
    }

    protected void cleanup(Context context) throws IOException,
        InterruptedException {
        // Close multiple outputs!
        mos.close();
    }
}
```

References

- Donald Miner and Adam Shook, *MapReduce Design Patterns*.
 - <http://oreil.ly/mapreduce-design-patterns>
 - <https://github.com/adamjshook/mapreducepatterns>